

Integrating a Multipath Transport Protocol in my Application

Challenges to Get Benefits

Quentin De Coninck (UCLouvain, Belgium) - November 24, 2022 - WOS'22
<https://qdeconinck.github.io>

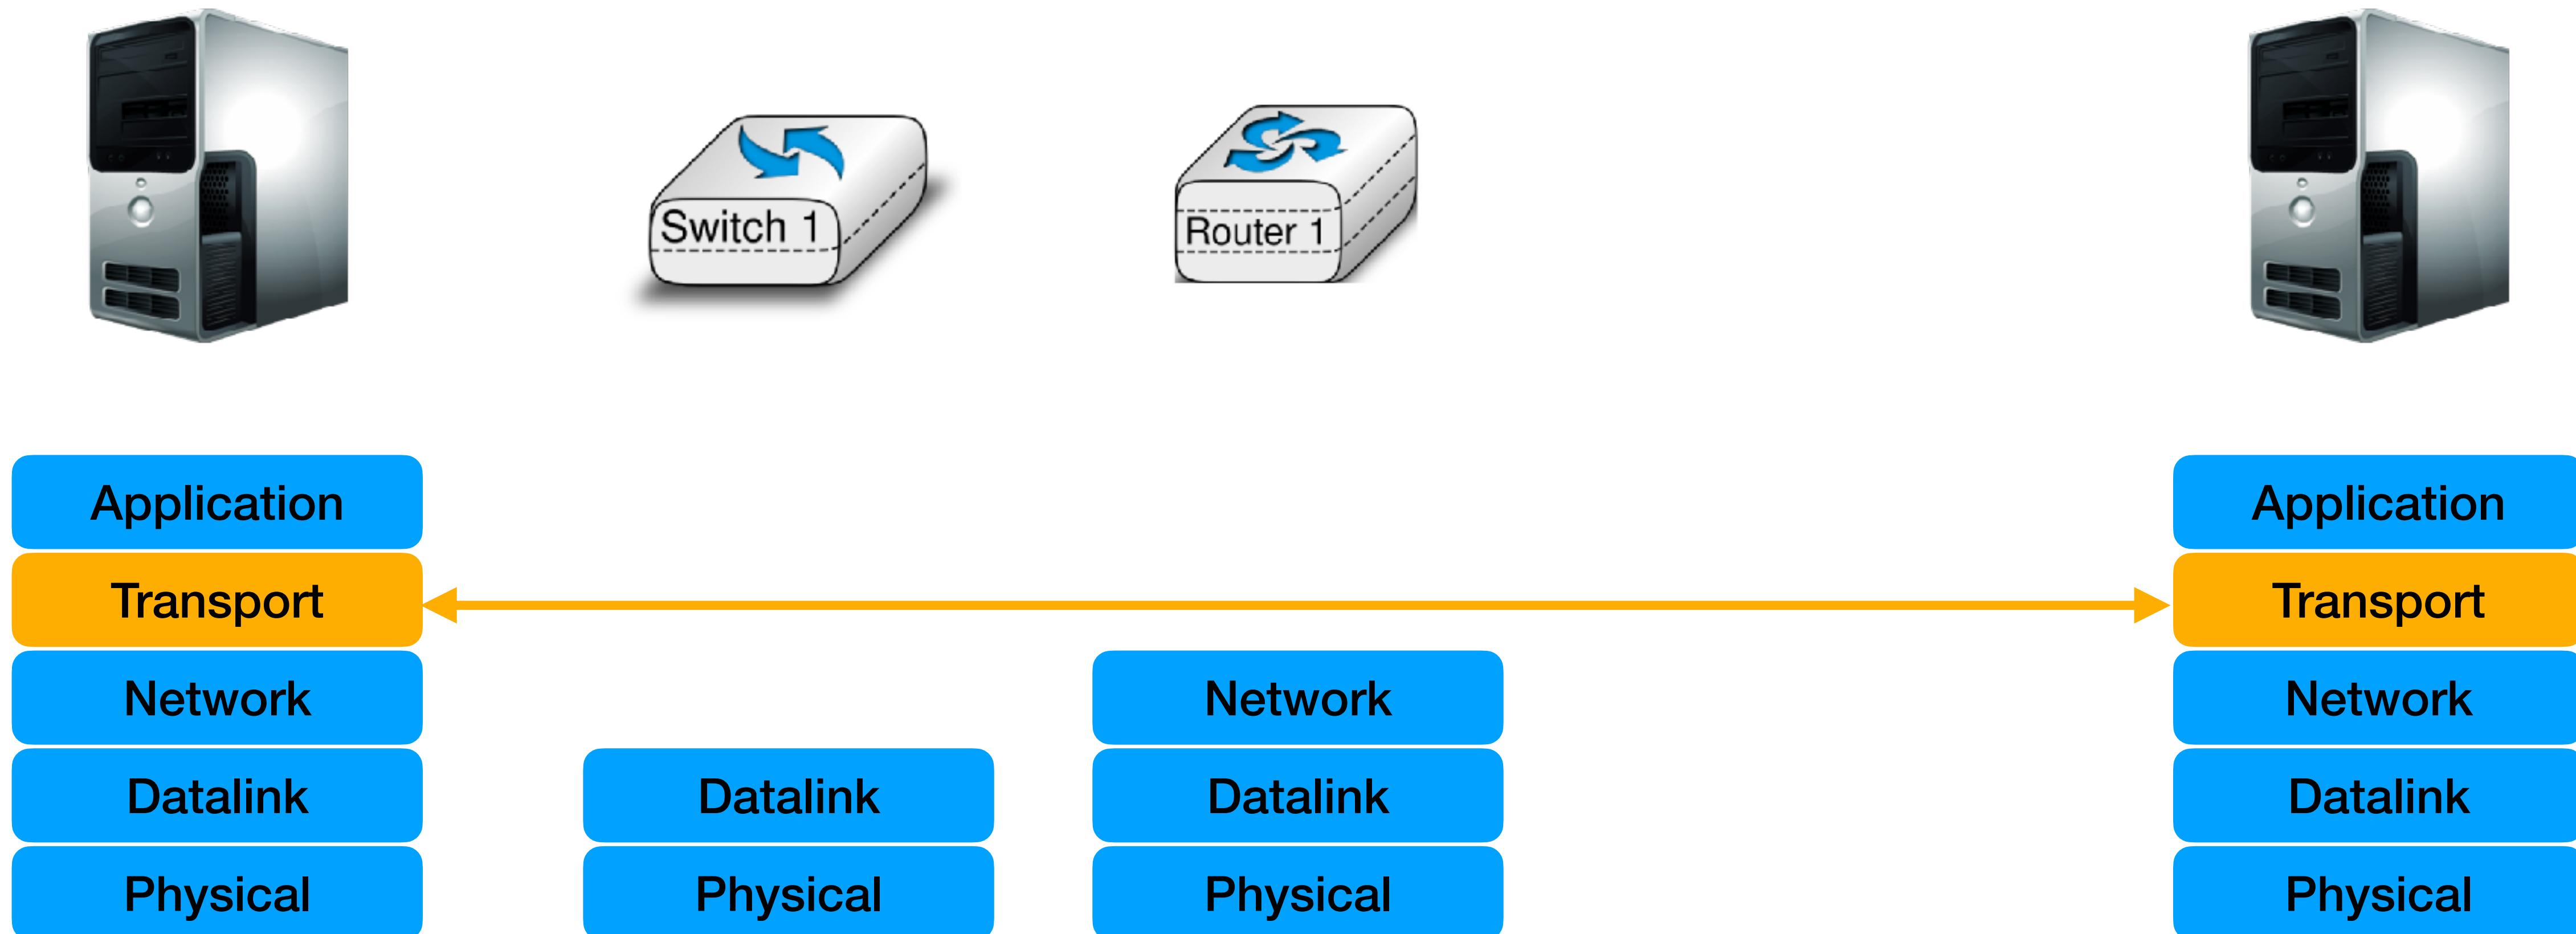
Quick Outline

- State-of-the-art: (Multipath) TCP
- Multipath in two simple applications
- (Multipath) QUIC
- Applications to Video Streaming

The Common (Transport) Protocol: TCP

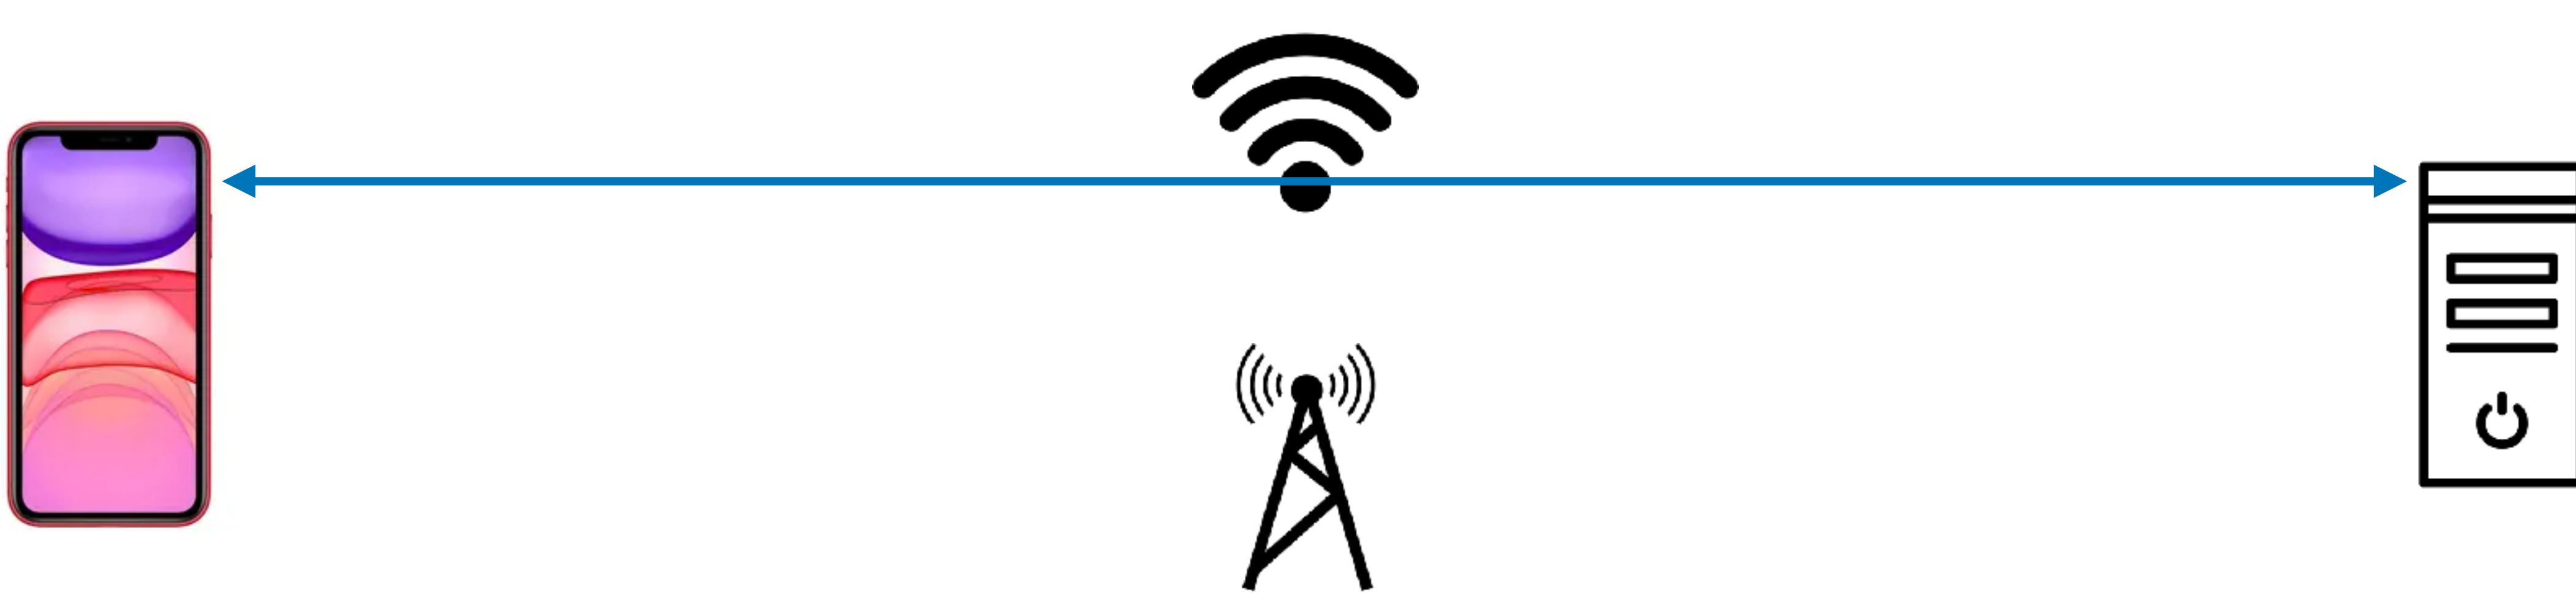
UDP being the other major one

End-to-End Principle



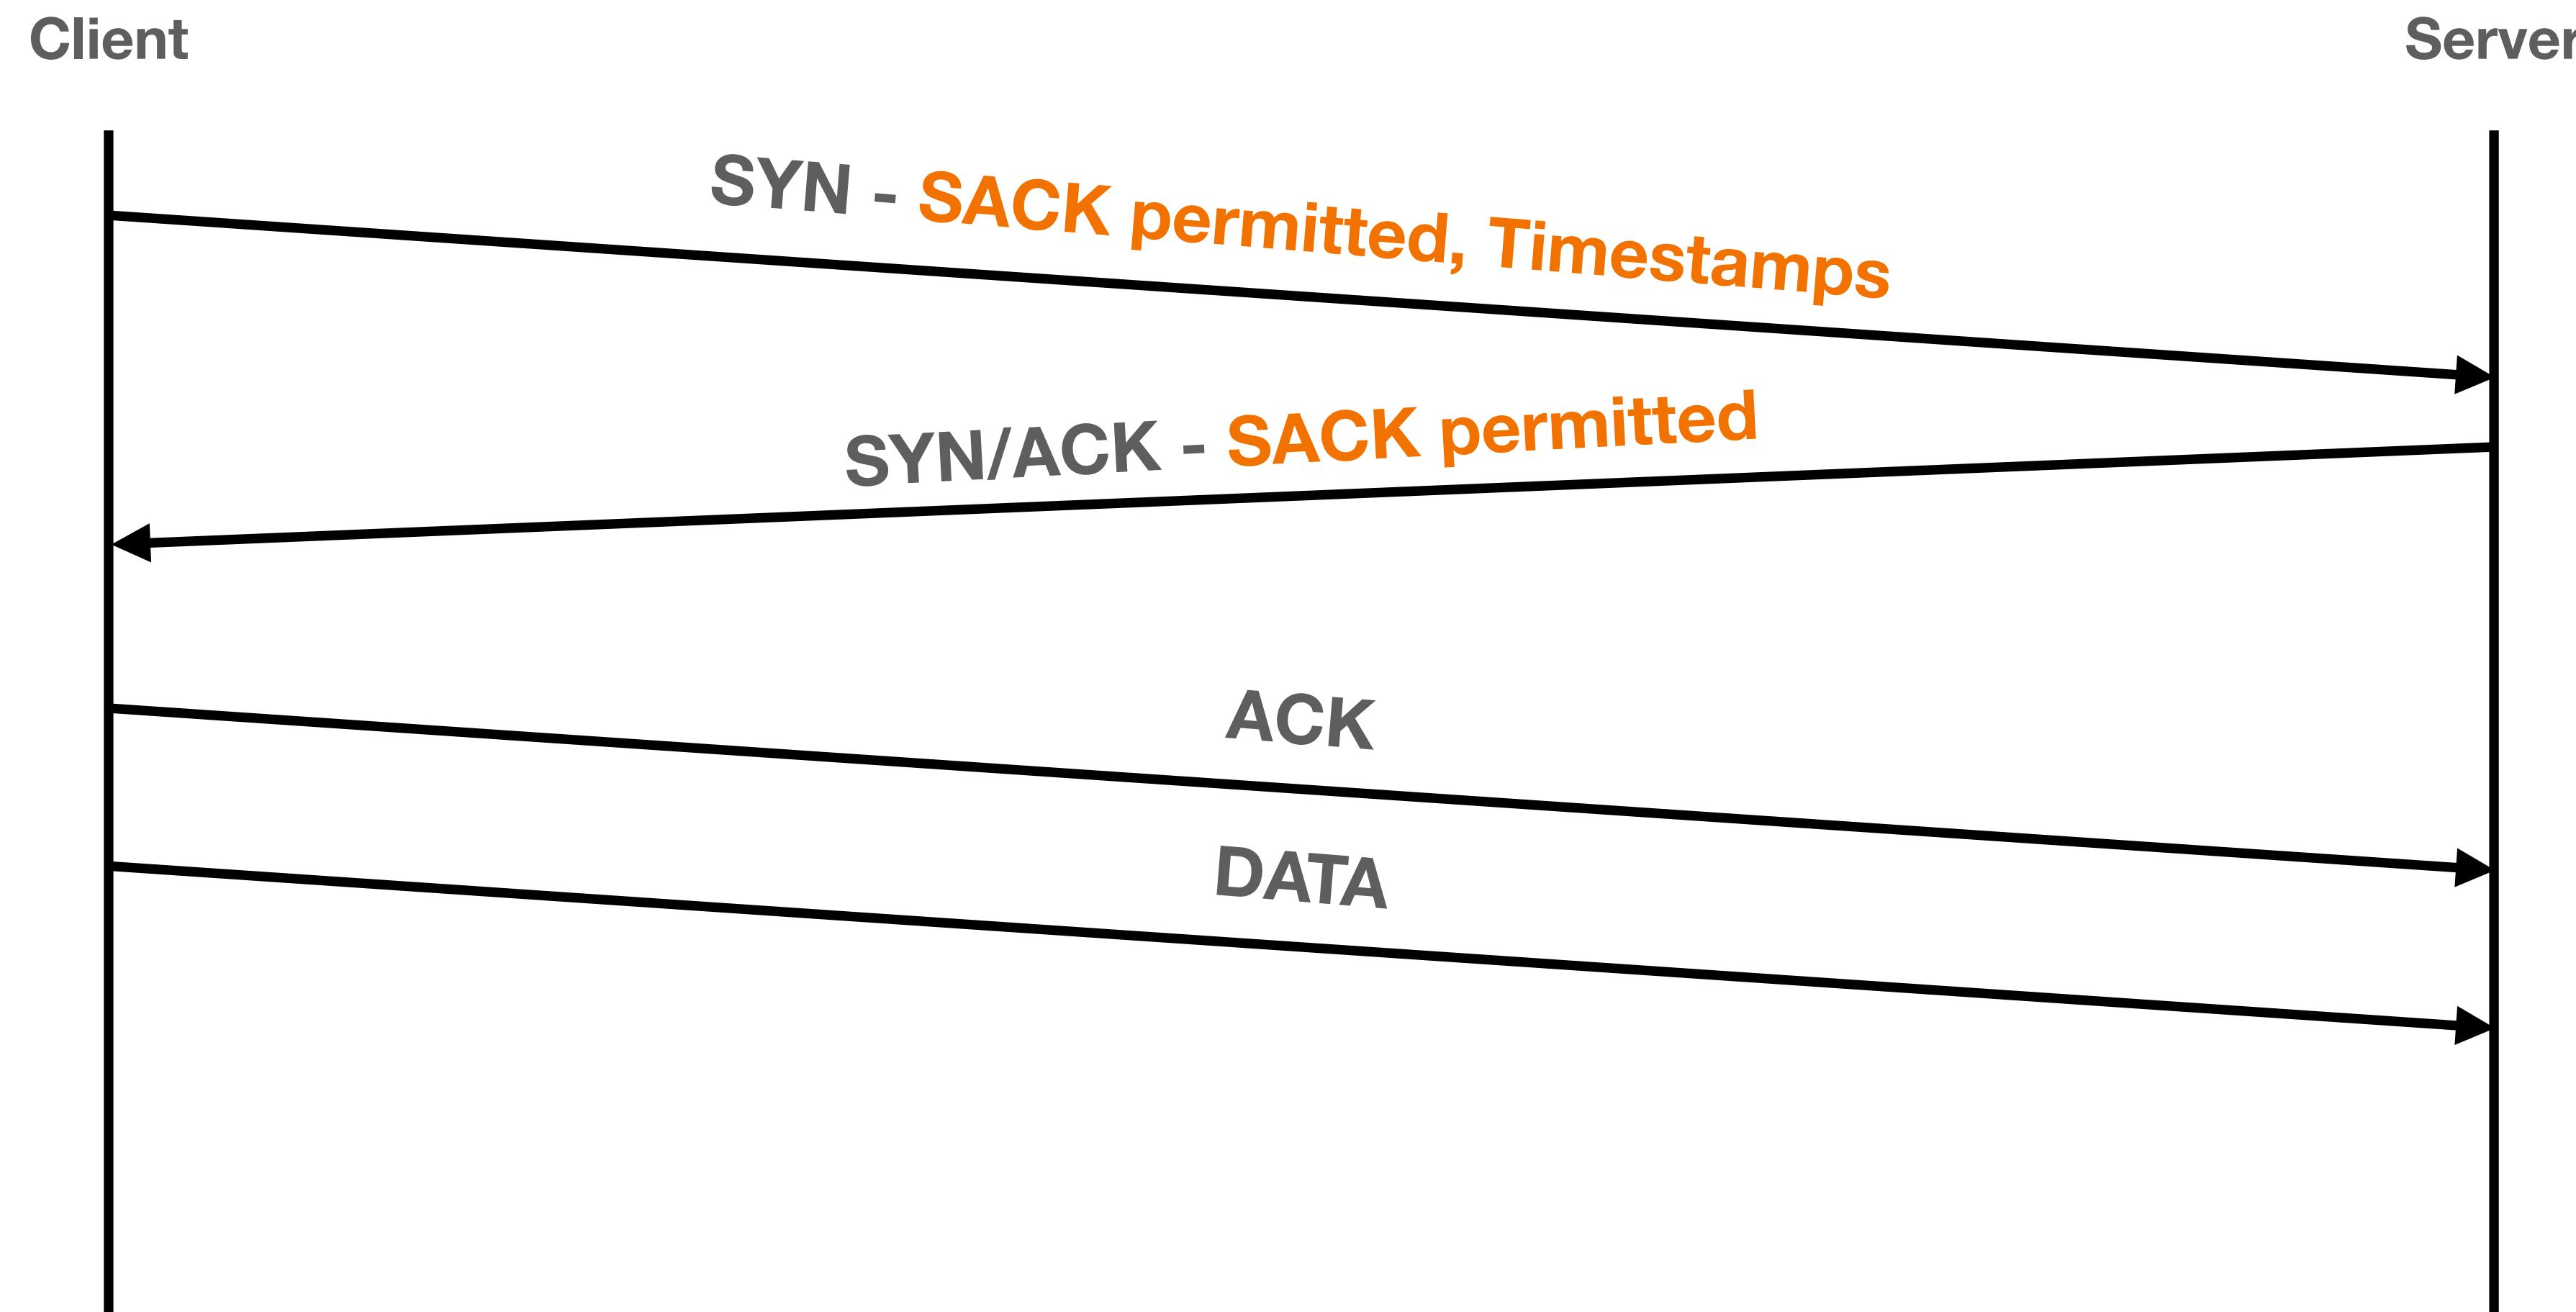
TCP and the Bytestream Service

Reliable, in-order delivery



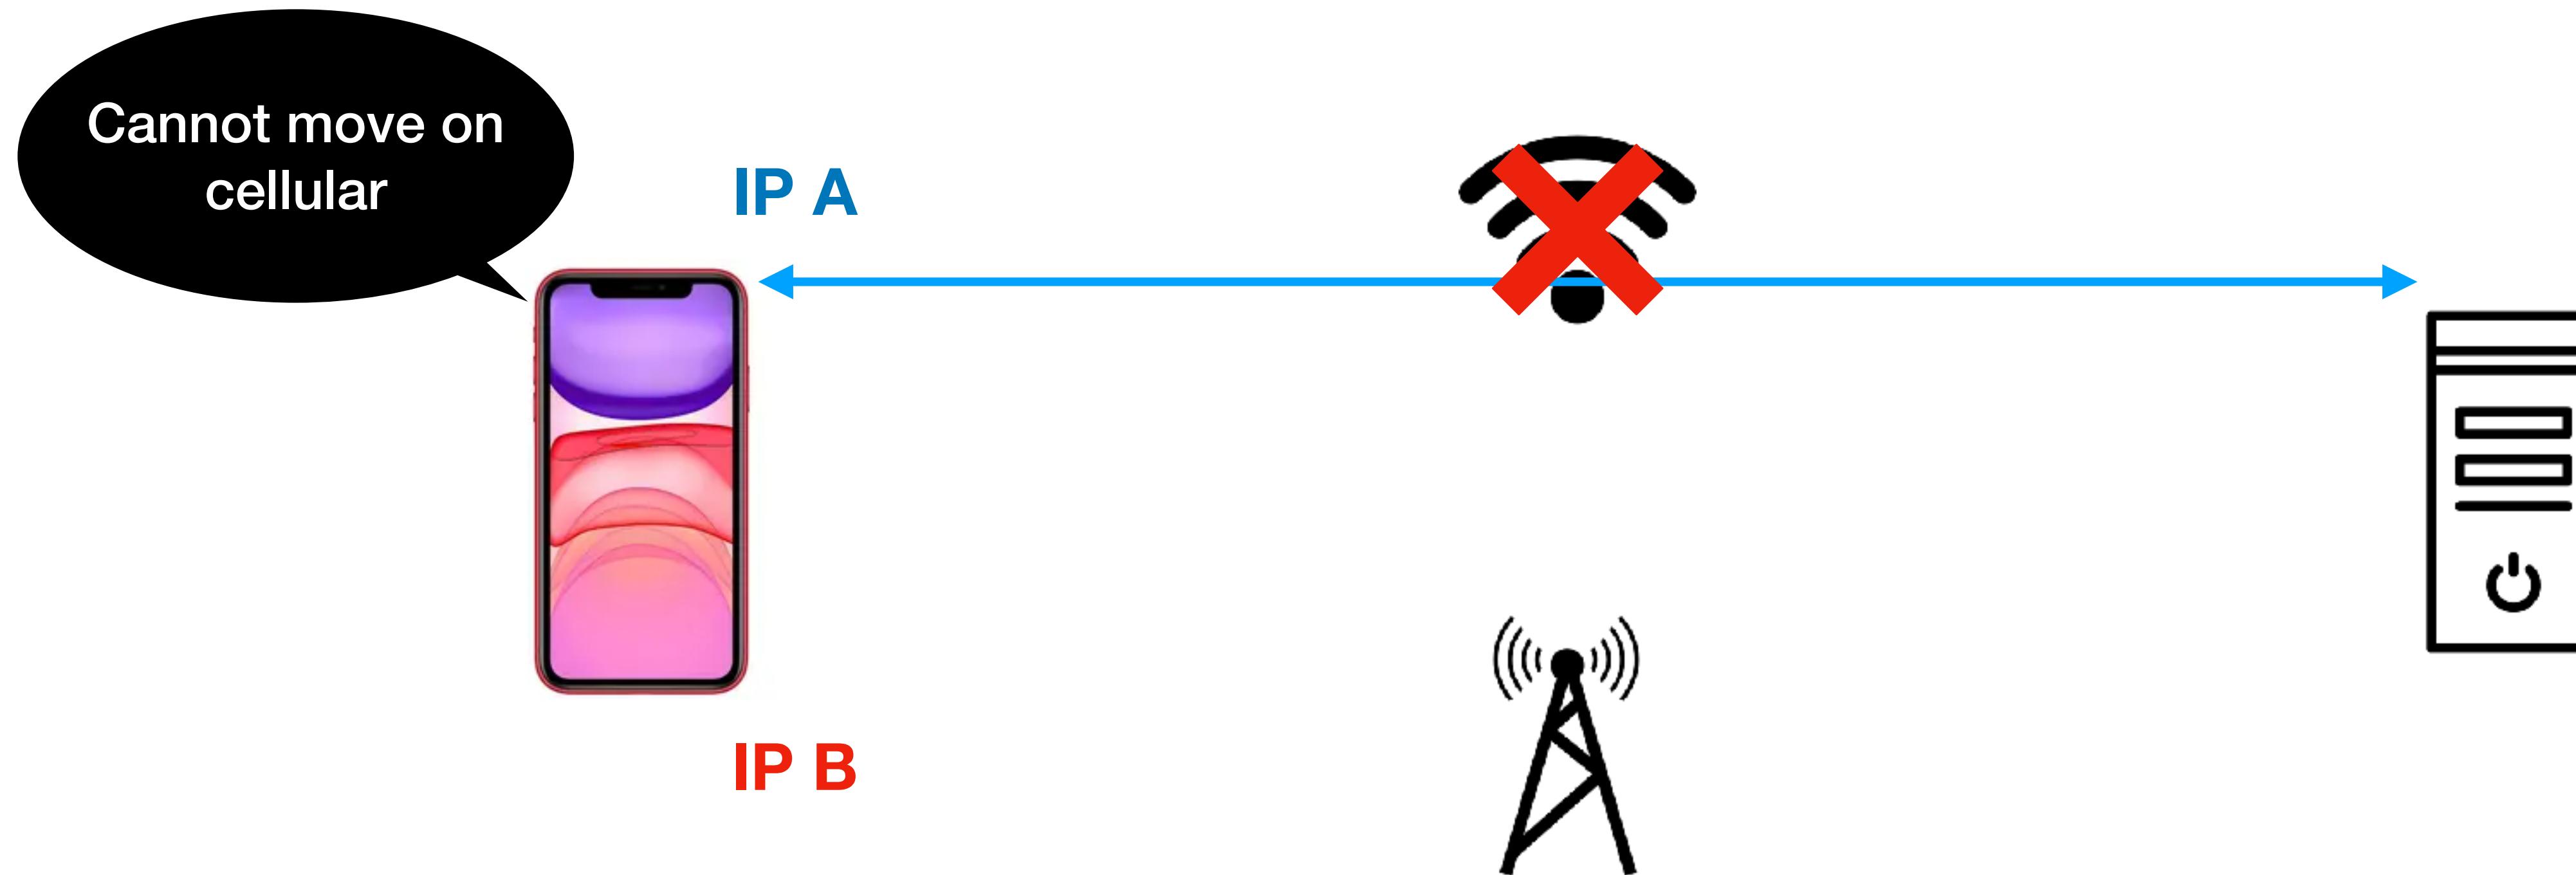
Establishing TCP Connections

a.k.a., the TCP Joke



The Mismatch with Mobile Devices

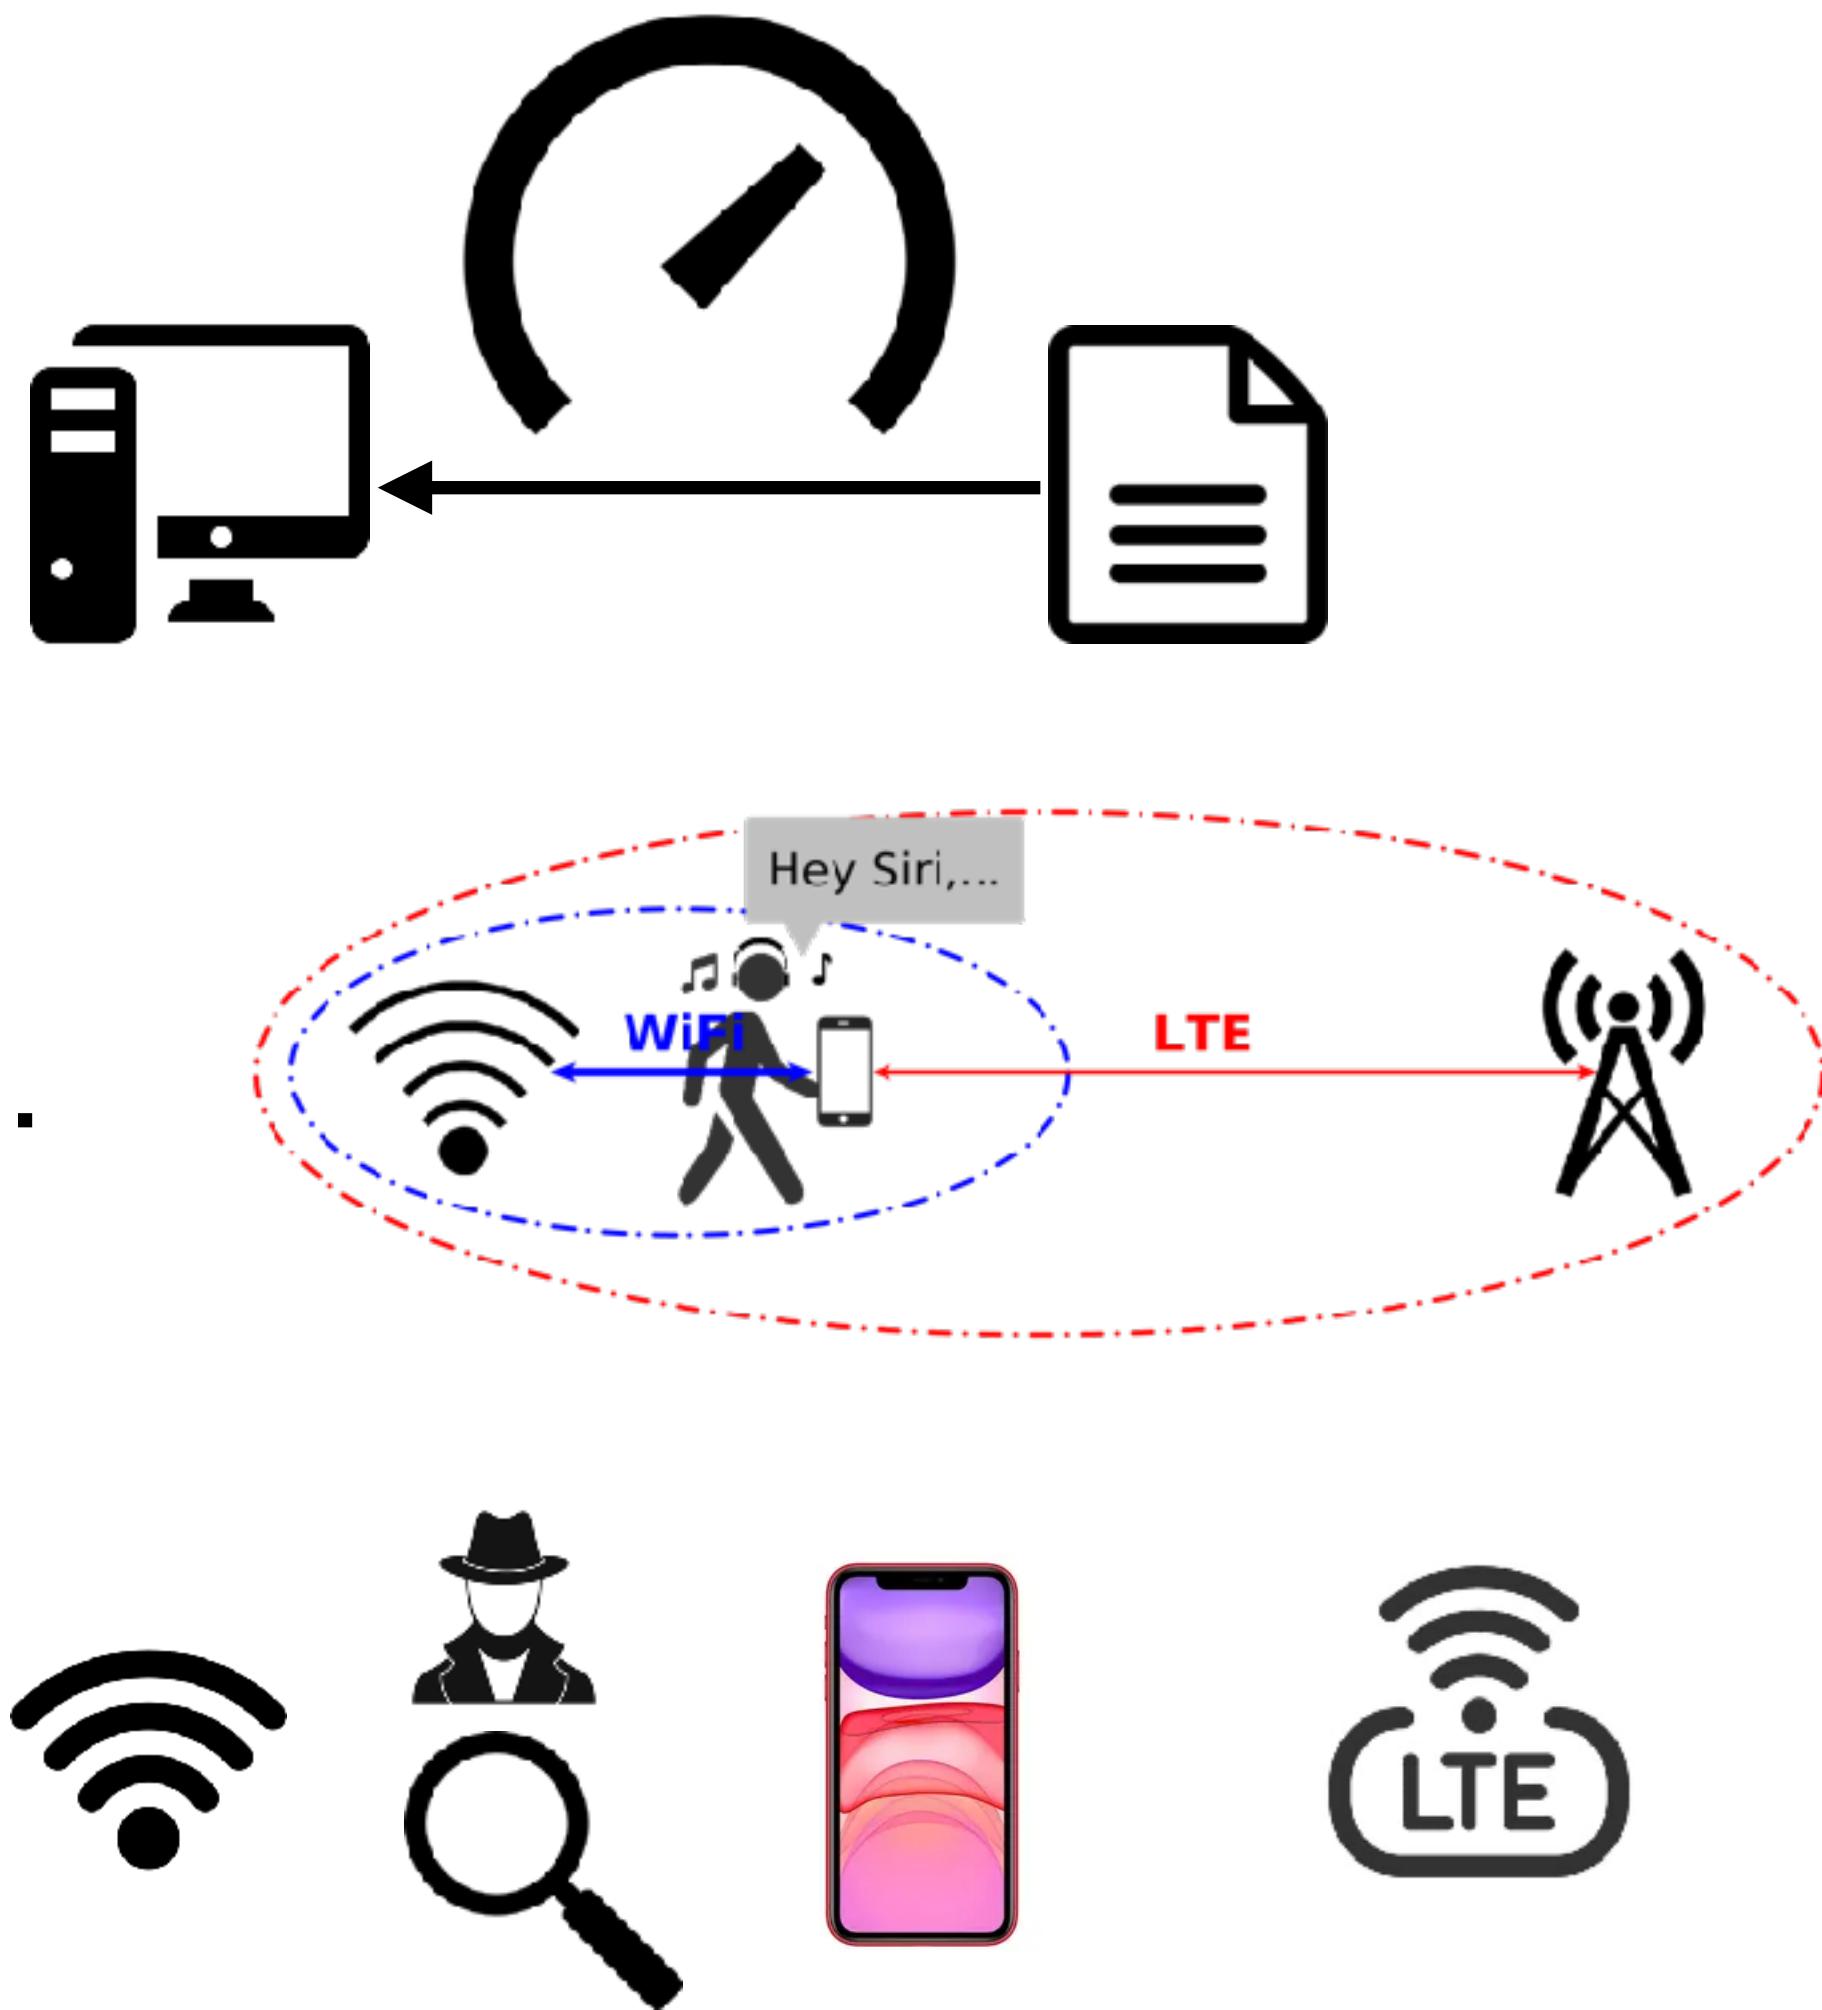
TCP Connections Are Bound to their 4-tuple (IP_{src} , port_{src} , IP_{dst} , port_{dst})



Losing the network path = Broken connectivity

Why Multipath?

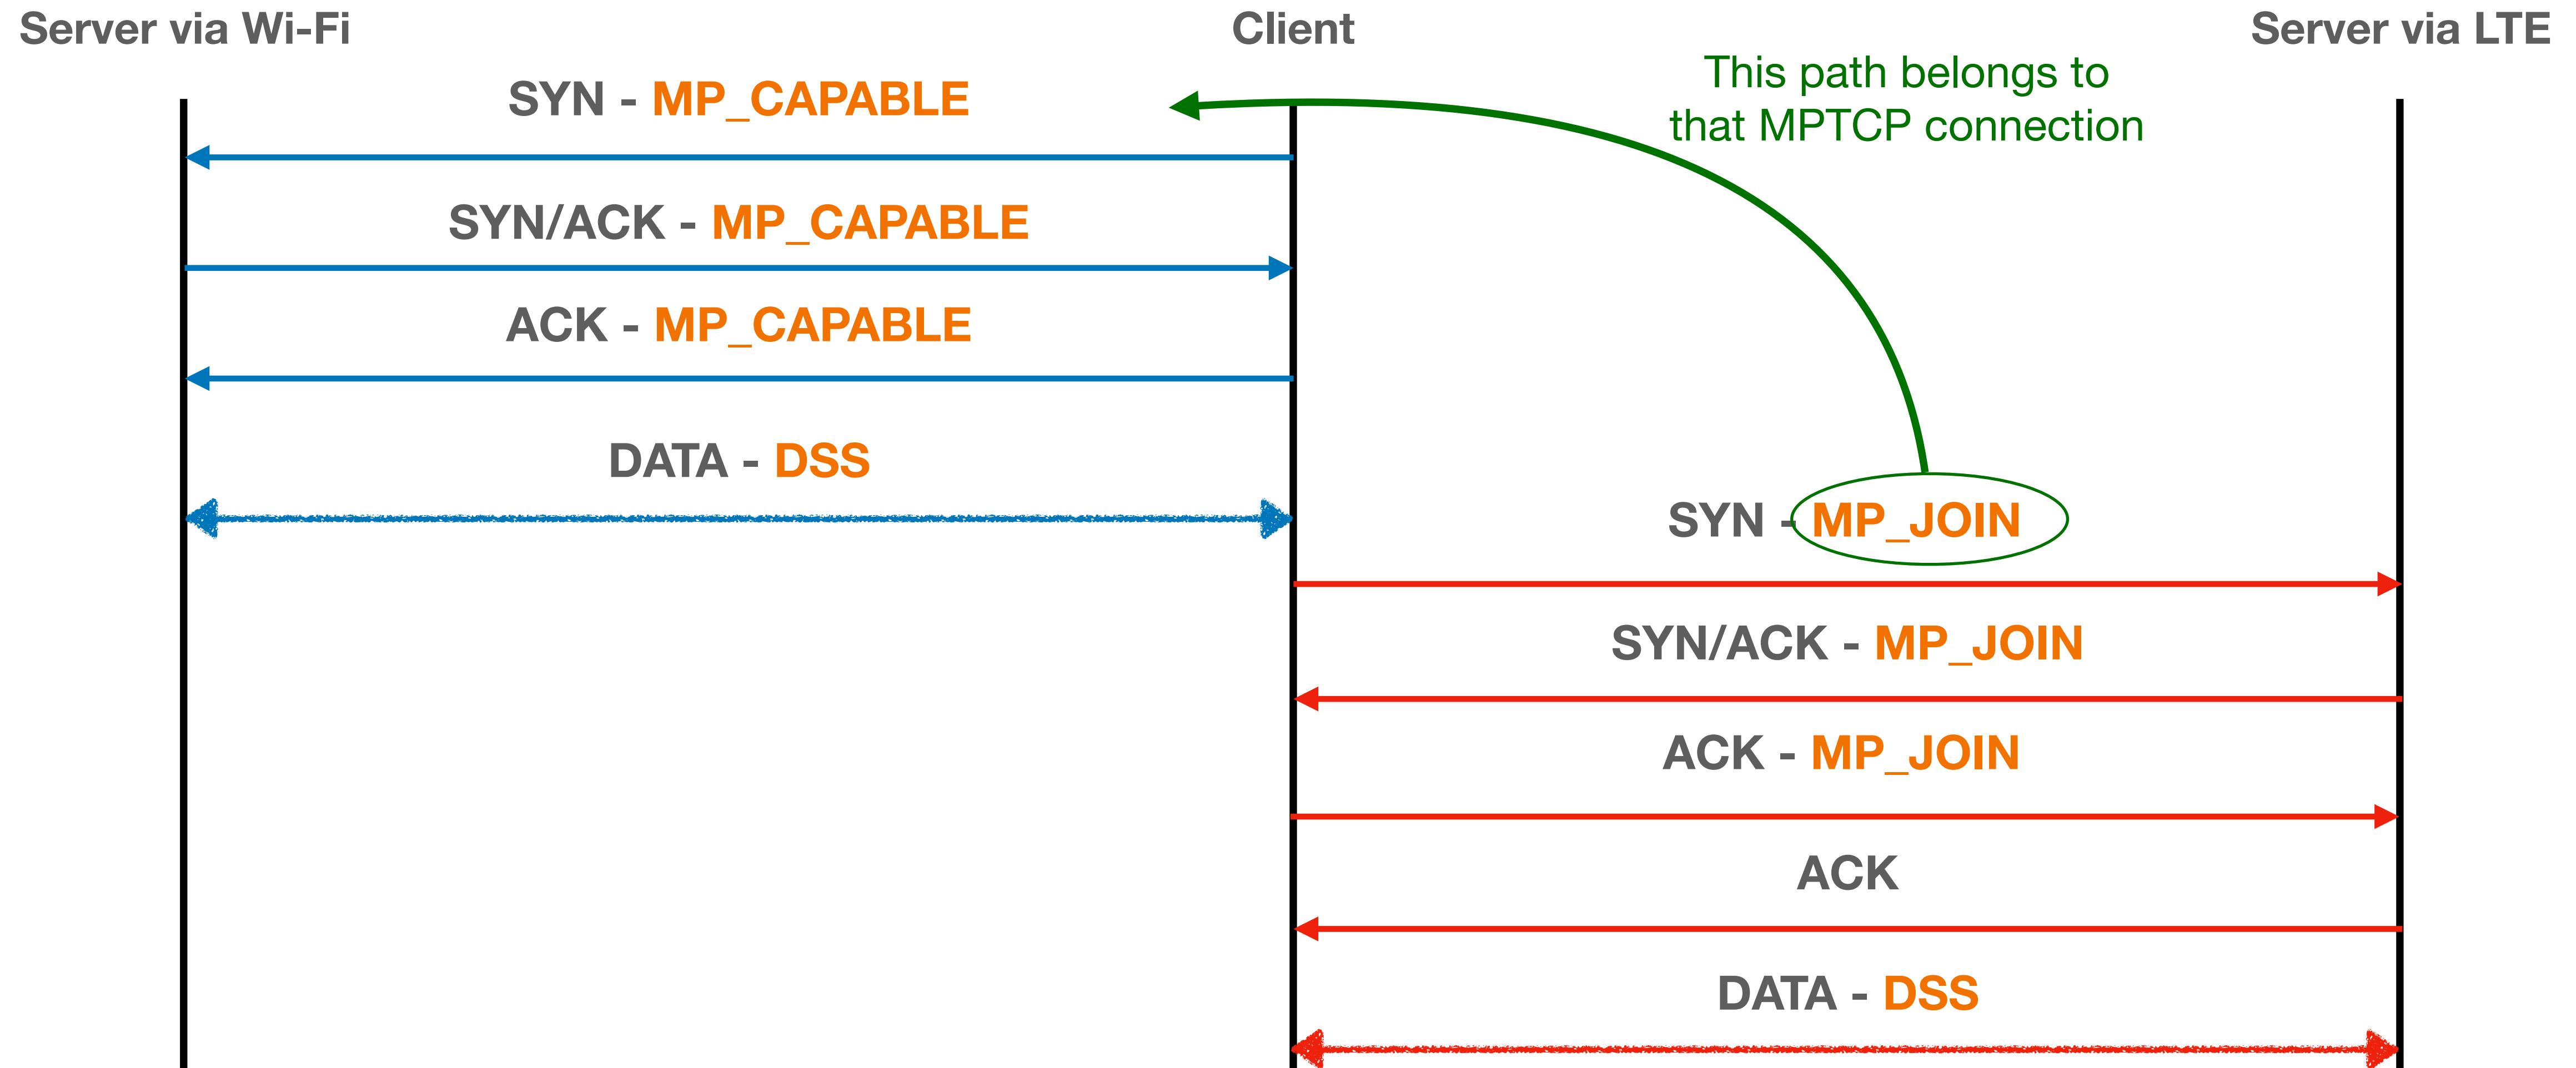
- Bandwidth Aggregation
 - E.g., download your files faster
- Keeping Low Latency
 - E.g., non-stable condition, network handover,...
- Network Management
 - E.g., control where data is sent (privacy,...)



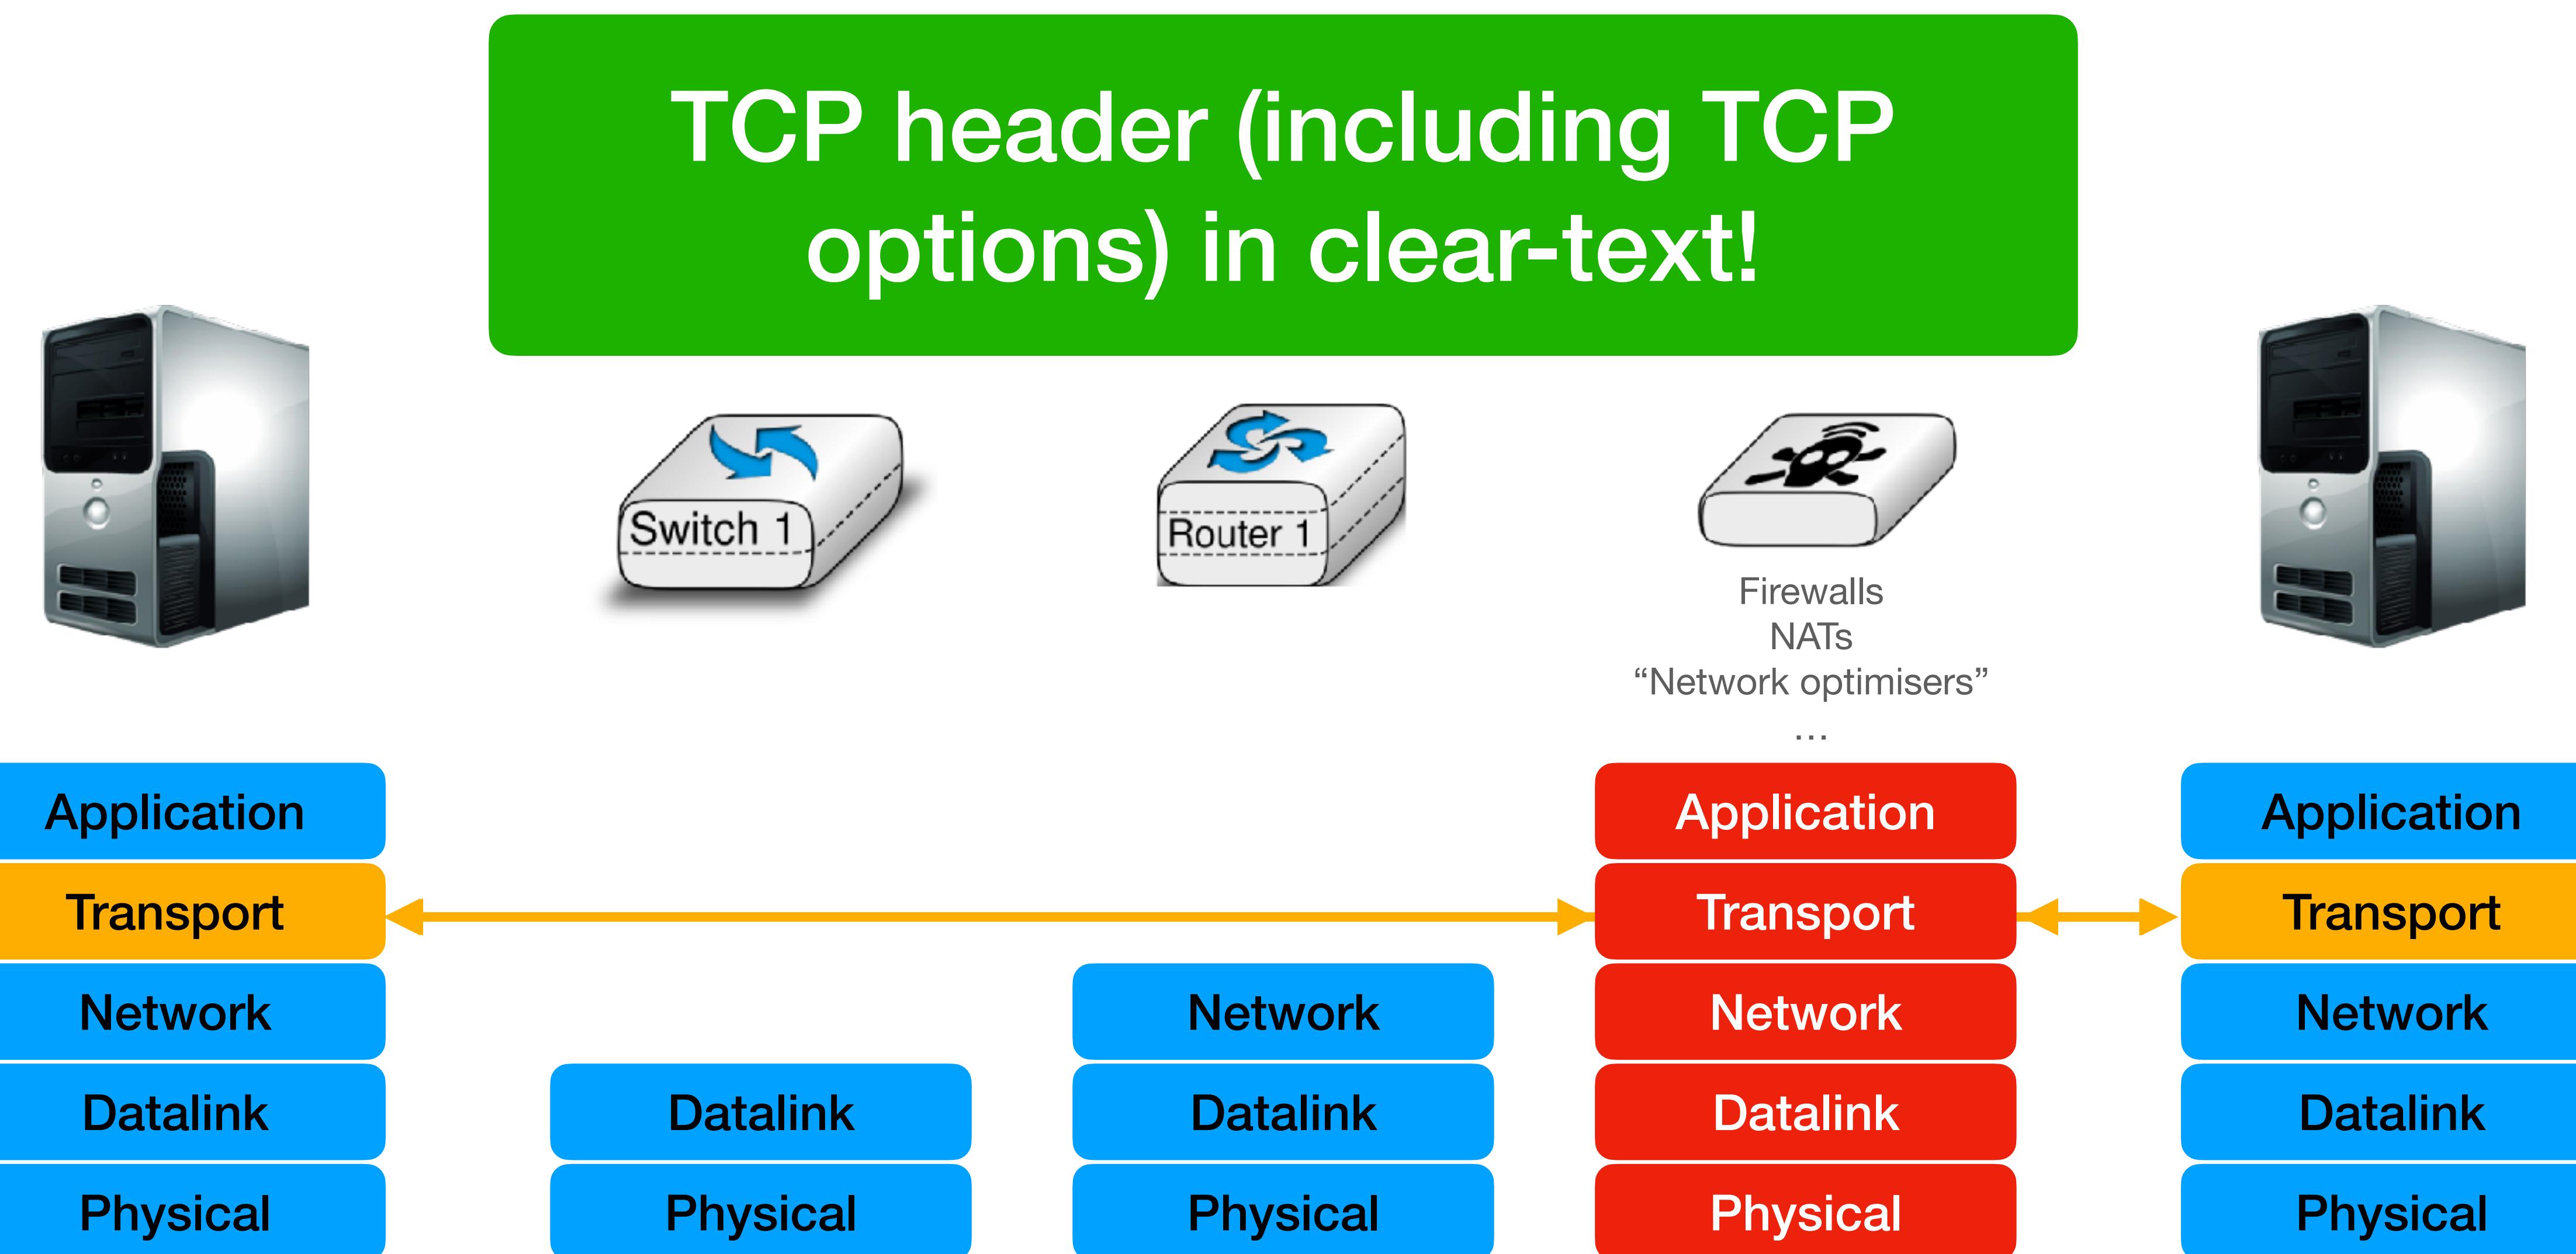
Bringing Multipath into TCP

Multipath TCP

As a TCP extension, using TCP options

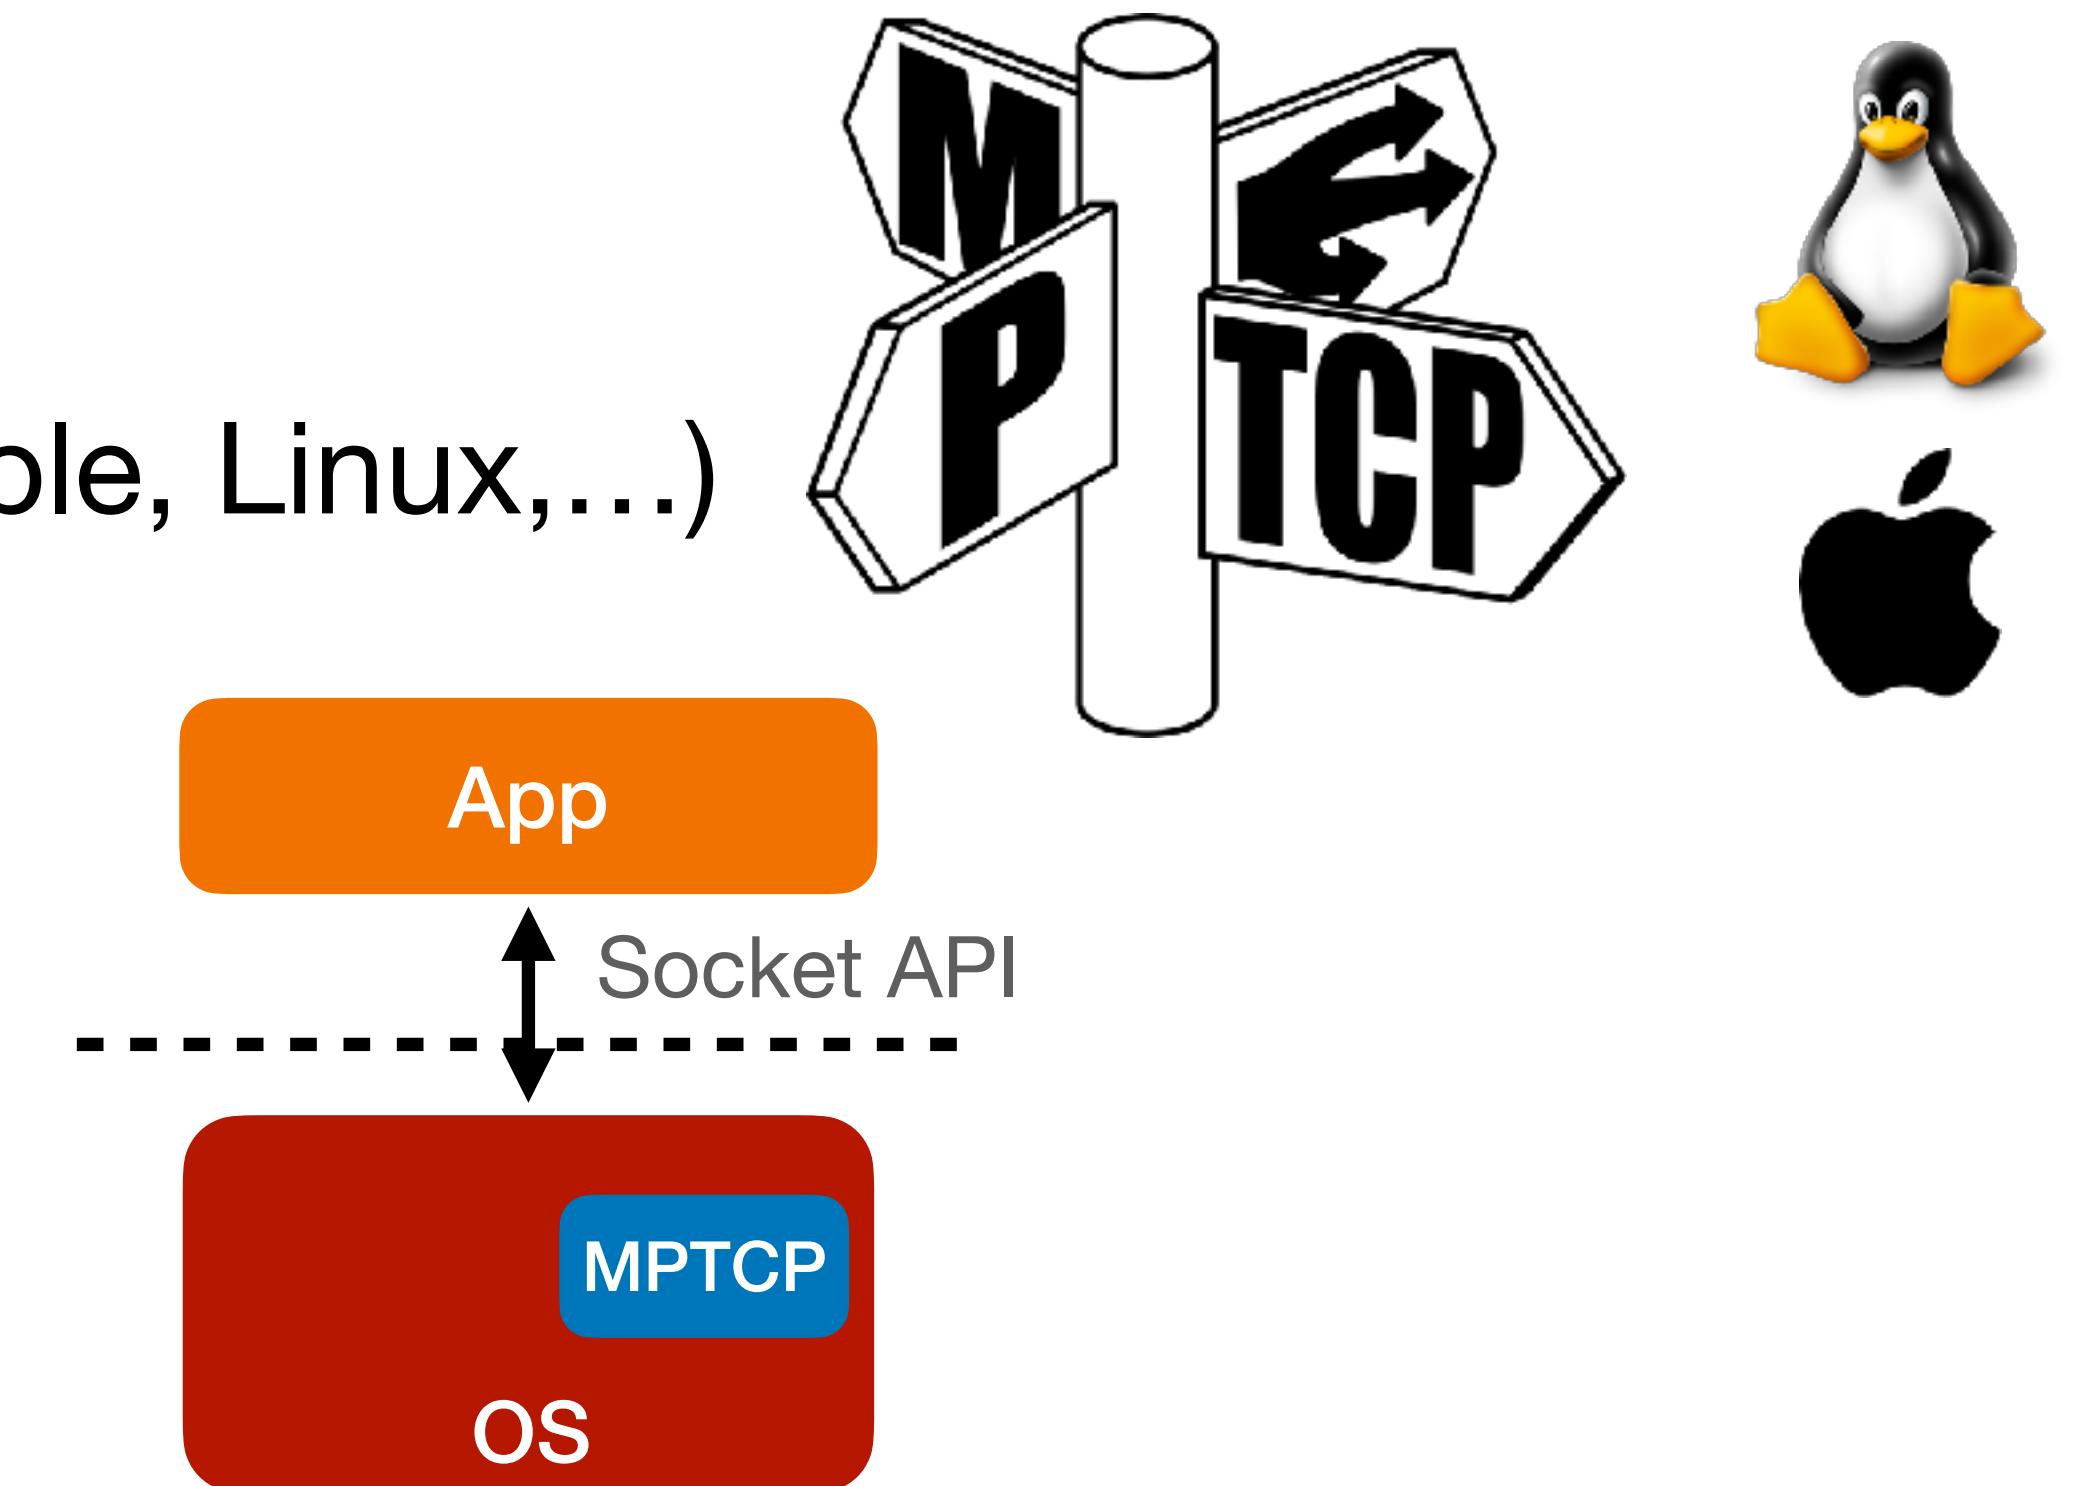


End-to-End Principle: The Reality



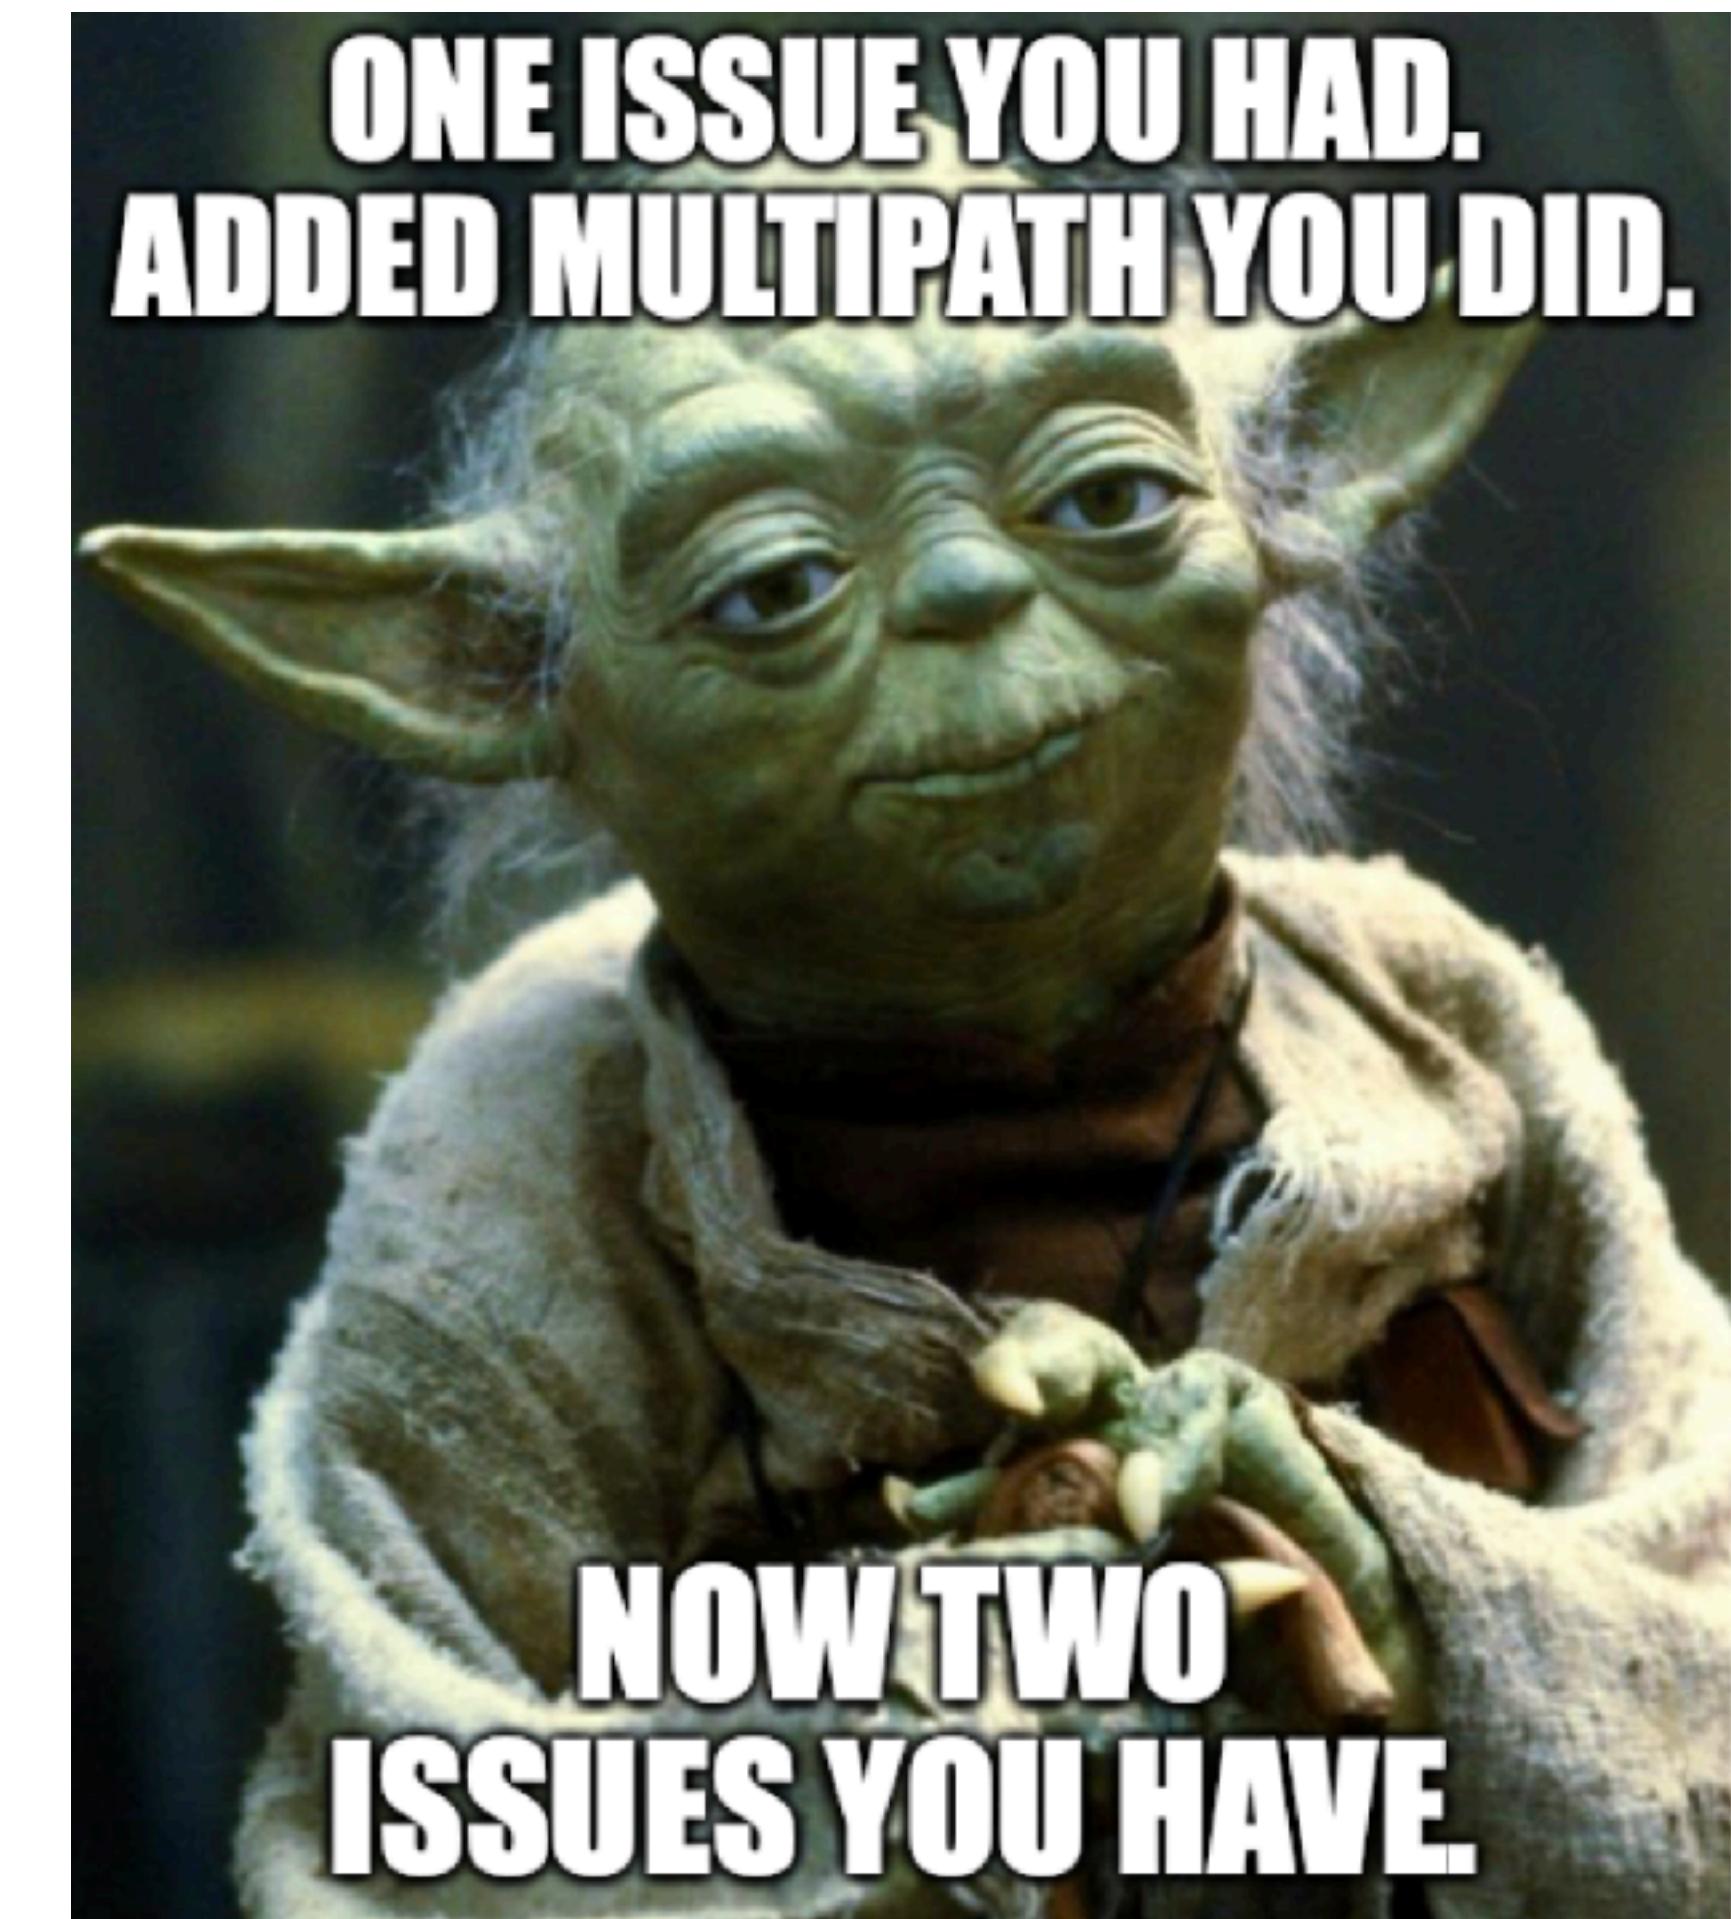
A Deployable Multipath TCP

- Design objective: Multipath TCP must work where TCP works
 - Each MPTCP path behaves like a regular TCP connection
 - Network anomaly? Fallback to plain TCP
- Several **kernel-space** implementations (Apple, Linux,...)
- App can do MPTCP without modification!



Multipath Algorithms

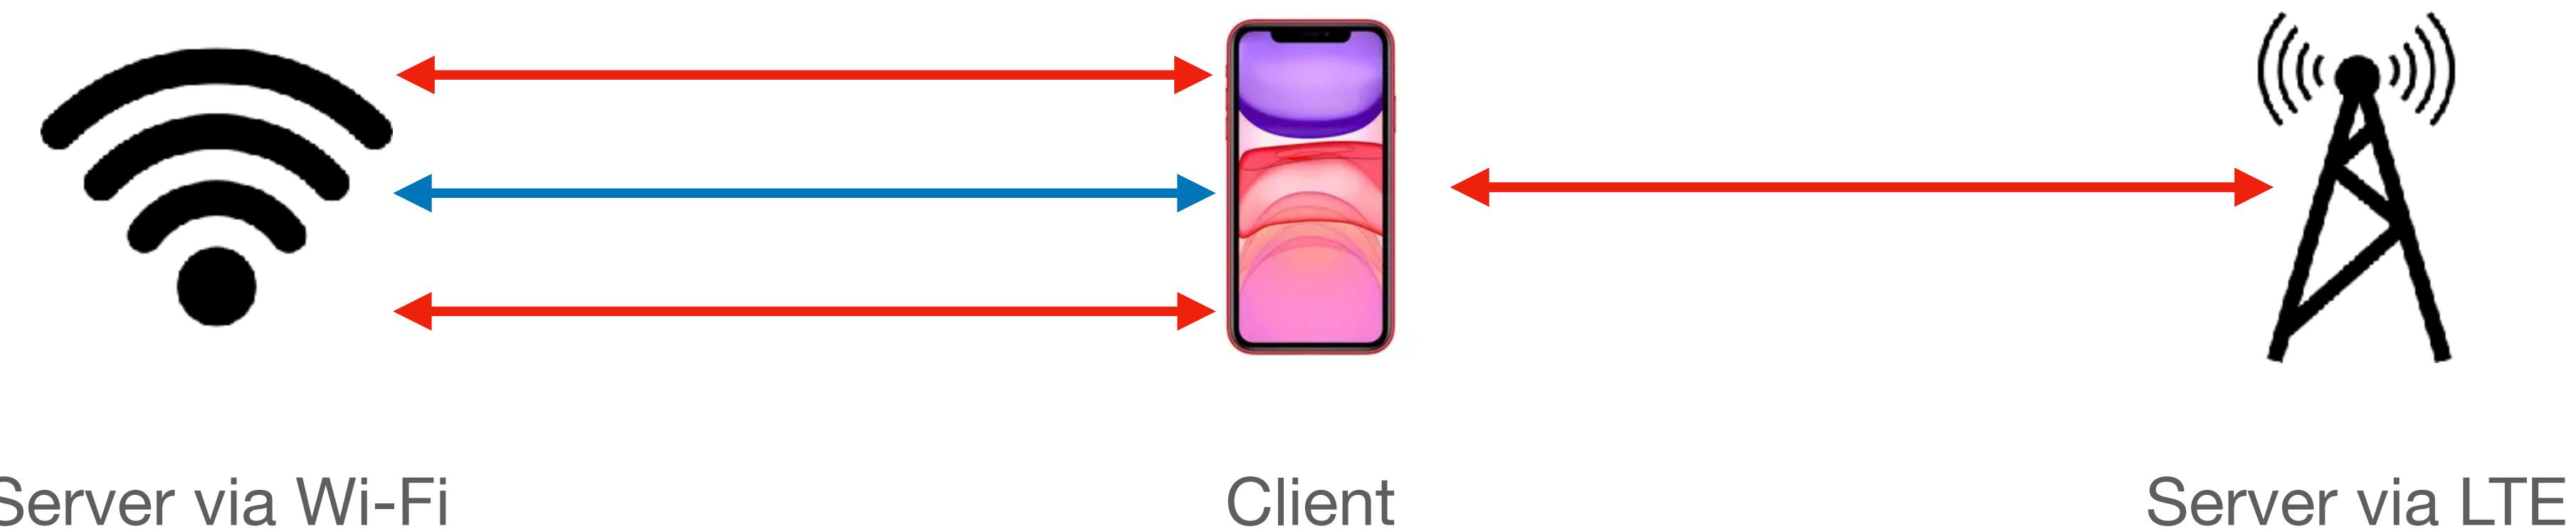
- Path management
 - Creating paths
 - Packet scheduling
 - Using paths
- (Congestion control algorithm)
 - Per-path state



Path Management

Creating Paths

Fullmesh: create one path per network interface as soon as possible



Client

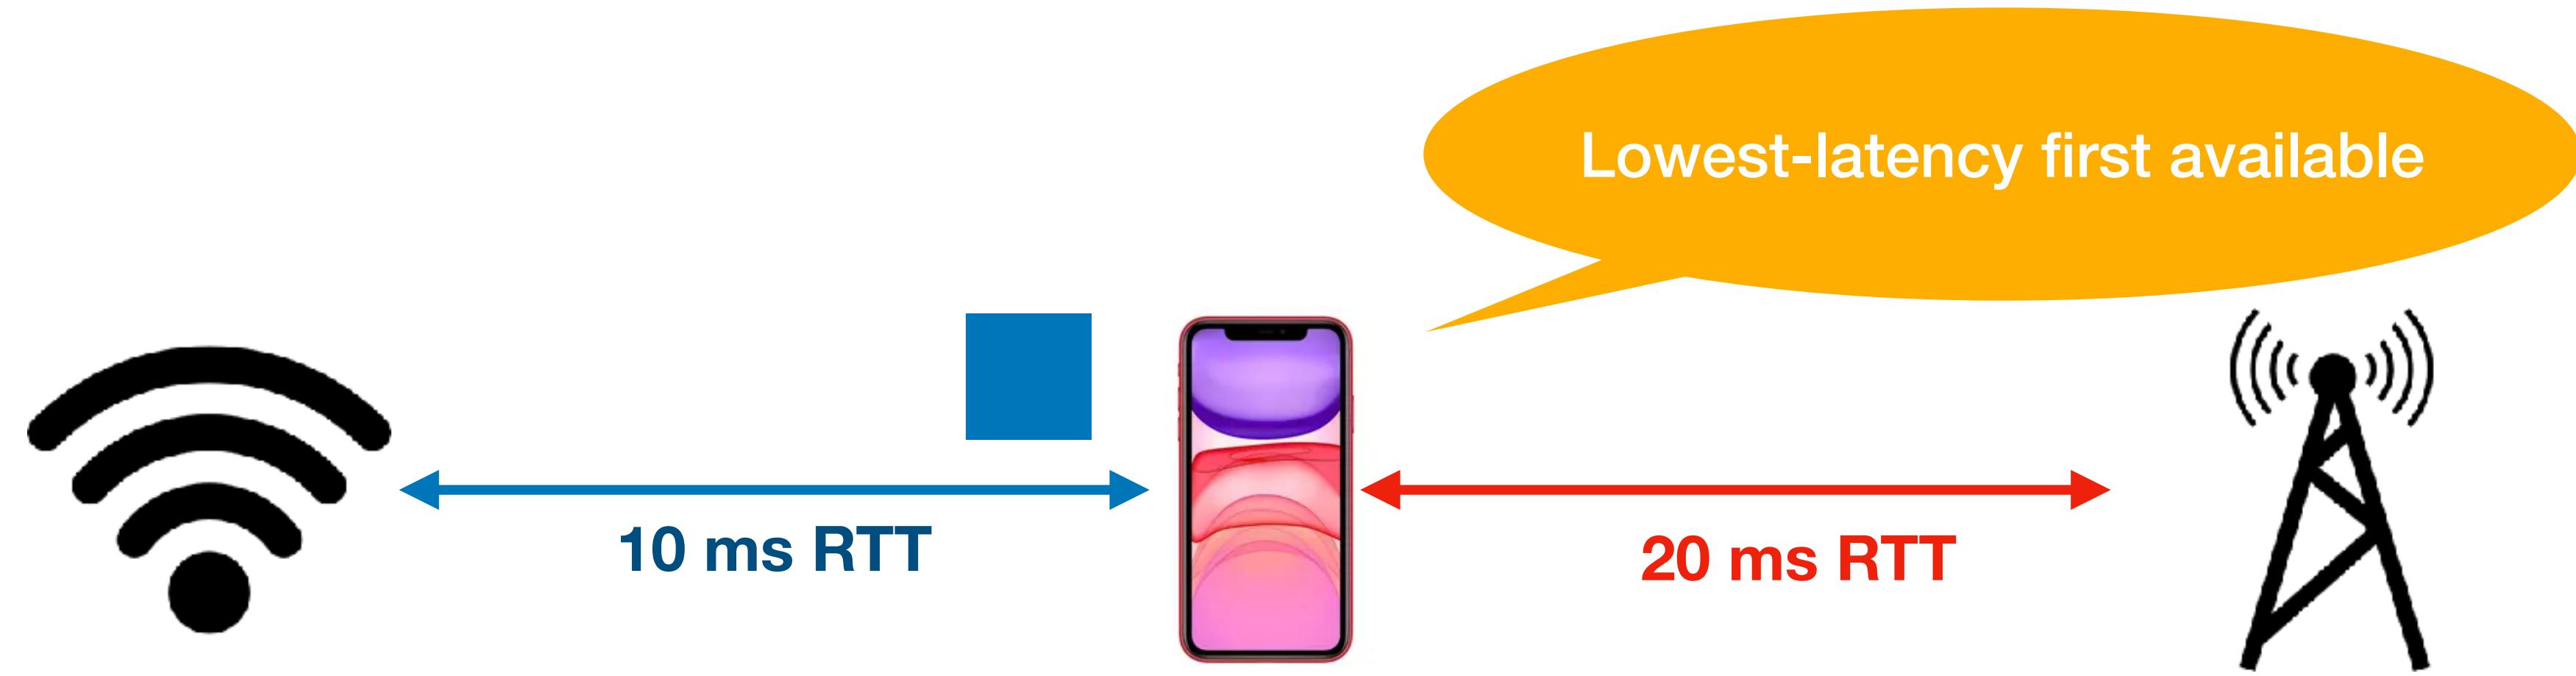
Server via Wi-Fi

Server via LTE

Client decision

Packet Scheduler Using Paths

In-kernel
implementation

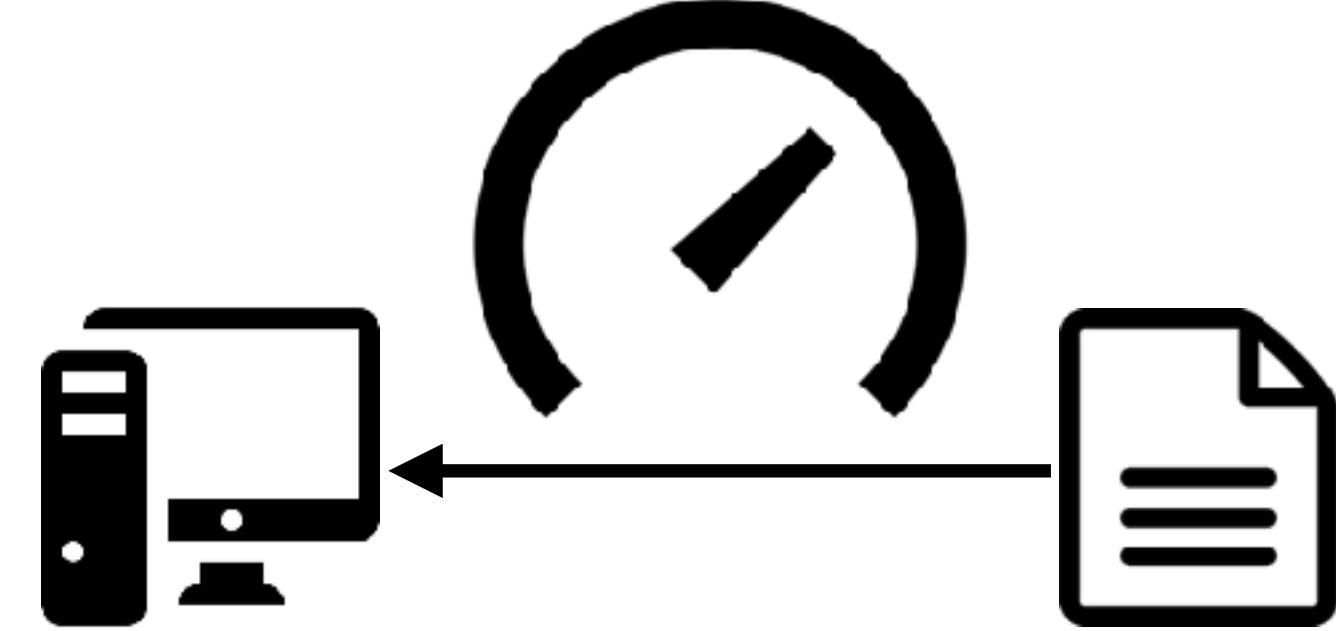


Sender decision

Upload → mostly client decisions
Download → mostly server decisions

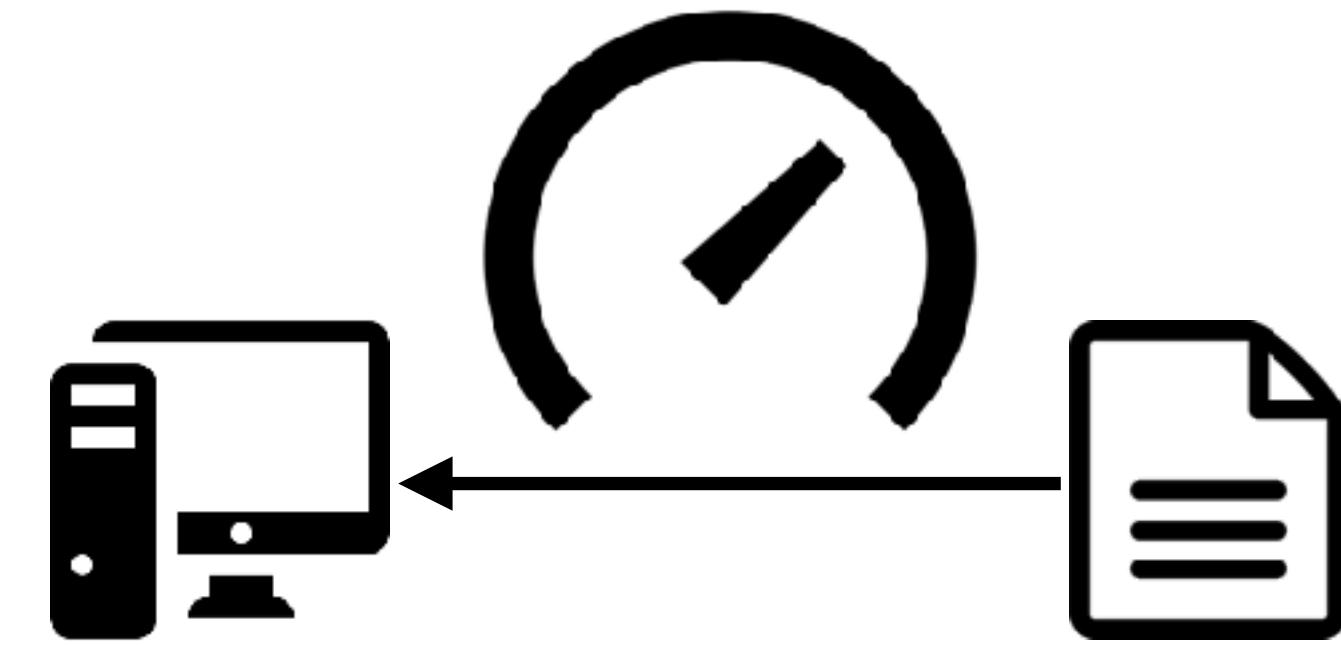
Our Two Basic Applications

- File downloader
 - Bandwidth Aggregation driven
 - E.g., Korean Telecom GiGA LTE
- Interactive Application (under mobile scenario)
 - Request-response based, latency-sensitive
 - E.g., voice-activated applications like Apple Siri



A Multipath-Enabled File Downloader

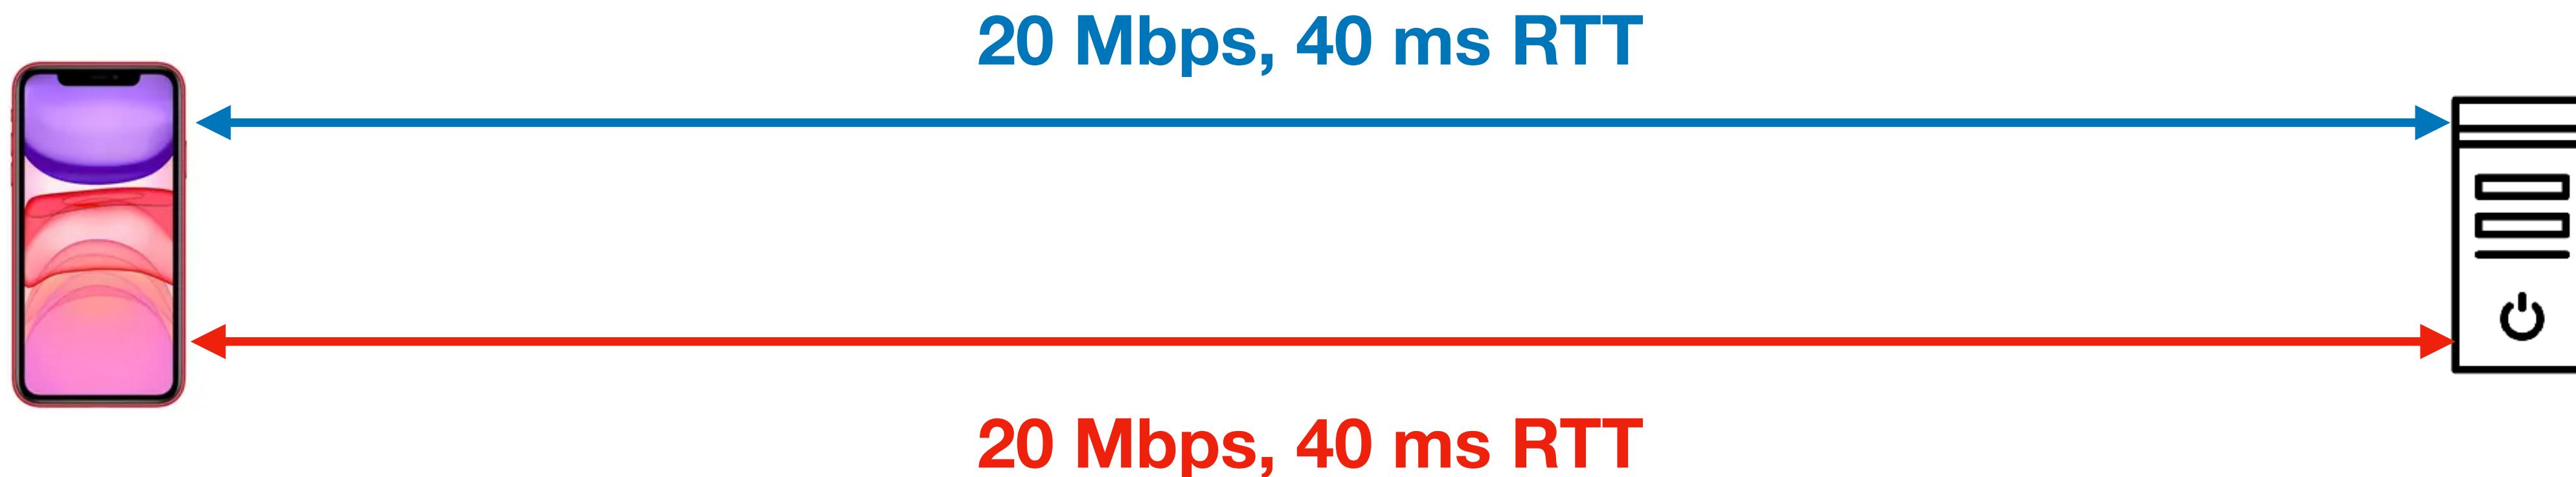
- User expects increased throughput
 - Regardless of file size
 - Regardless of network conditions
- Assume plain HTTP GET request for a file



A Multipath-Enabled File Downloader

Homogeneous case

Fullmesh path manager
Lowest-latency first scheduler



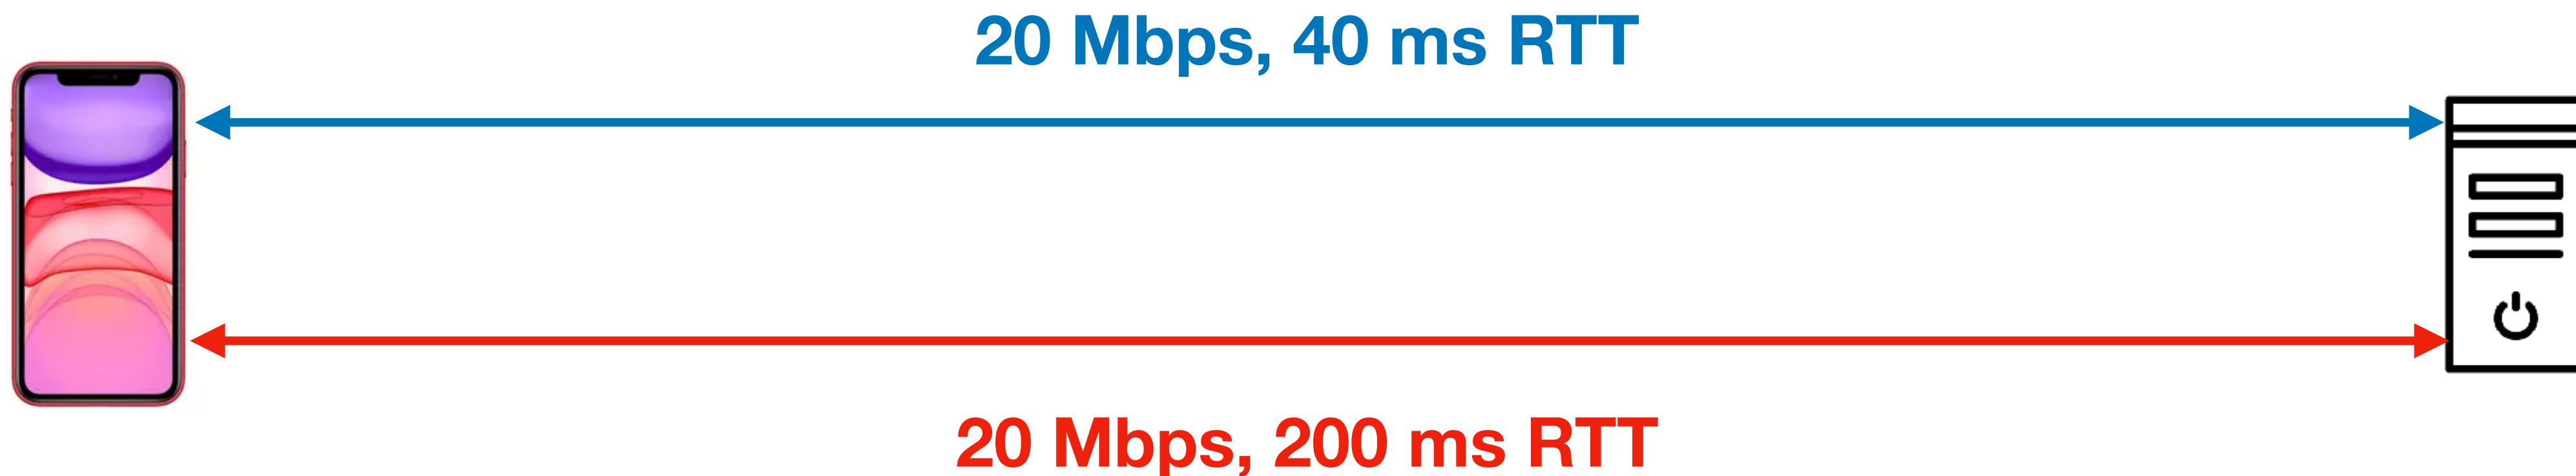
File size	20 MB	512 KB
Download time TCP (only blue path)	8.9 s	390 ms
Download time MPTCP (both paths)	4.65 s	250 ms ~1.9x faster

~1.56x faster

A Multipath-Enabled File Downloader

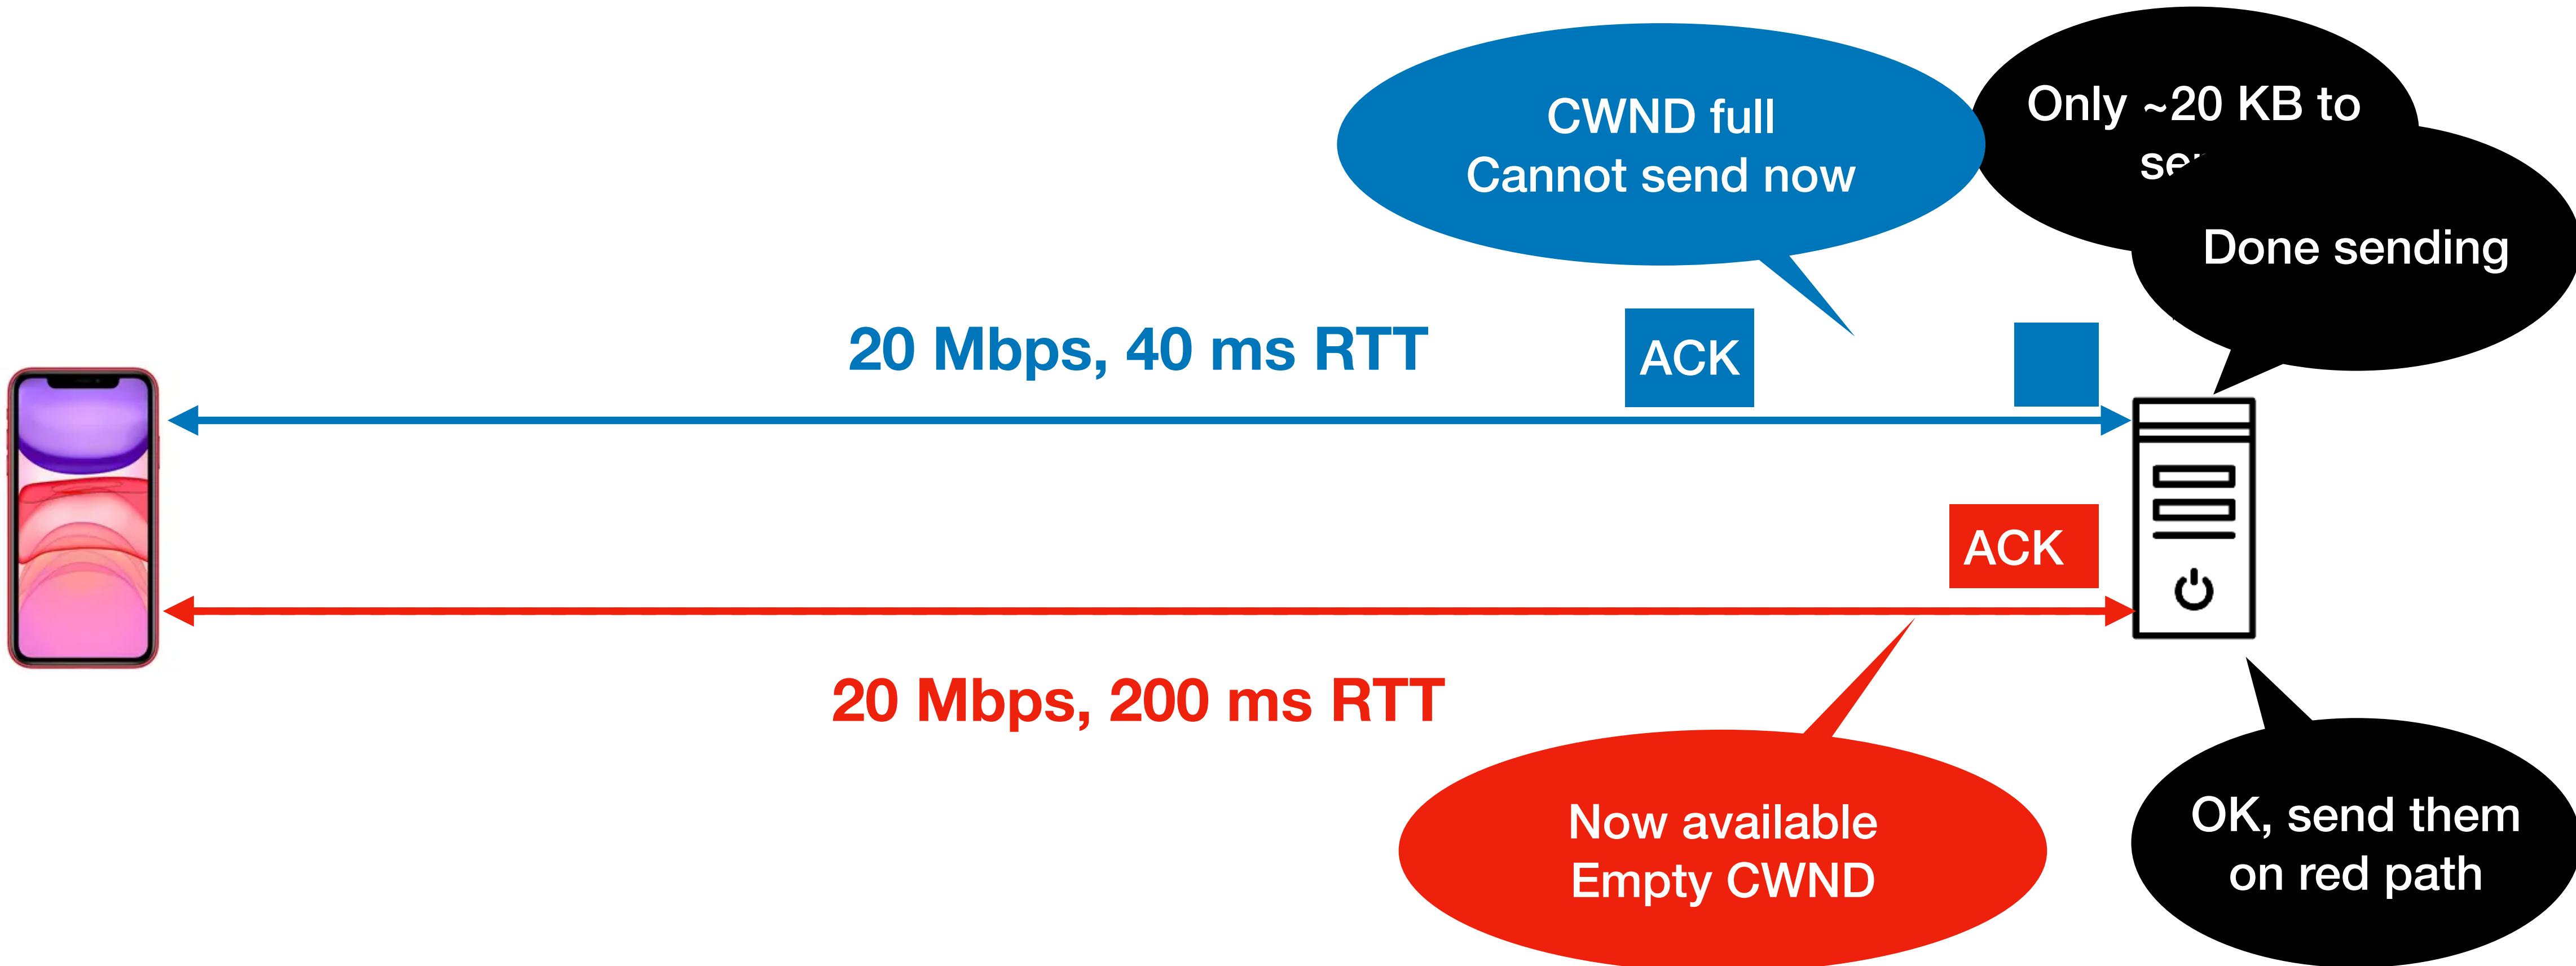
Heterogeneous case

Fullmesh path manager
Lowest-latency first scheduler



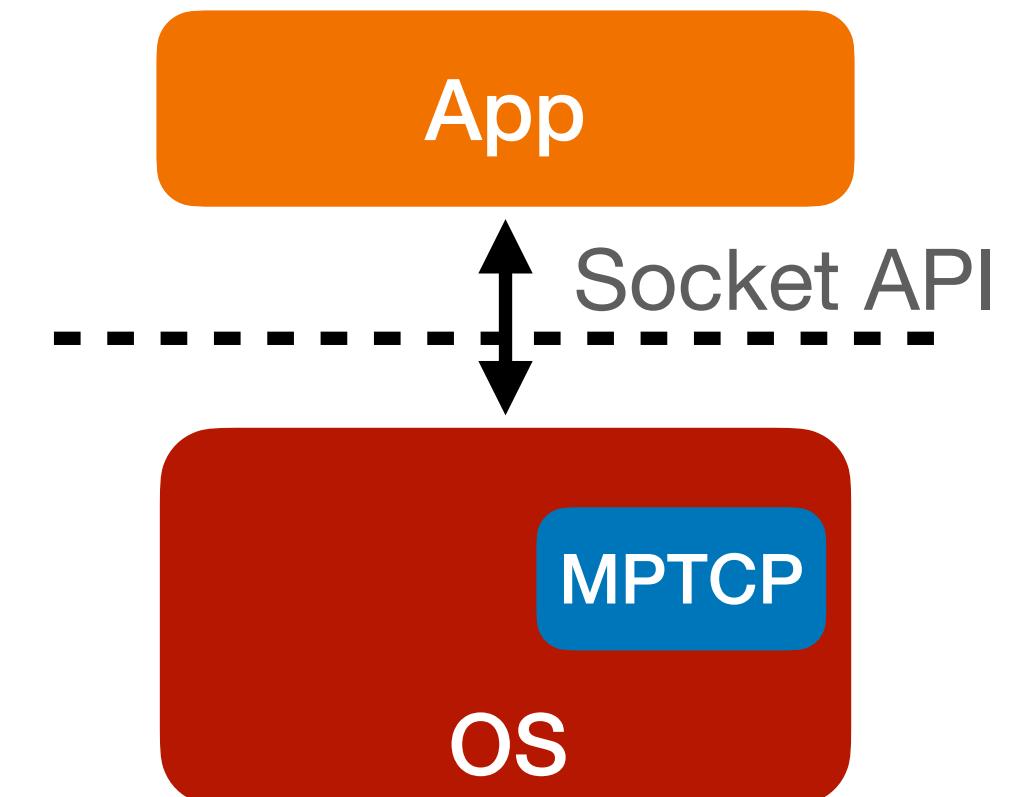
File size	20 MB	512 KB
Download time TCP (only blue path)	8.9 s	390 ms ~1.56x faster
Download time MPTCP (both paths)	5.69 s	480 ms ~1.25x slower

What Happened?

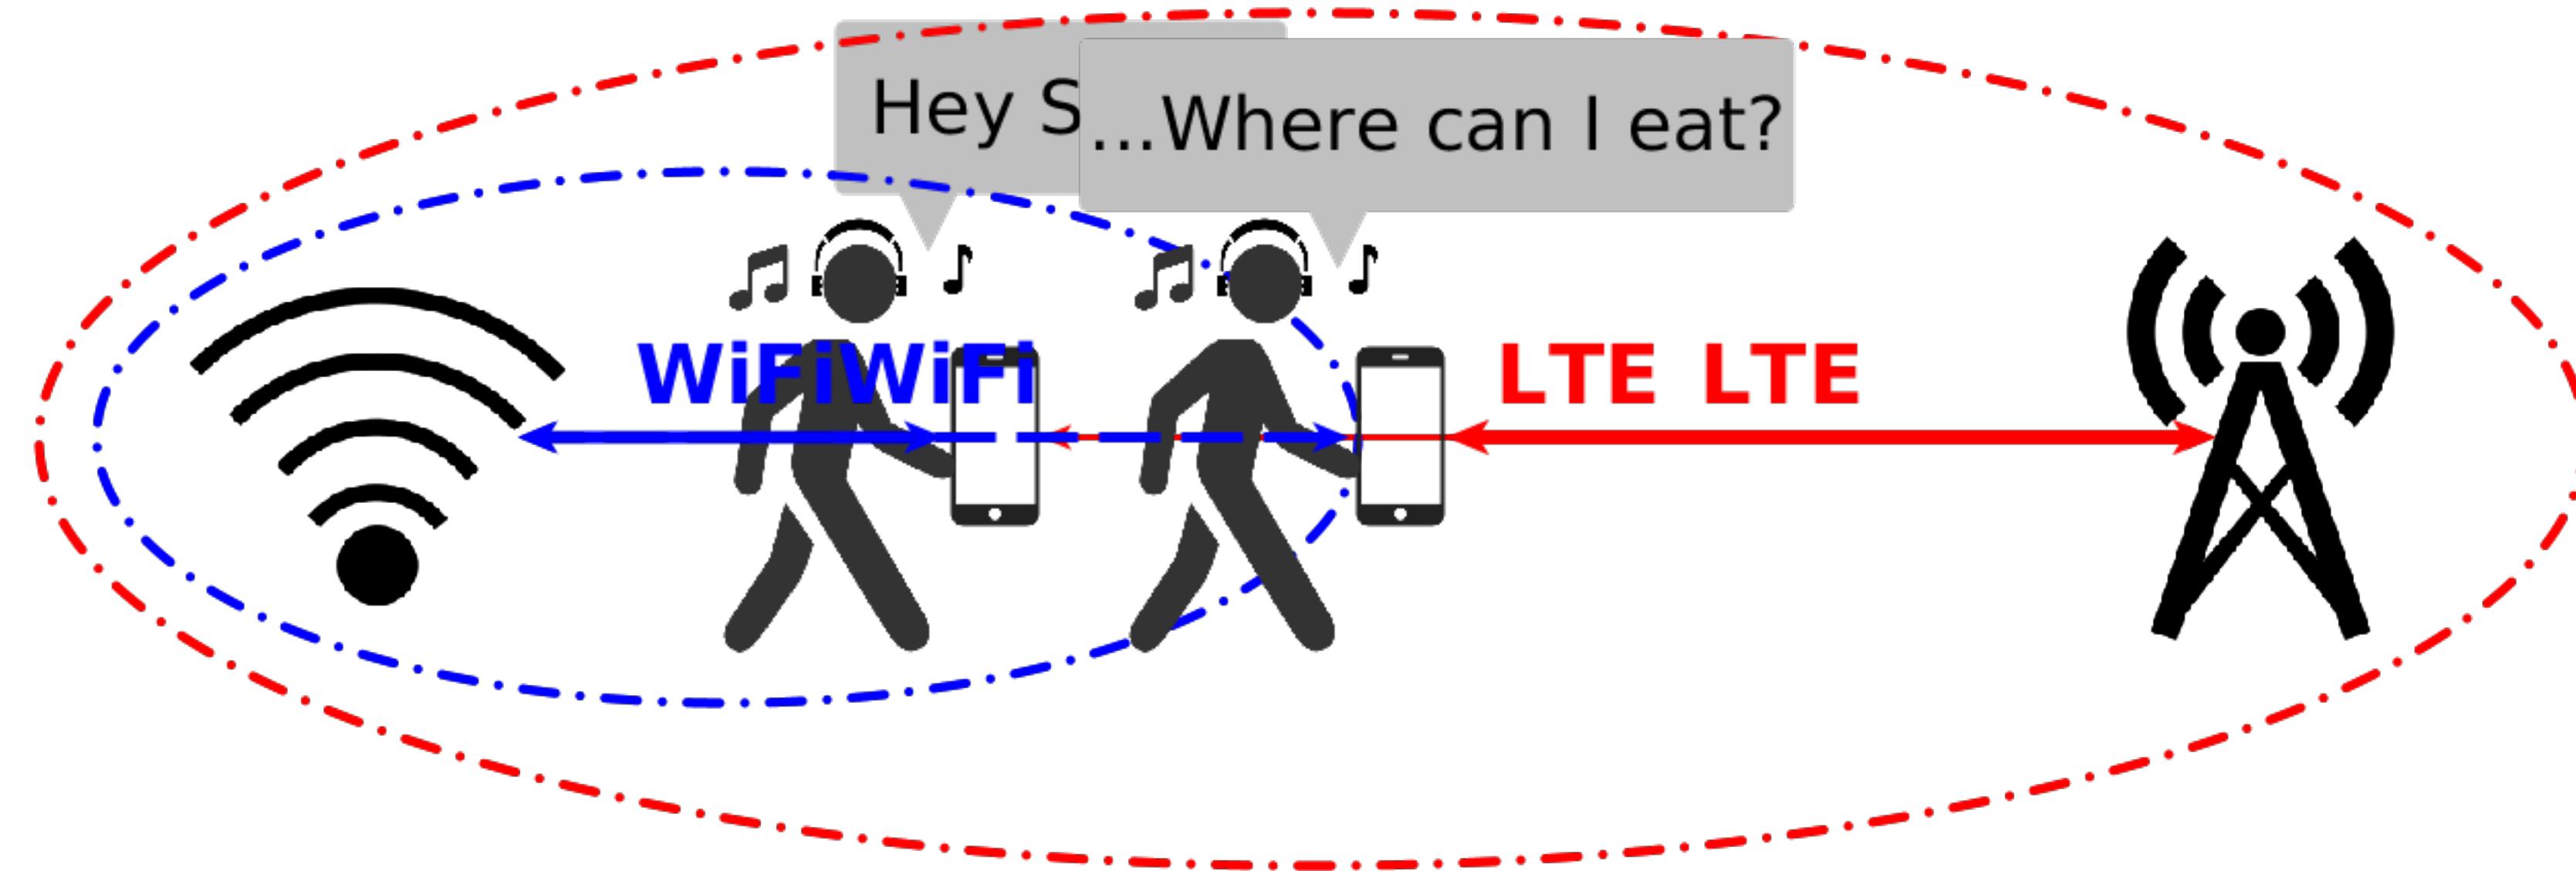


Can We Do Better?

- If known transfer size, avoid using slow path at the end of the transfer
 - But application should give data to the kernel...
- Other heuristics to avoid using slow path if it may block the fast one
 - BLEST [Networking'16], ECF [CoNext'17],...



Interactive Applications under Mobility



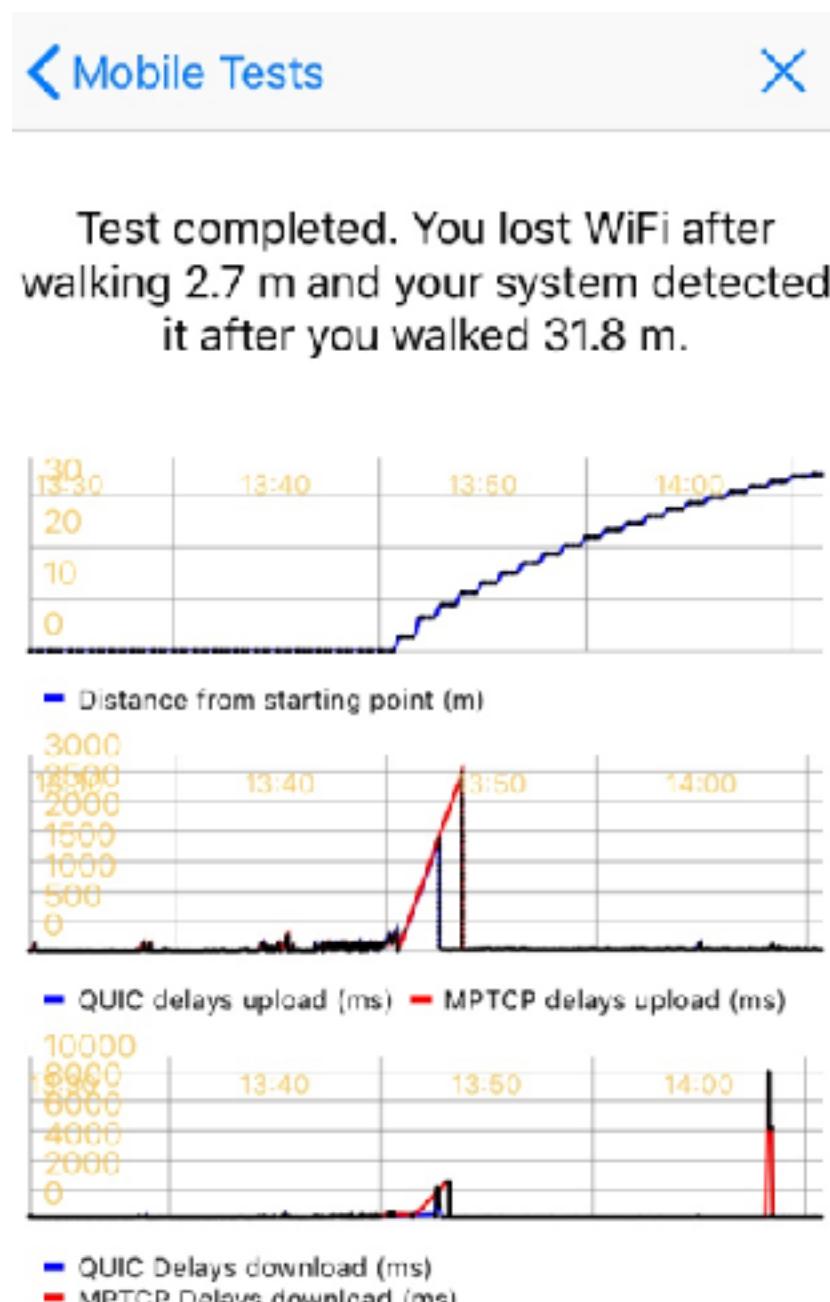
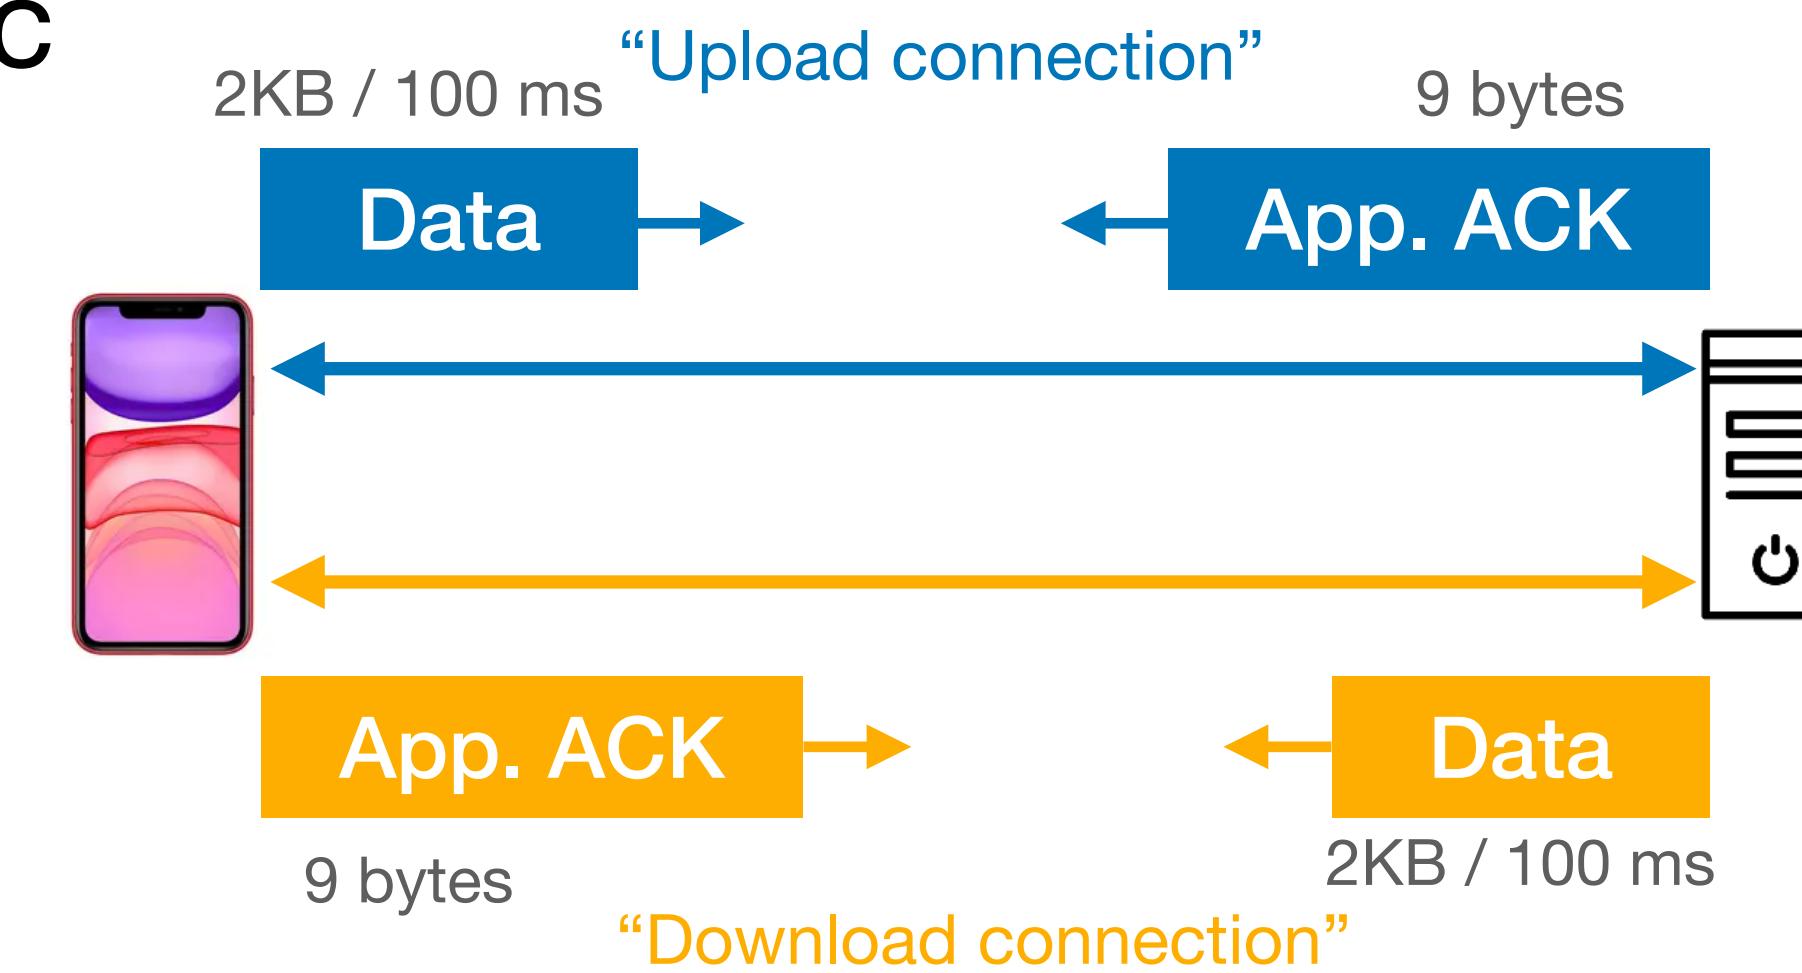
Interactive Application under Mobility

Evaluating Application Perceived Latency

- Active measurements through MultipathTester iOS application
- Run traffic until client device declares Wi-Fi as lost
 - 1100 test runs, +250 unique devices
 - Request-response like traffic

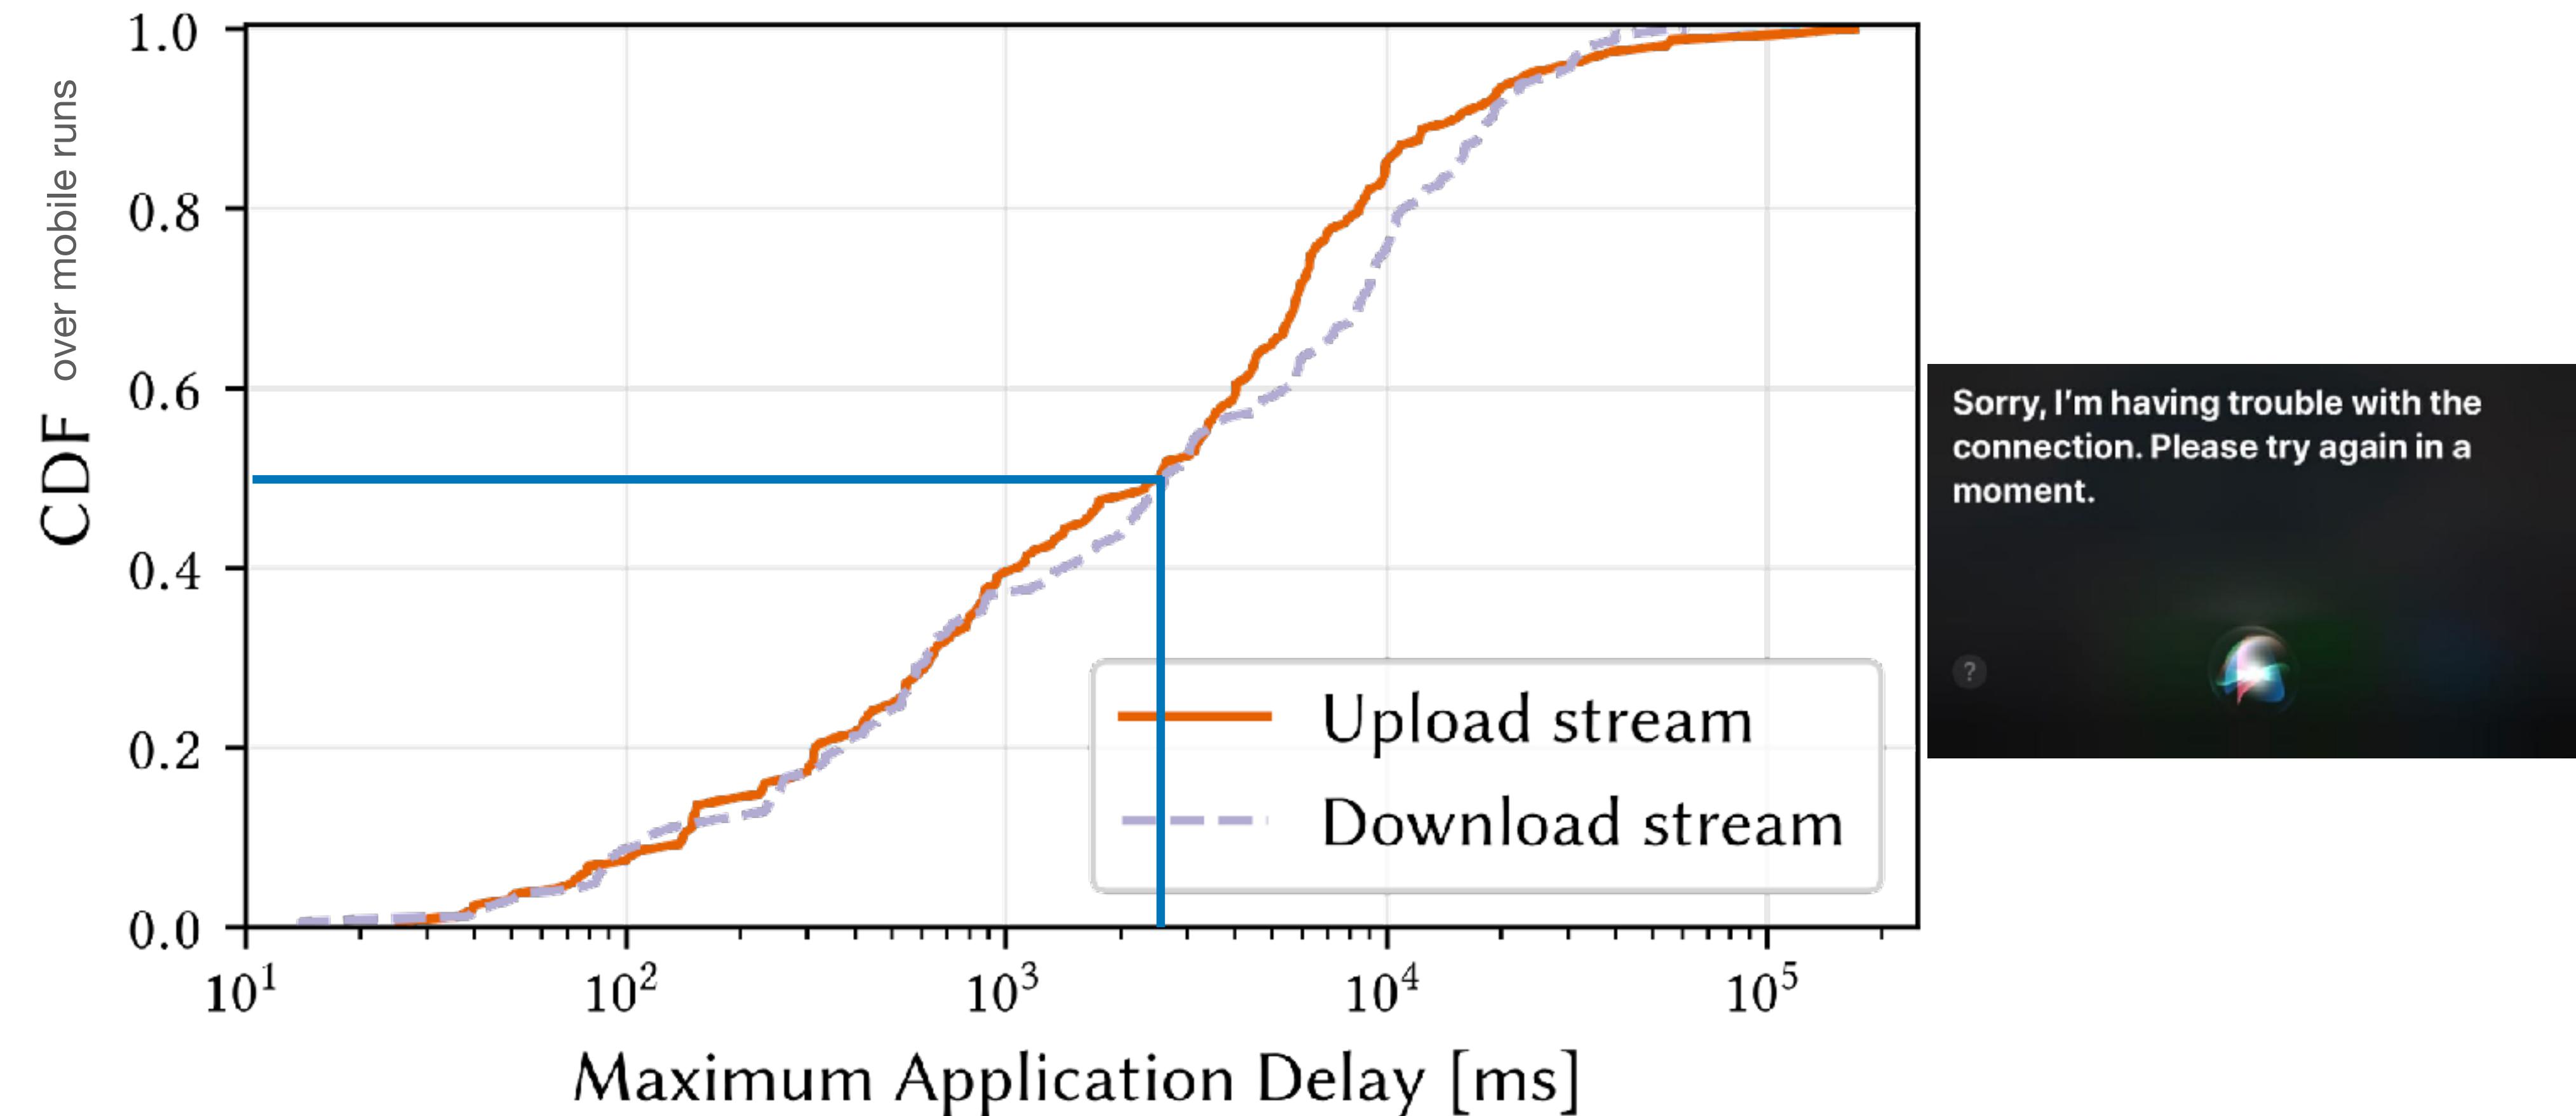


Apple's MPTCP



Linux's MPTCP

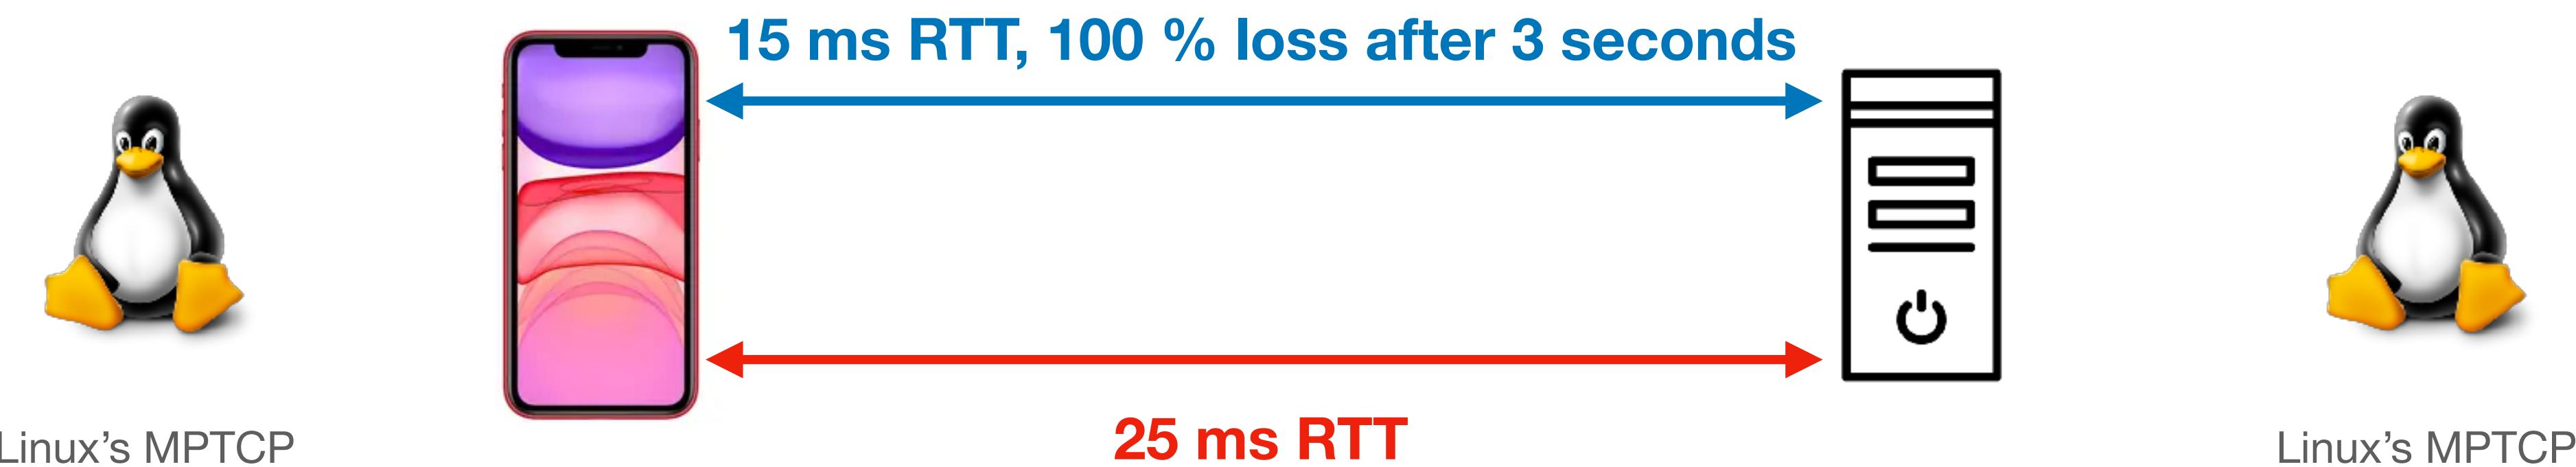
Applications Perceive Network Handover



Possible unresponsiveness for several seconds

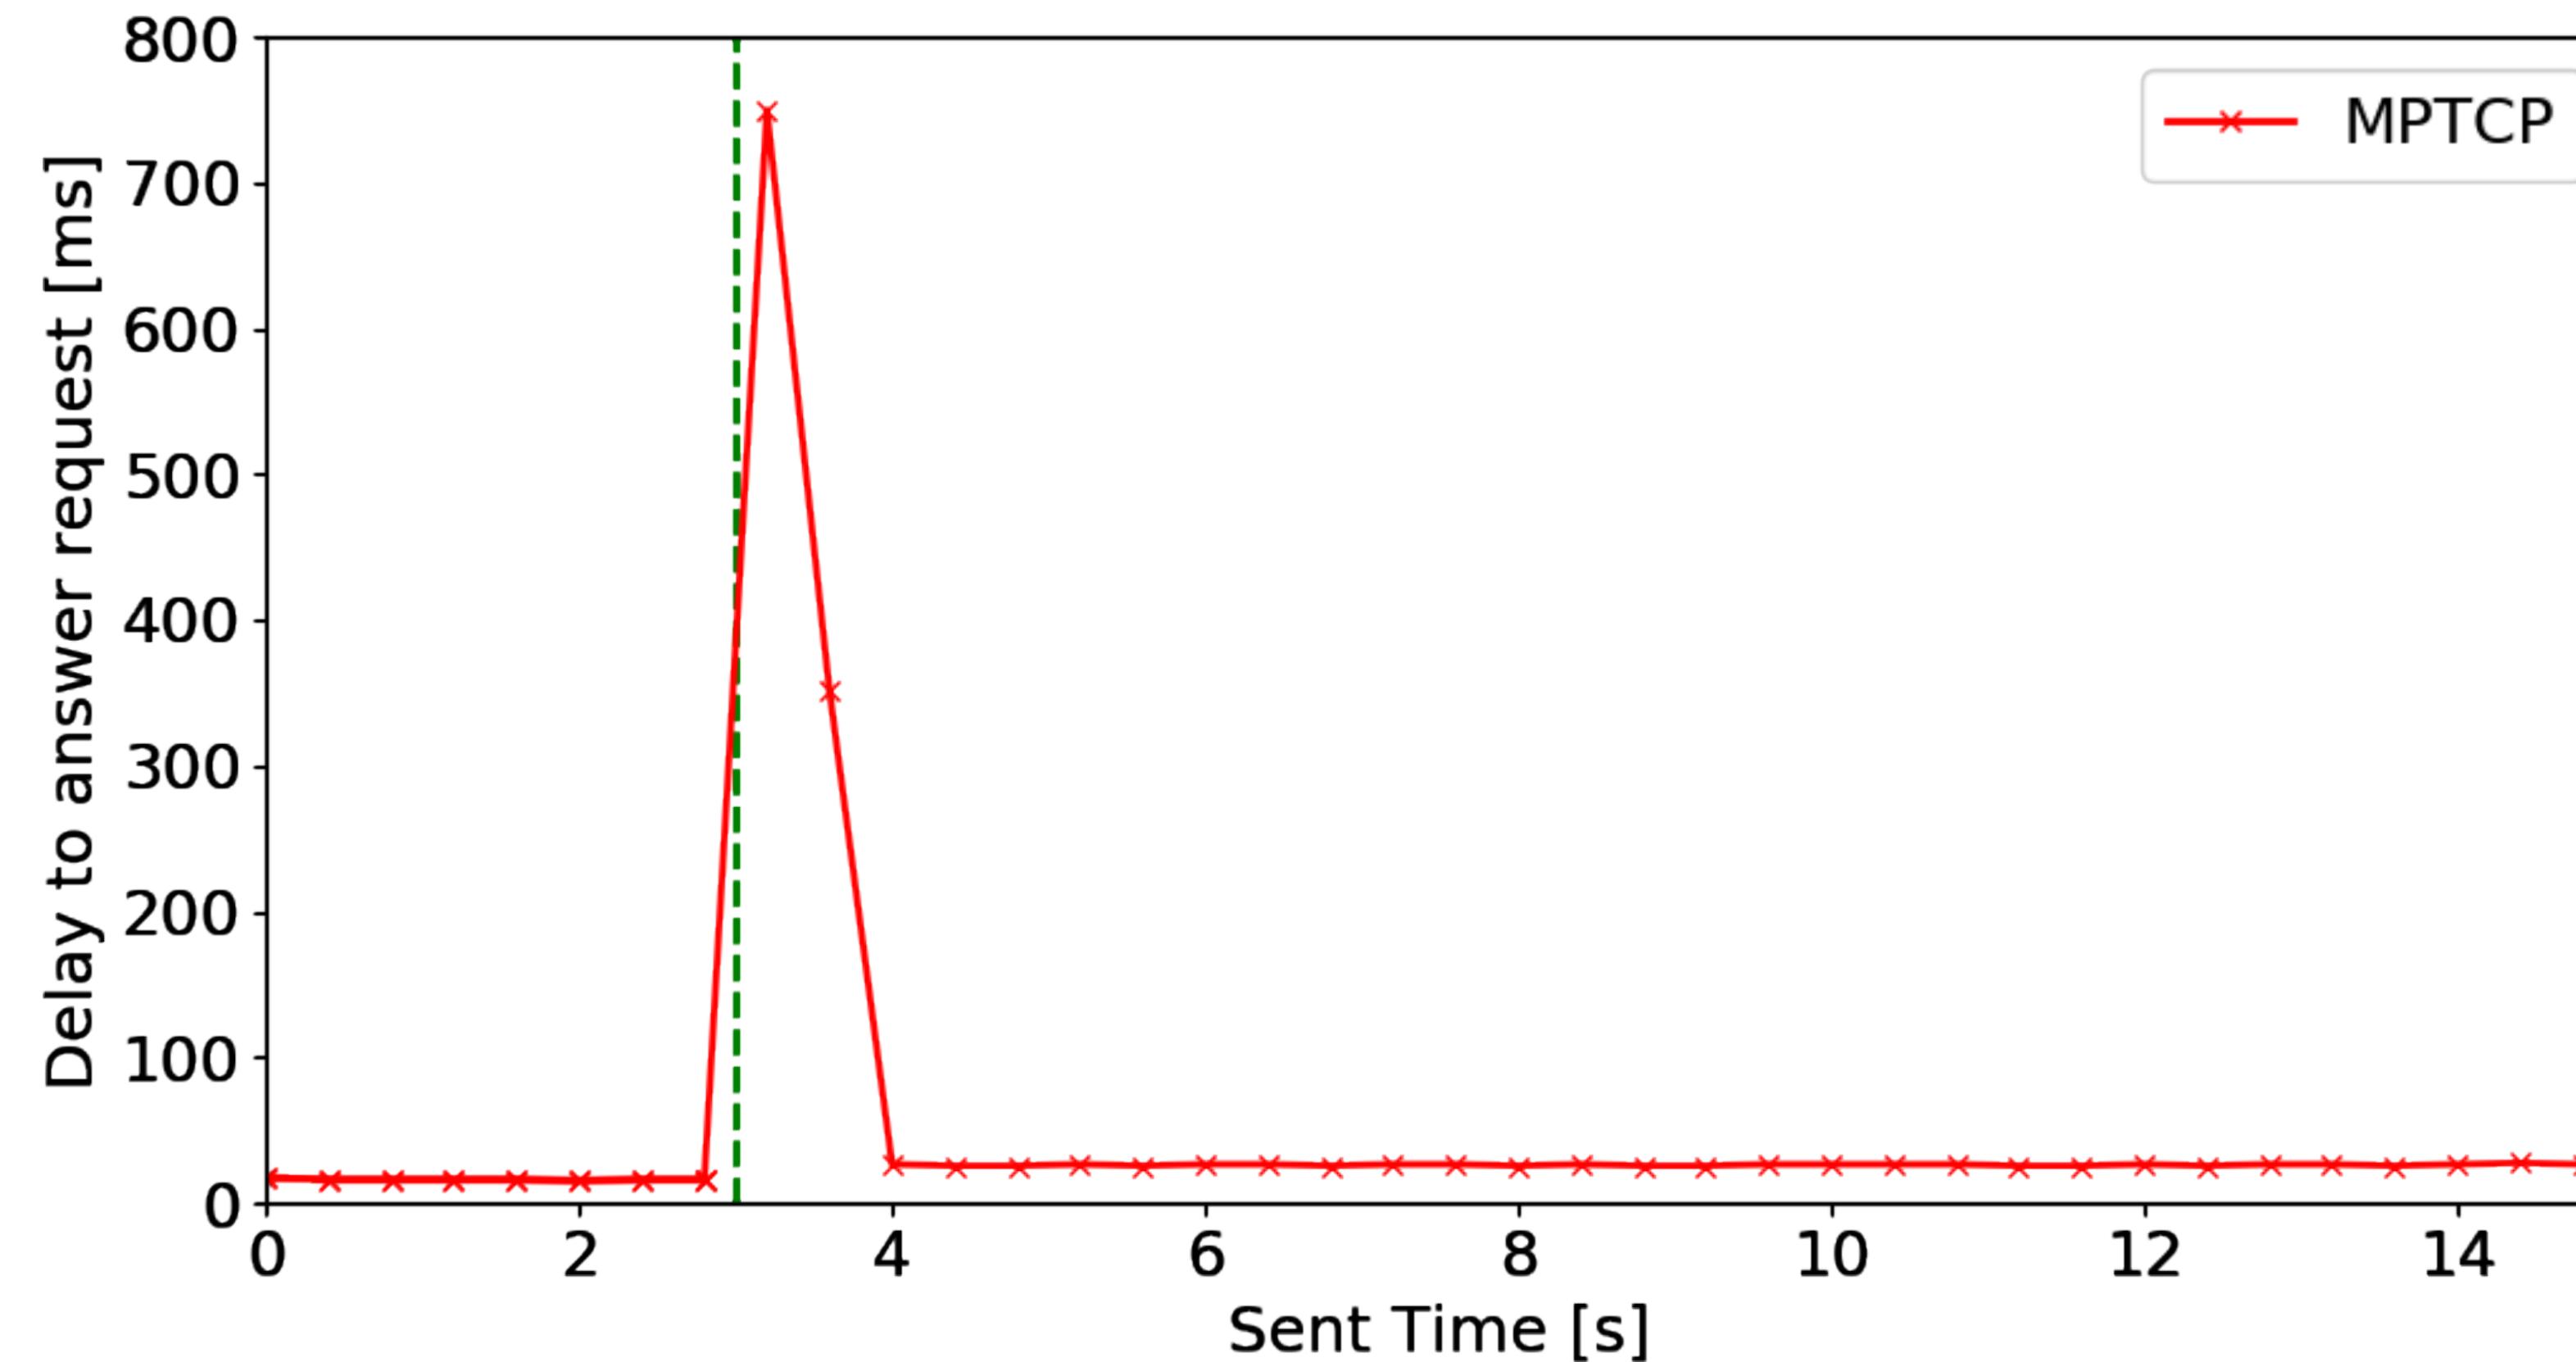
Understanding The Root Cause

- Consider a simpler request/response traffic
 - 750 bytes request/response every 400 ms
 - Measure the delay seen by the client



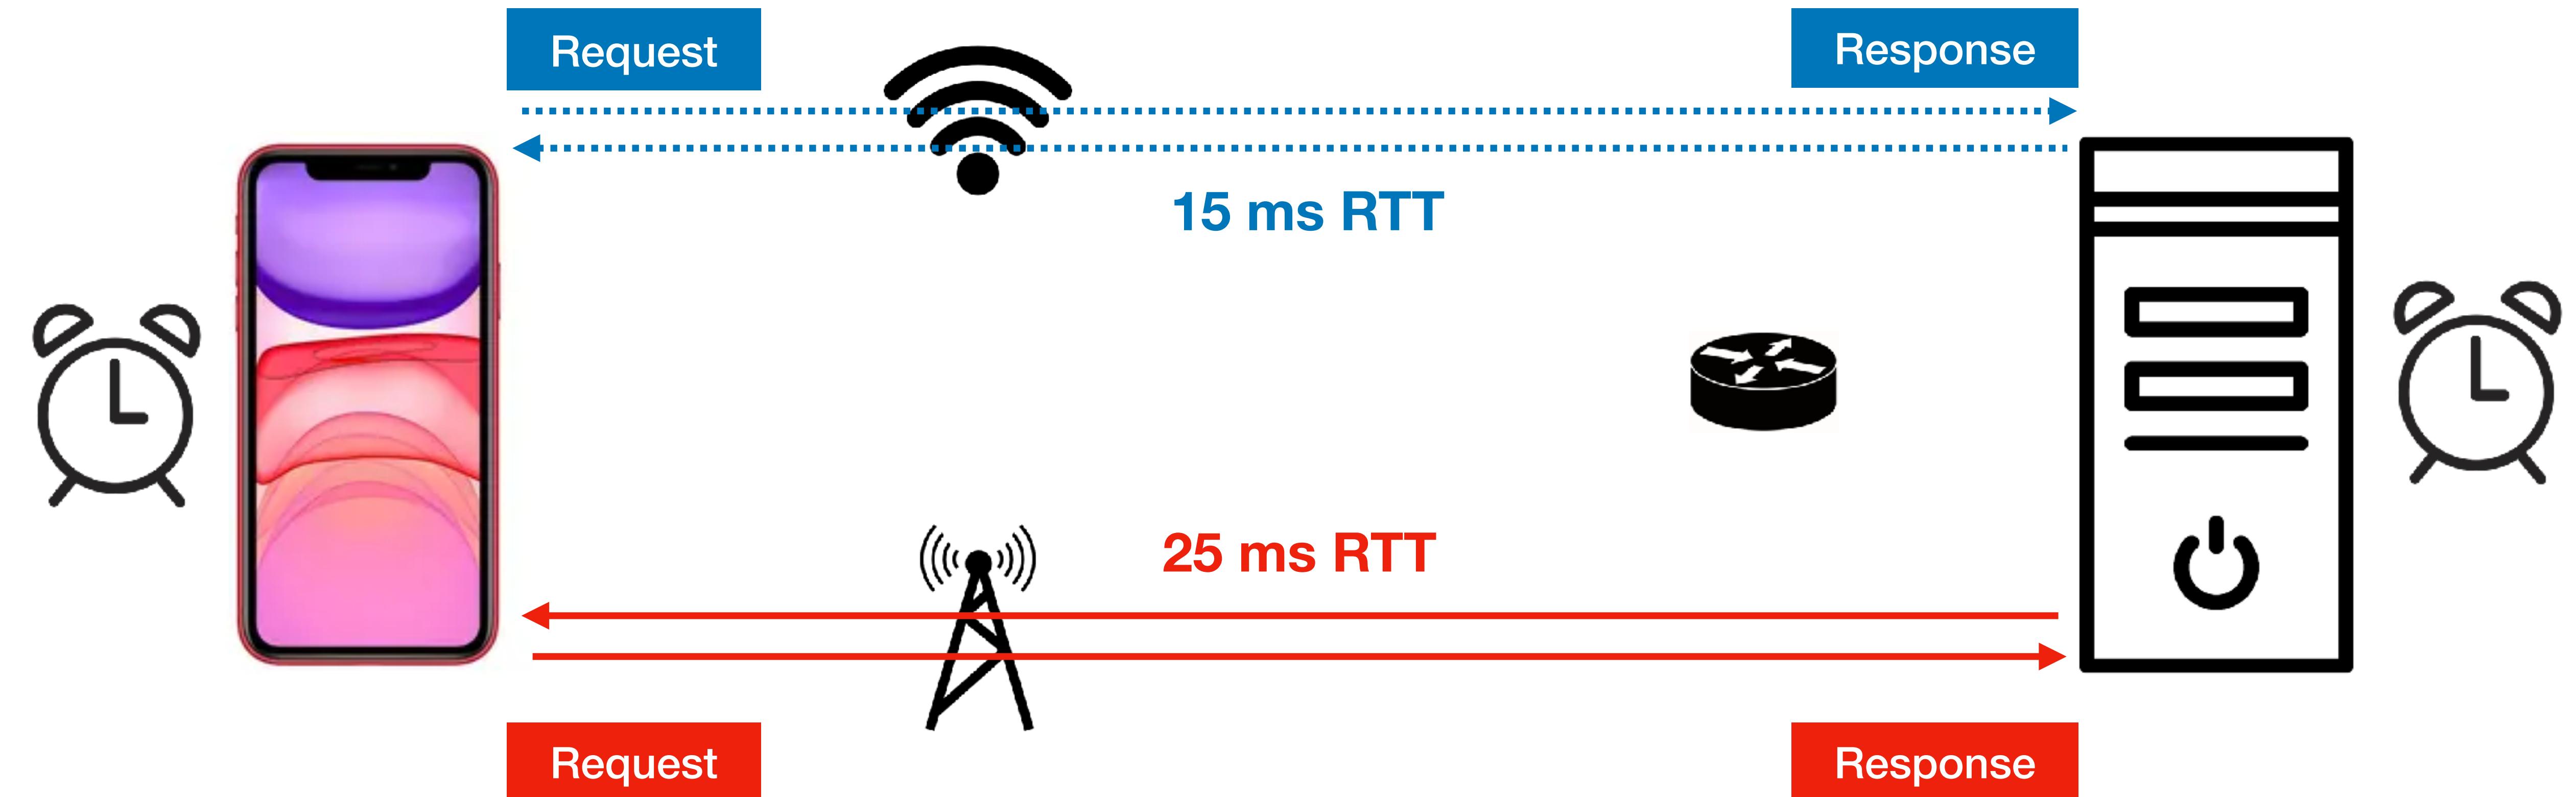
Multipath TCP Handover

Fullmesh path manager
Lowest-latency first scheduler



What Happened Here?

Fullmesh path manager
Lowest-latency first scheduler



Retransmission timeout at both sides

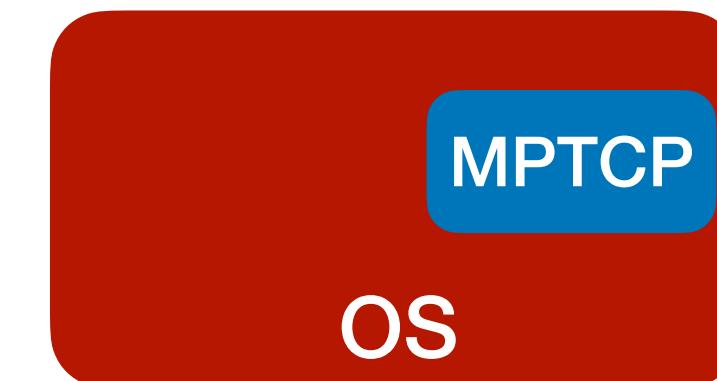
Could We Adapt Multipath Algorithms?



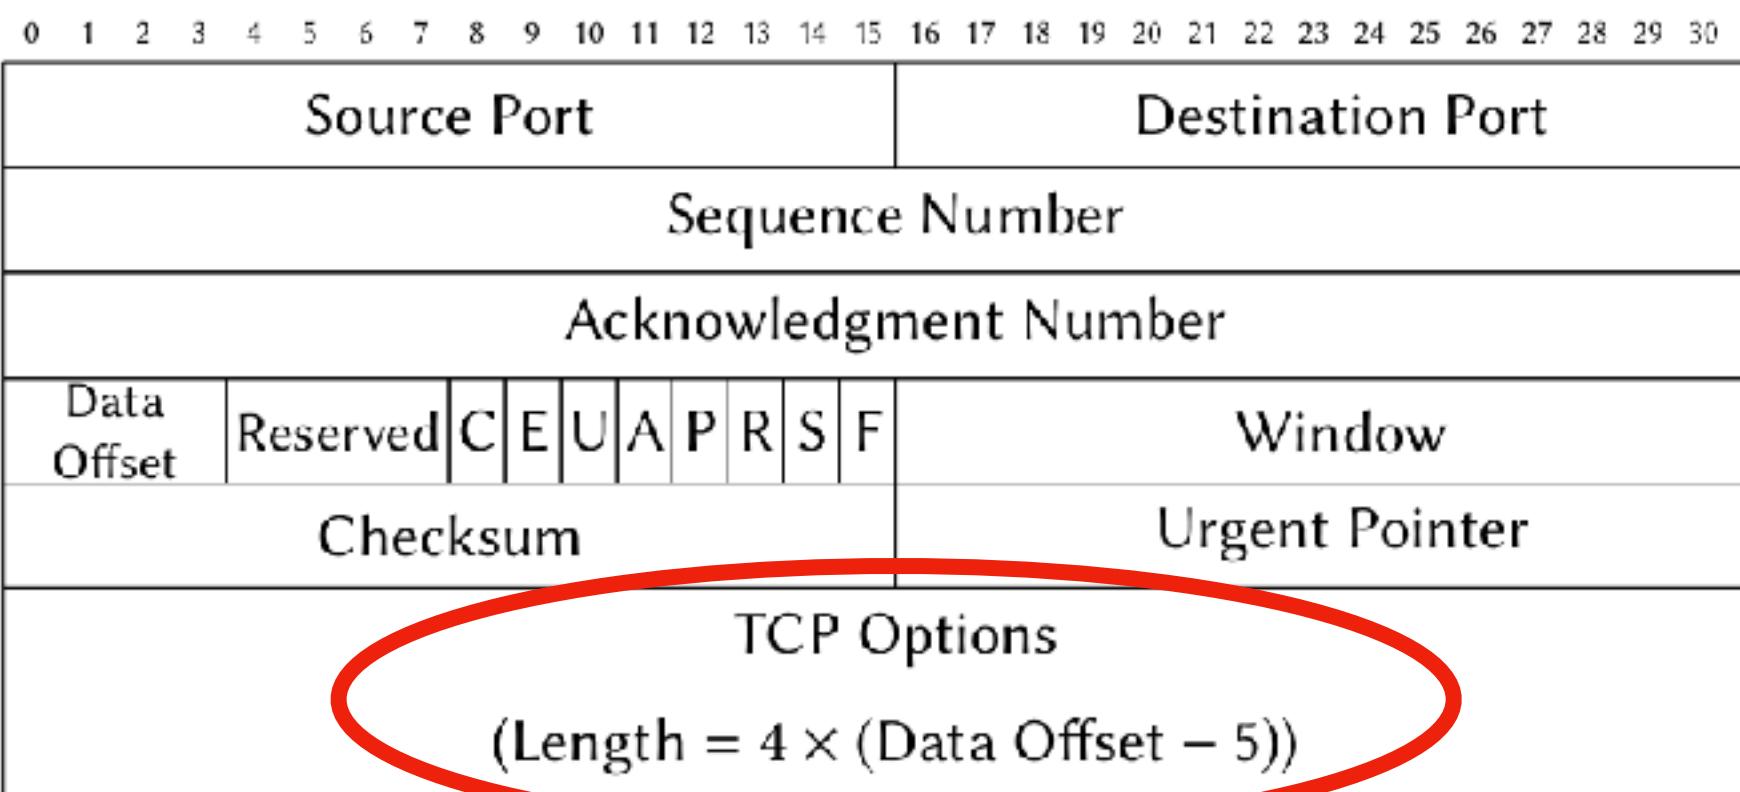
- Server-side scheduler [Networking'18]
 - Remove server-side RTO in mobile scenarios (~350 ms delay vs. 750 ms)
 - Relevancy for the file downloader case?

An Application-Tailored Multipath TCP?

- Deployment issues



- TCP Limitations



Max 40 bytes



The QUIC Transport Protocol

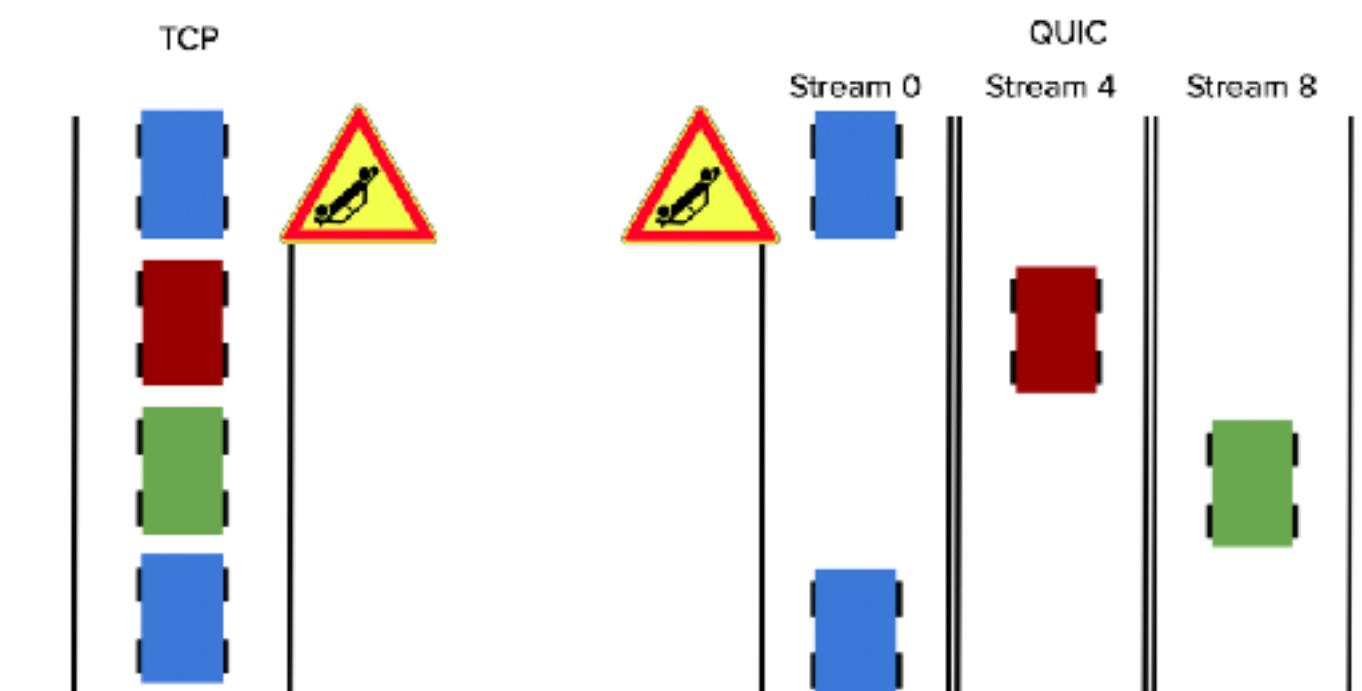
The QUIC Protocol

Since 2014

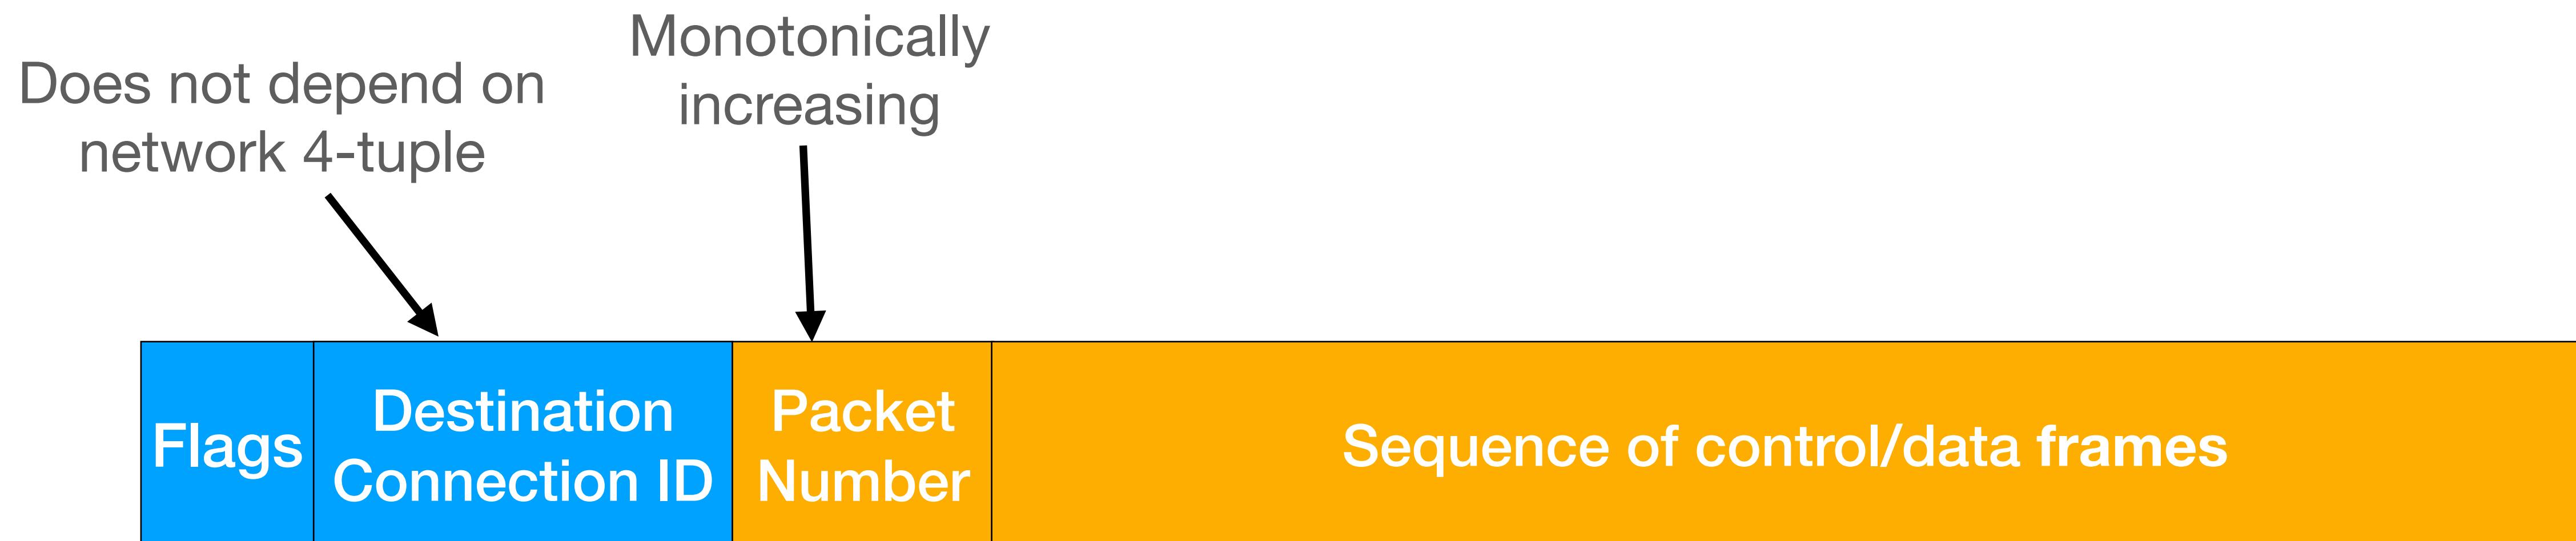
- TCP/TLS 1.3 replacement atop UDP
- Reliable
Encrypted
- Unreliable
- Unidirectional Connection IDs (IDs)



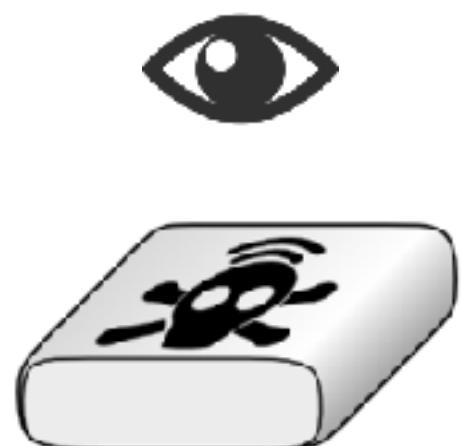
- Multistreaming support
- 0-RTT connection establishment (most of the time)
- Extension negotiation during handshake



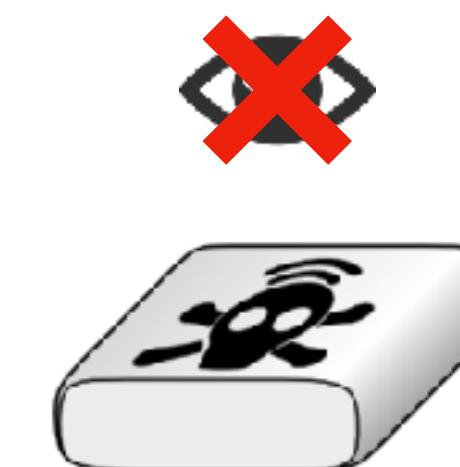
A QUIC Packet



Cleartext Header

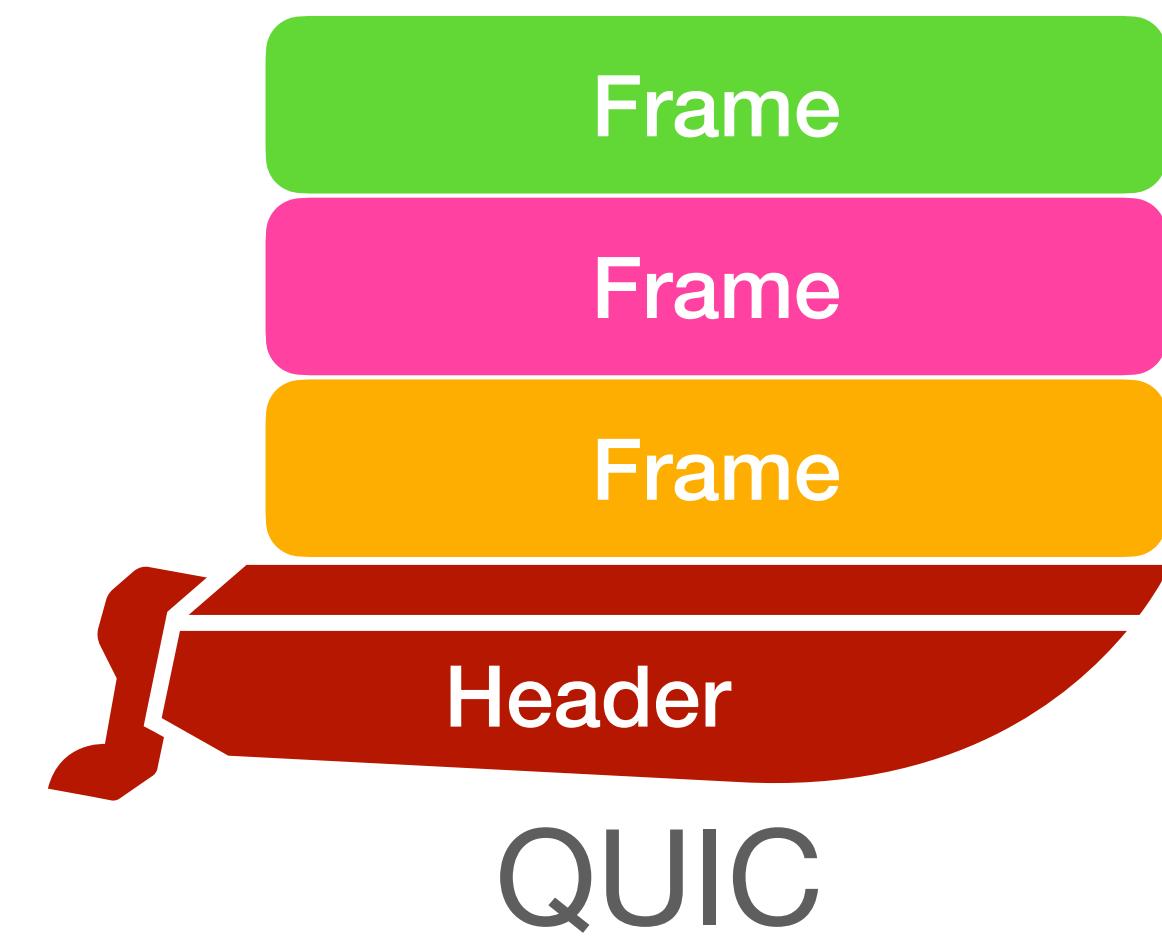
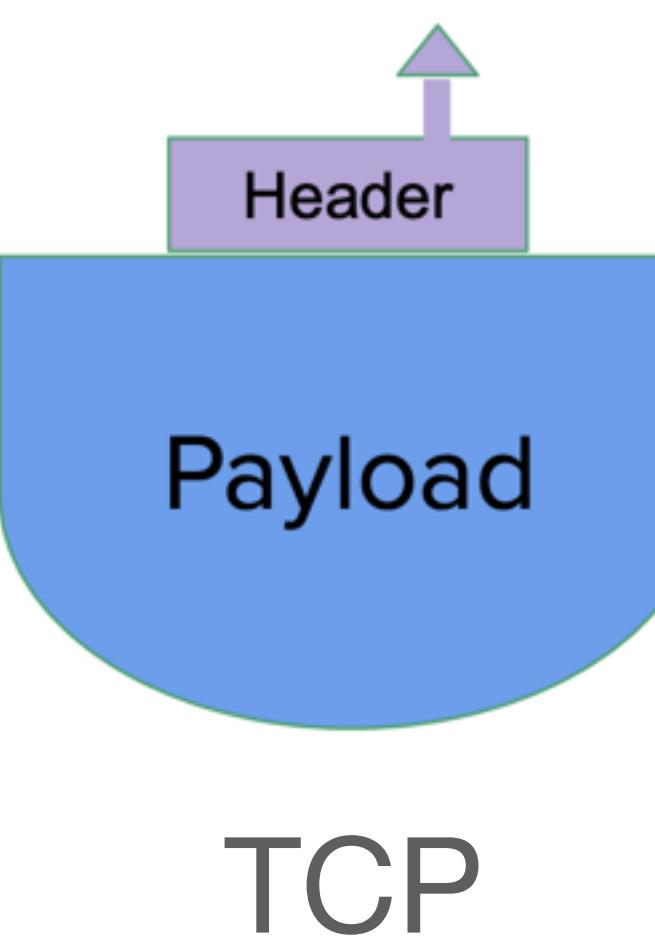


Encrypted Payload

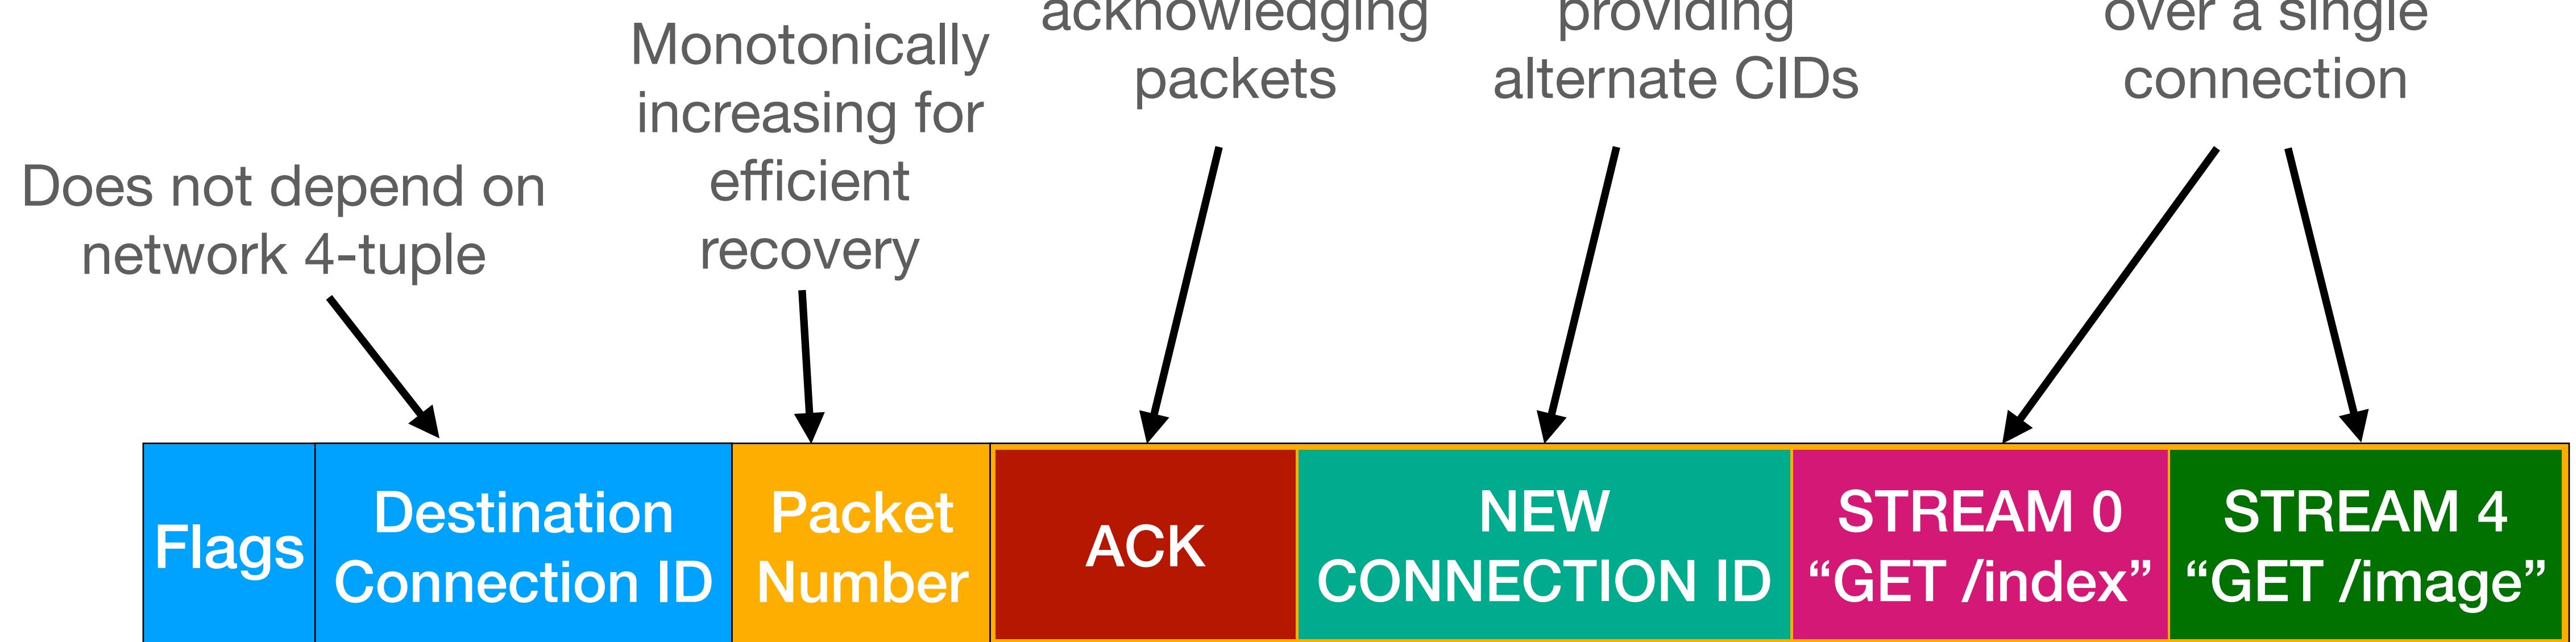


Easier deployment

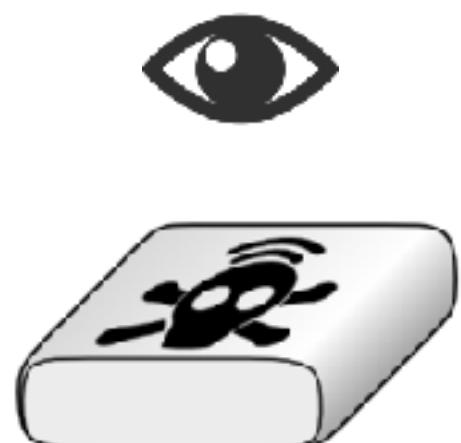
A QUIC Packet Seen by the Sea



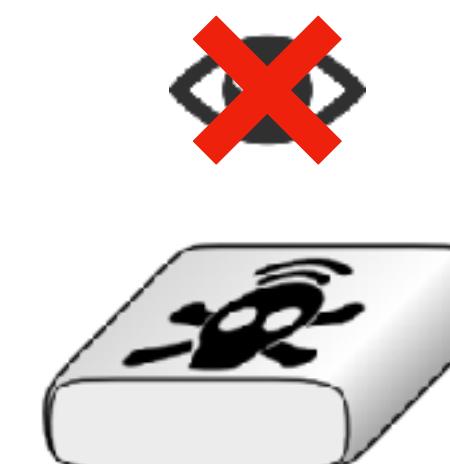
A QUIC Packet



Cleartext Header



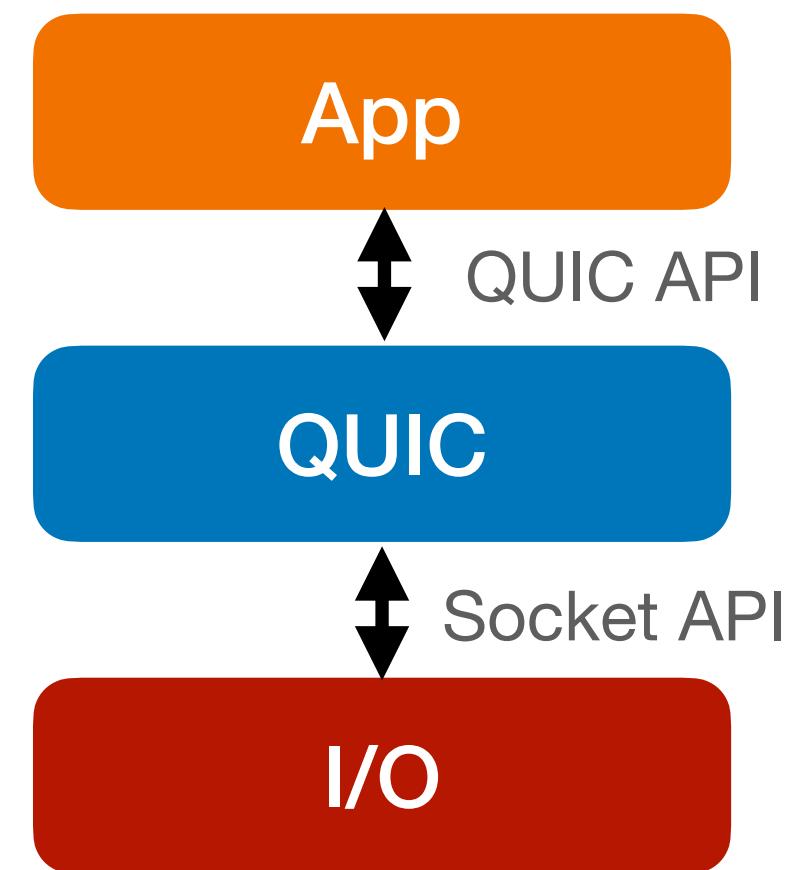
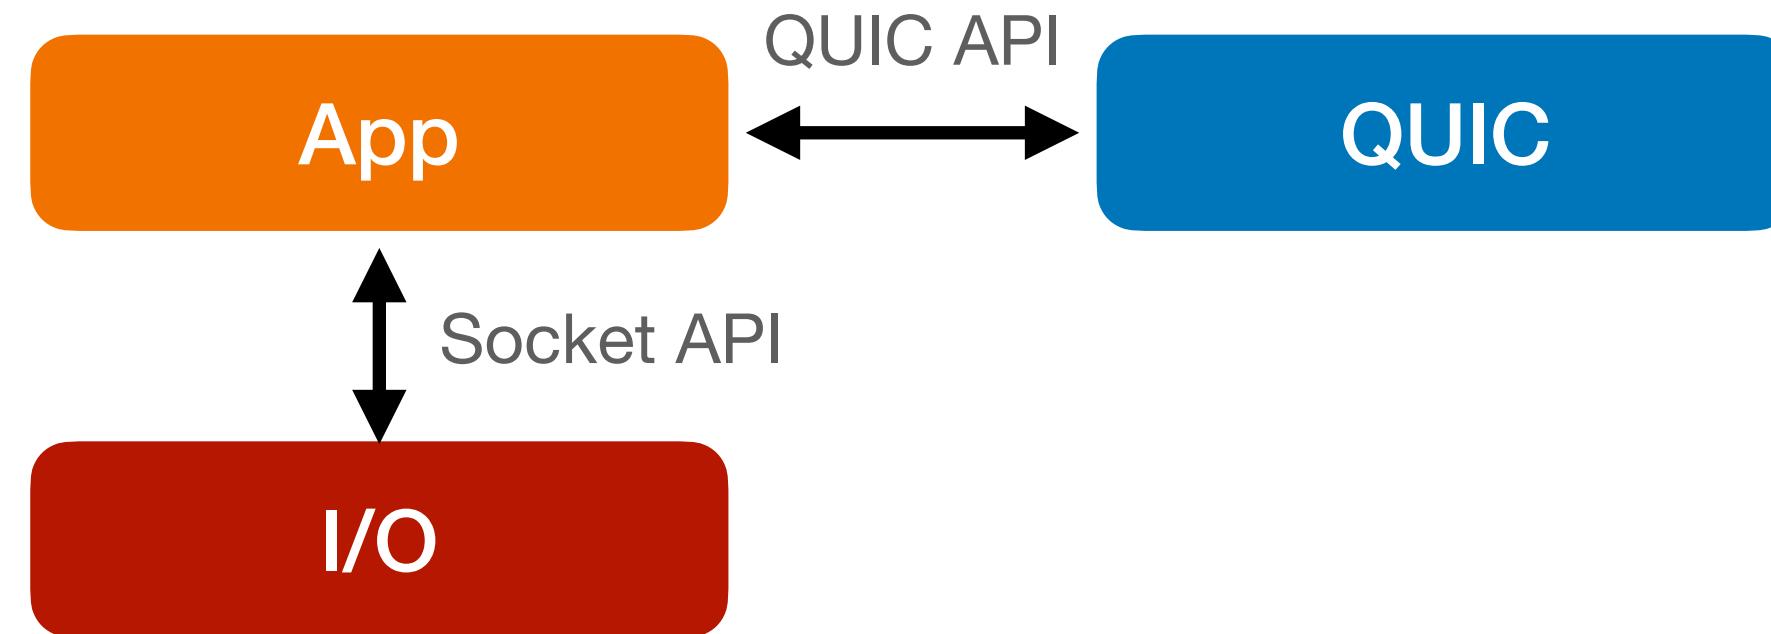
Encrypted Payload



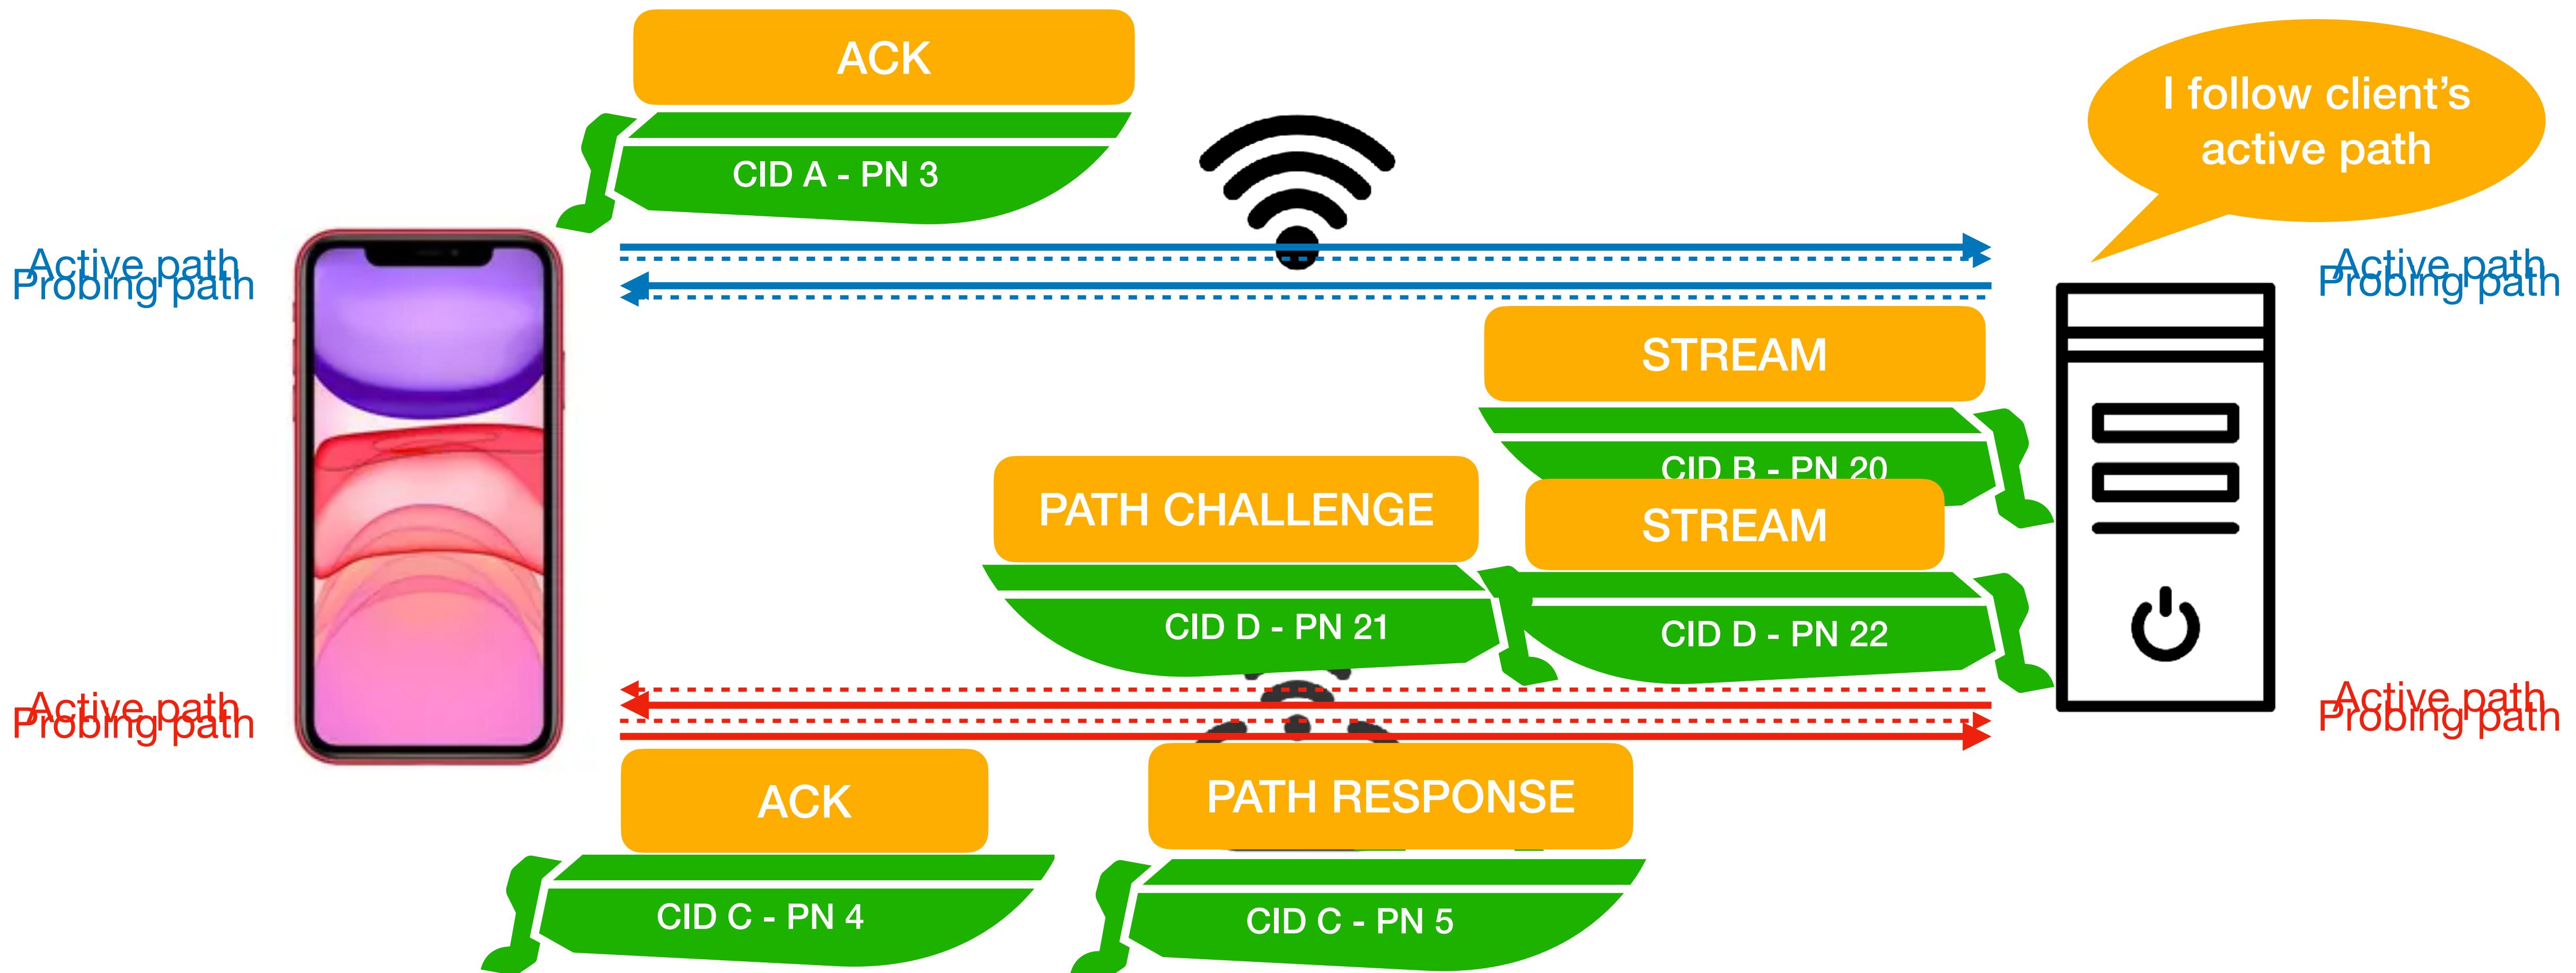
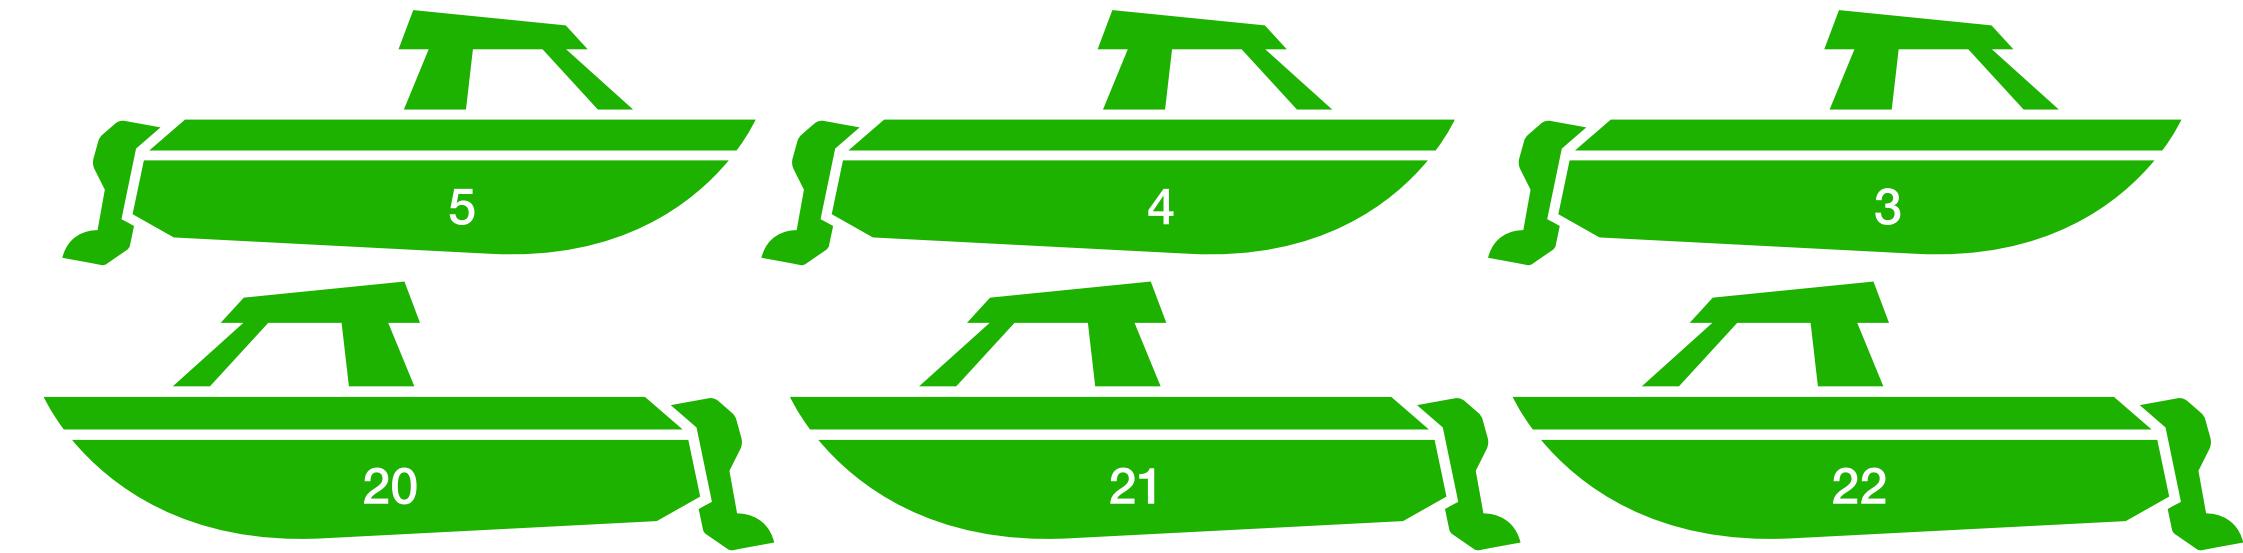
Easier deployment

QUIC Implementations

- +20 publicly known implementations
 - Meta's mvfst, Microsoft's MSQUIC, AppleQUIC, Cloudflare's quiche,...
 - User-space implementations
 - Can be easily integrated with applications
 - But apps should be adapted



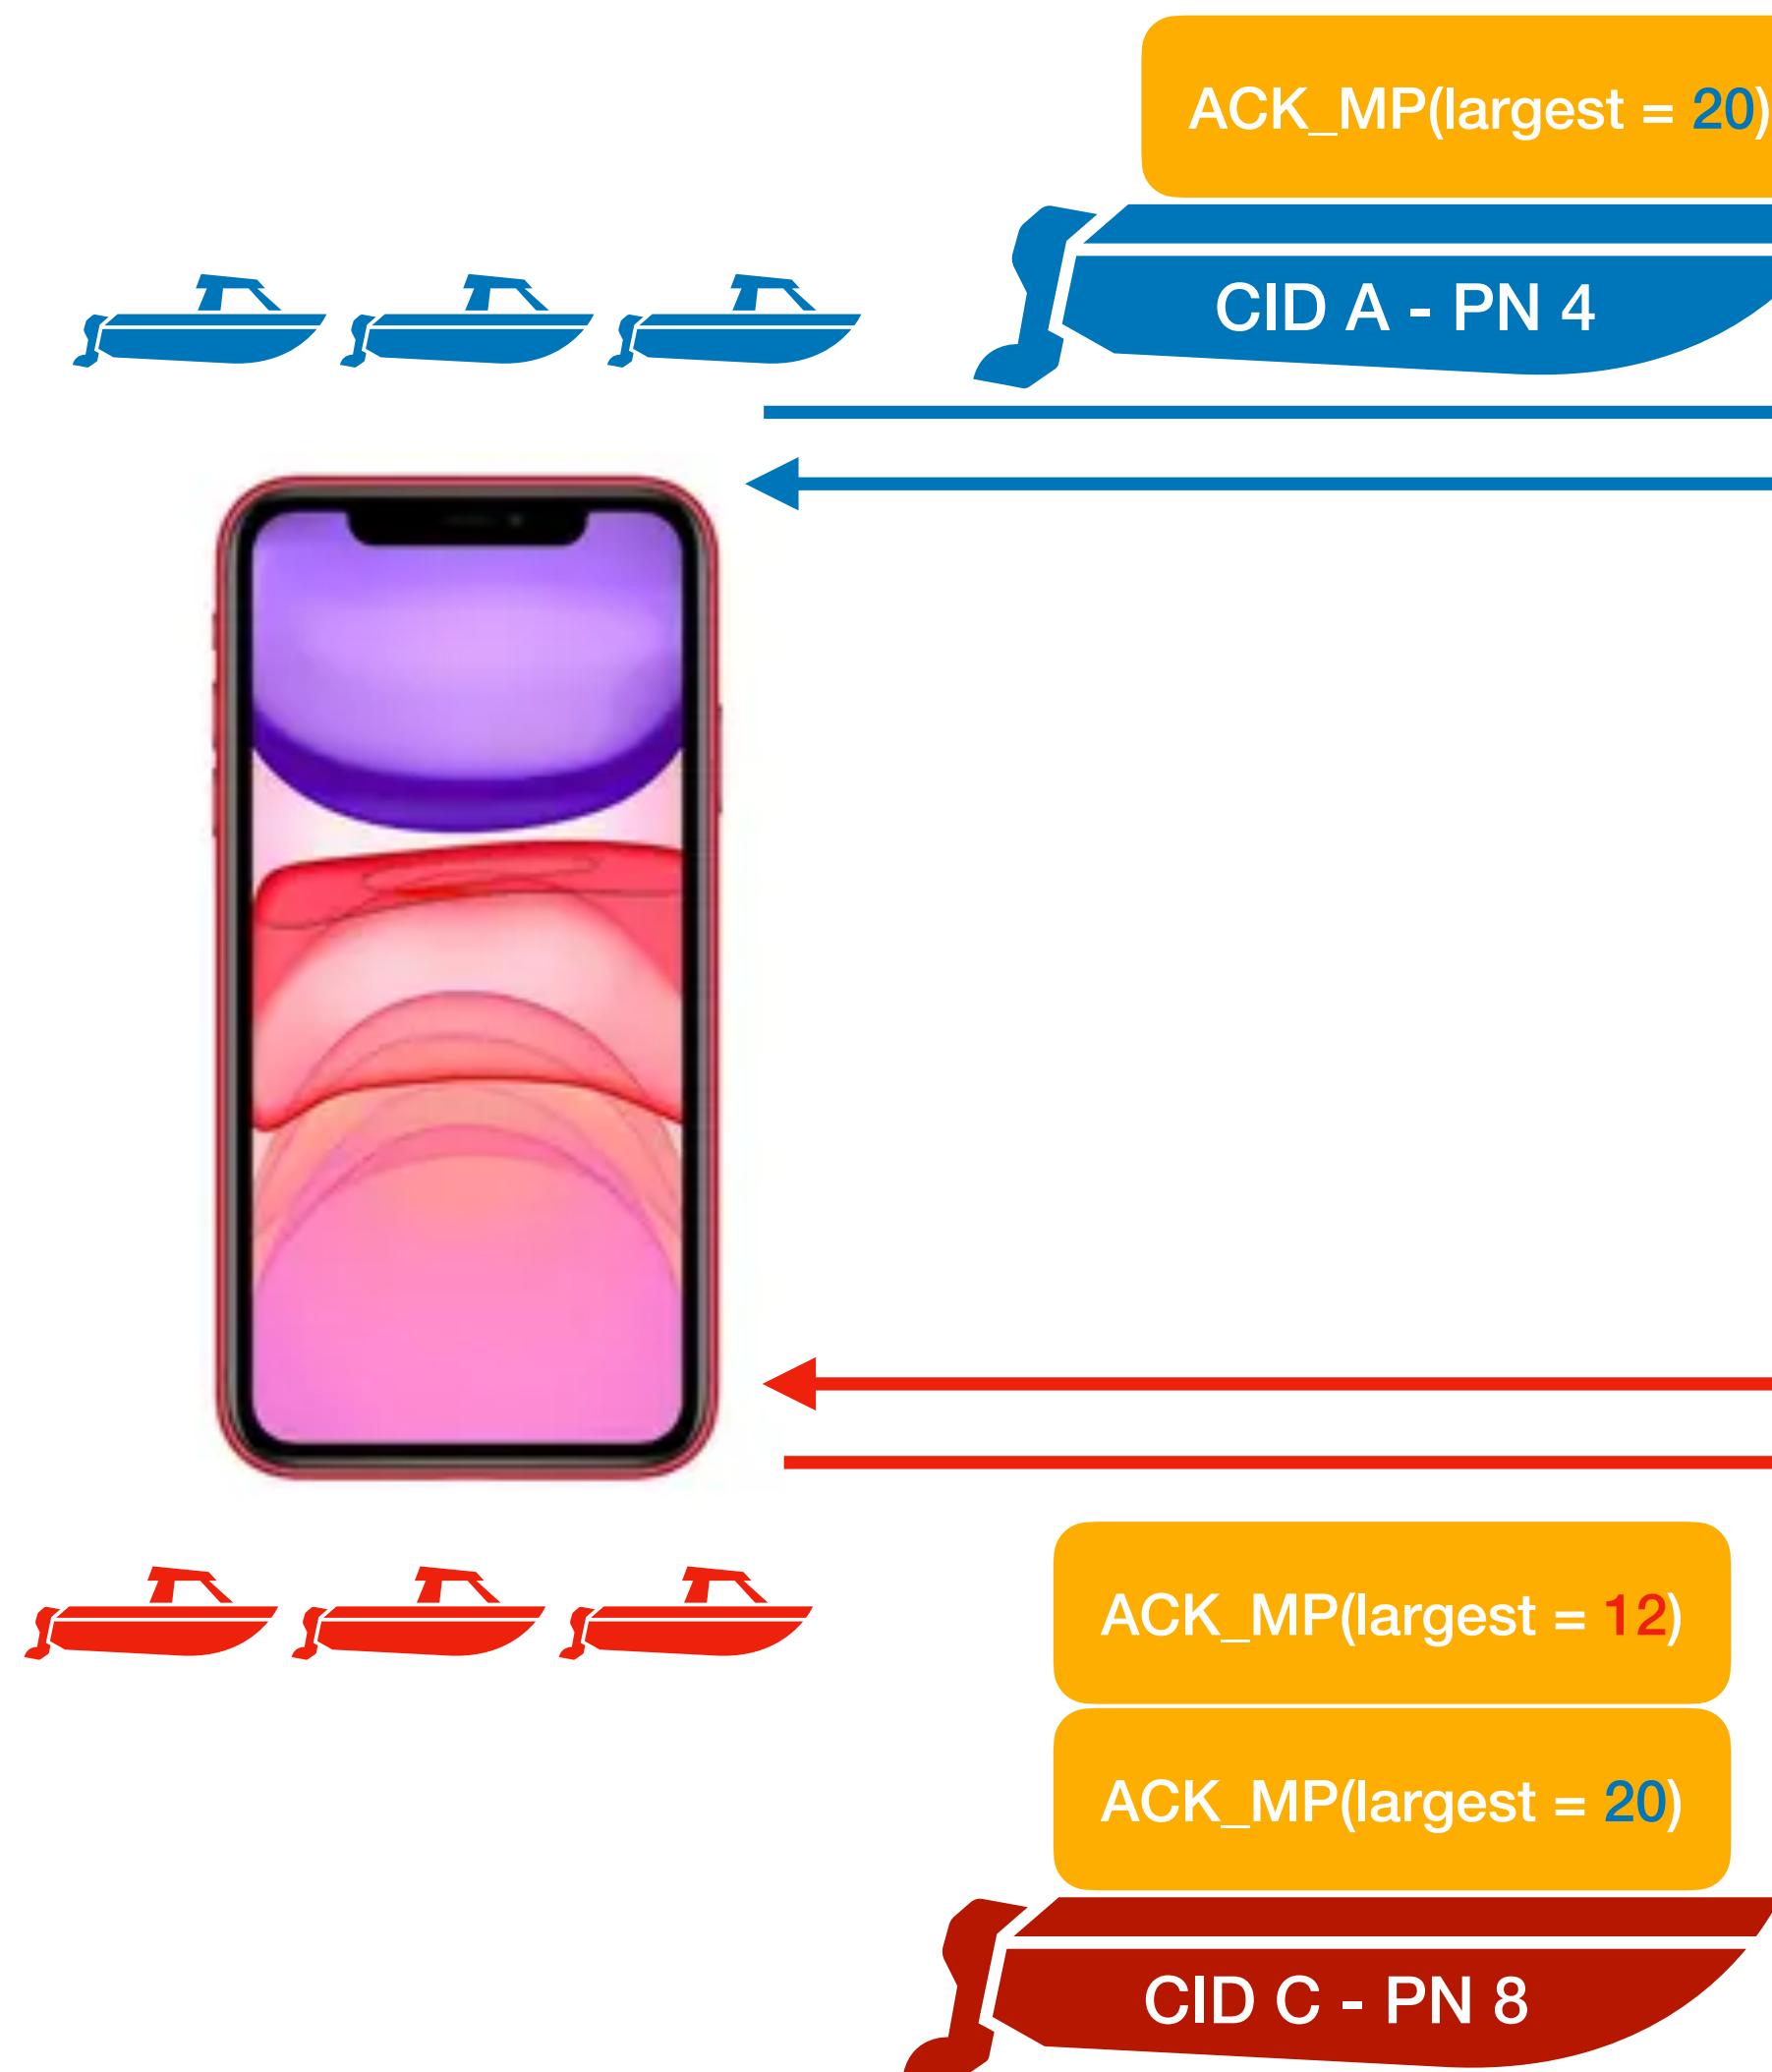
Connection Migration



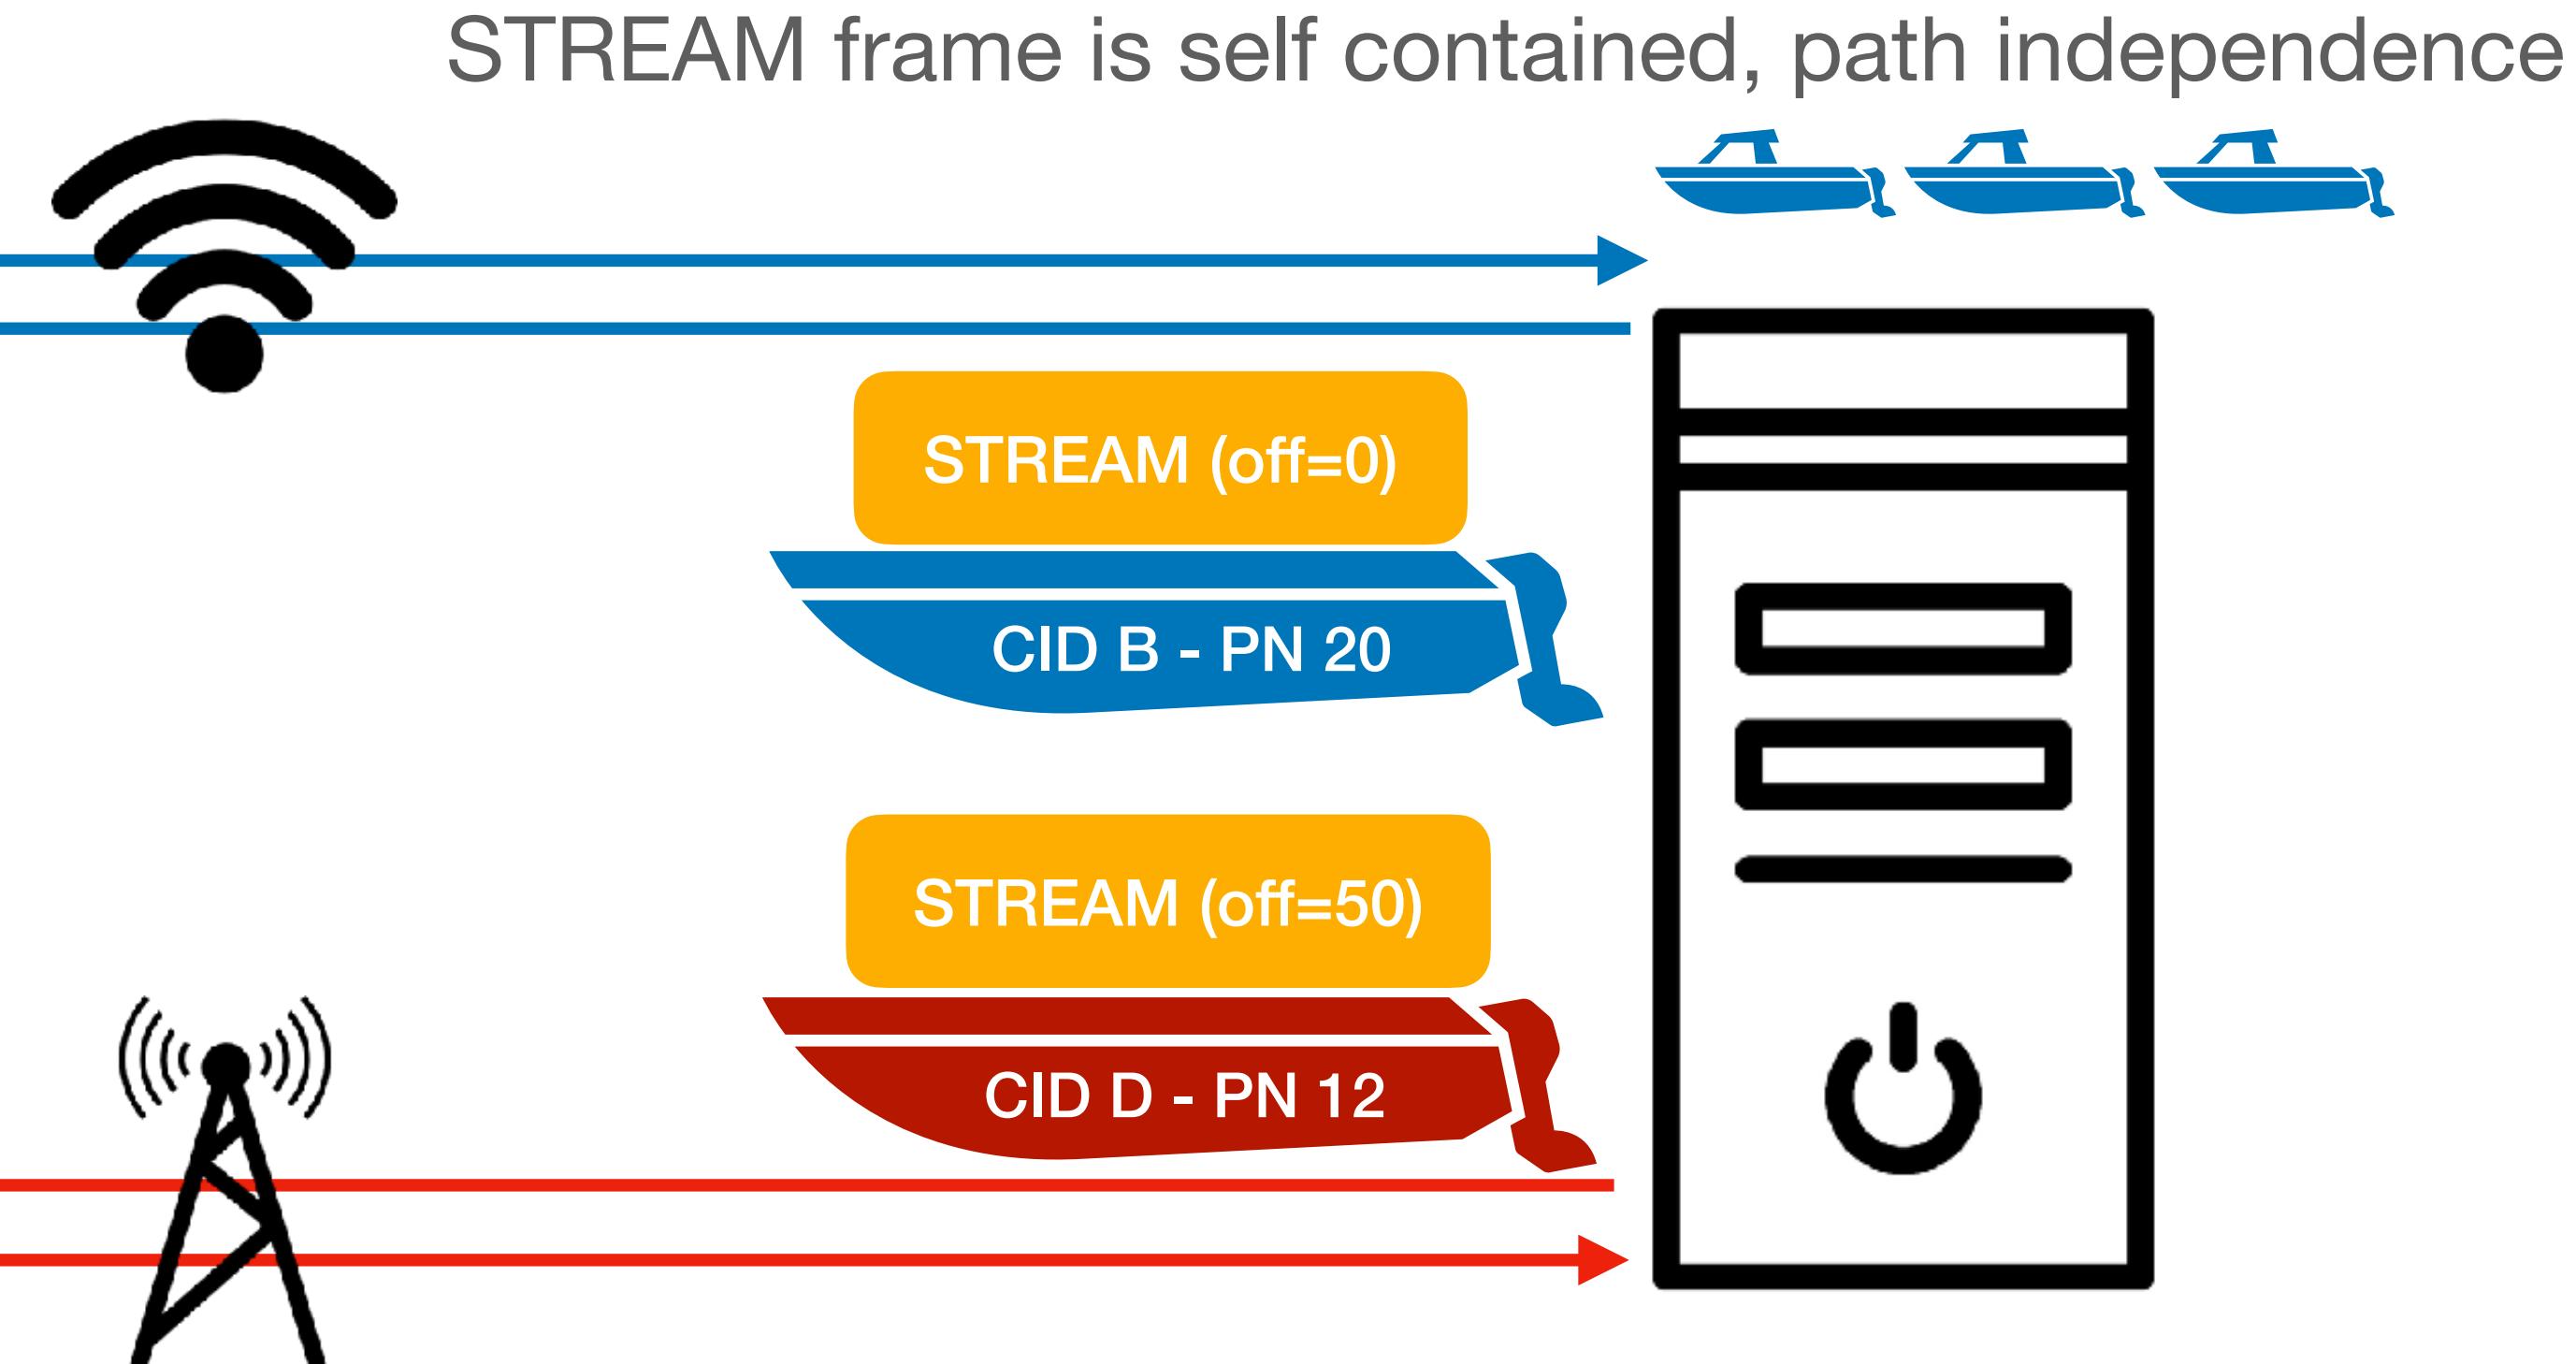
Limitation: only one “data path” at a time

Bringing Multipath to QUIC

Multipath QUIC



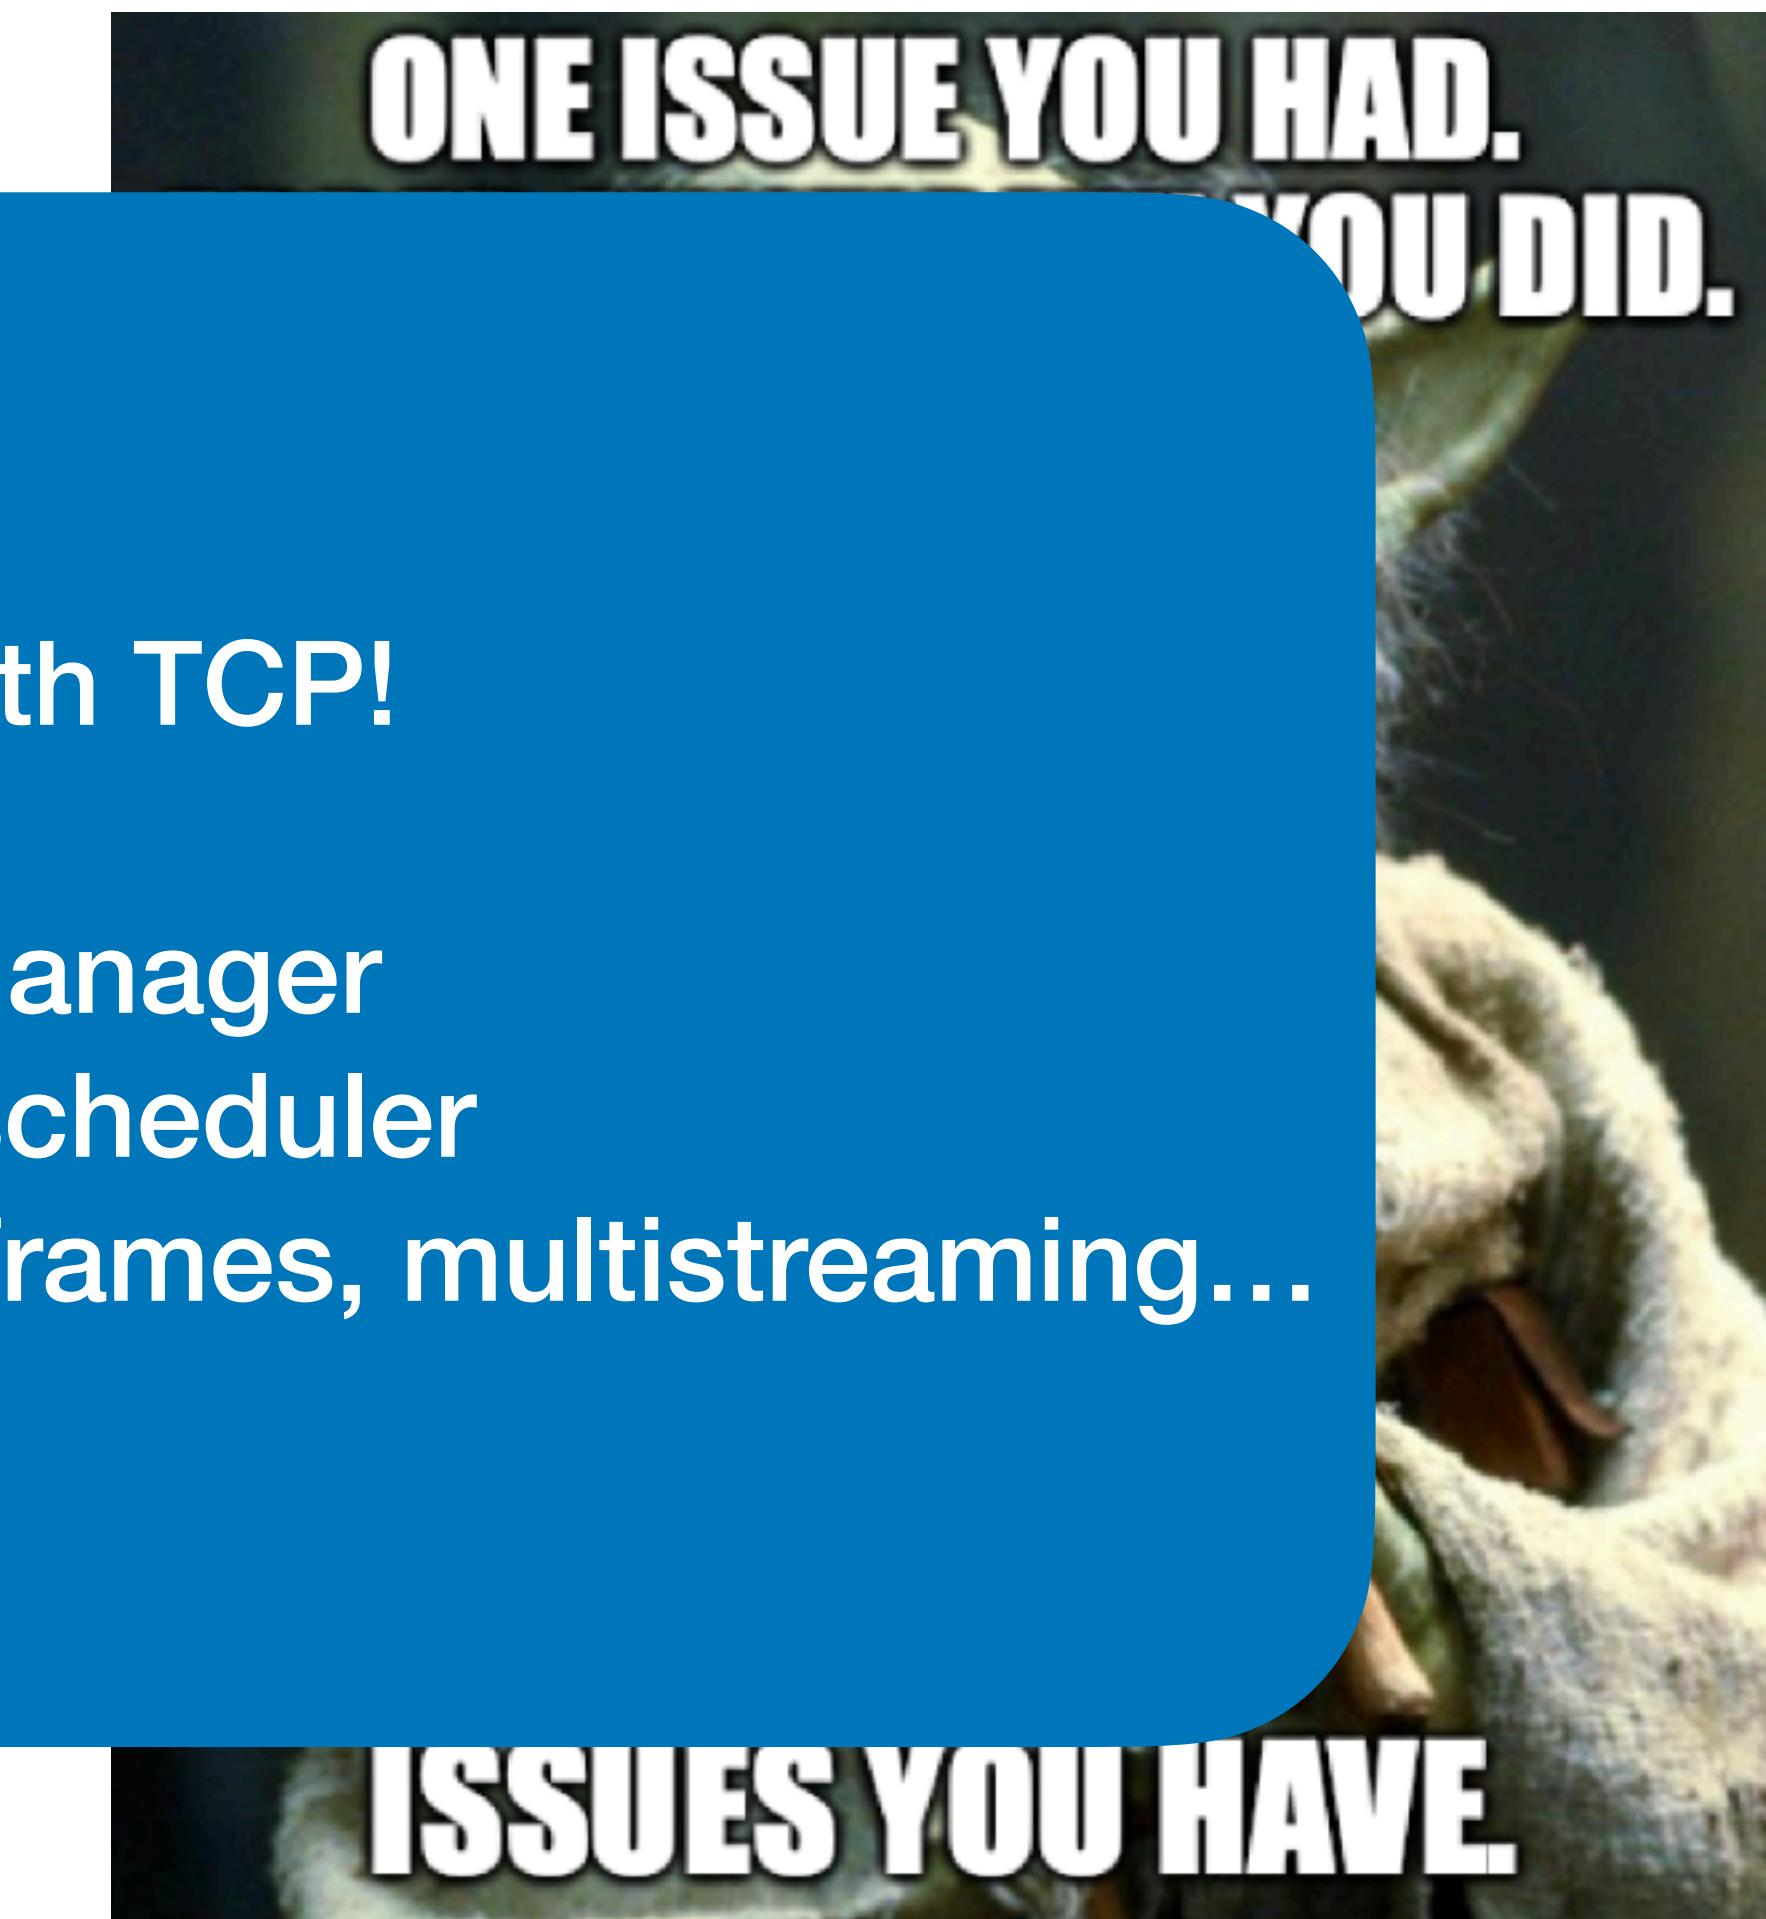
Path identification using Connection IDs
Each path has its own numbering space



Connection IDs provided through
`NEW_CONNECTION_ID` frames

Multipath Algorithms

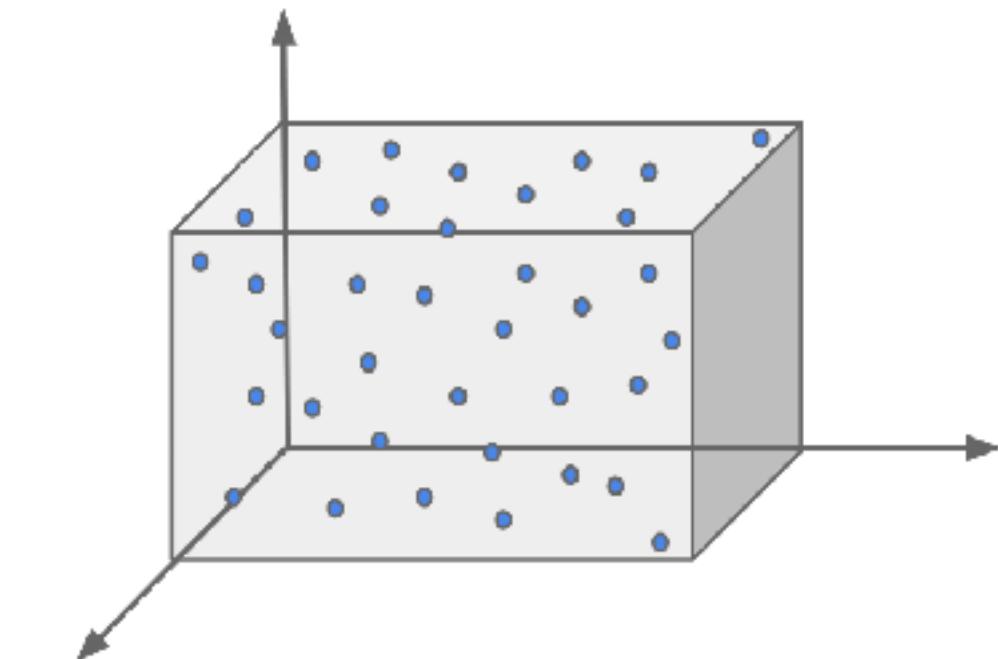
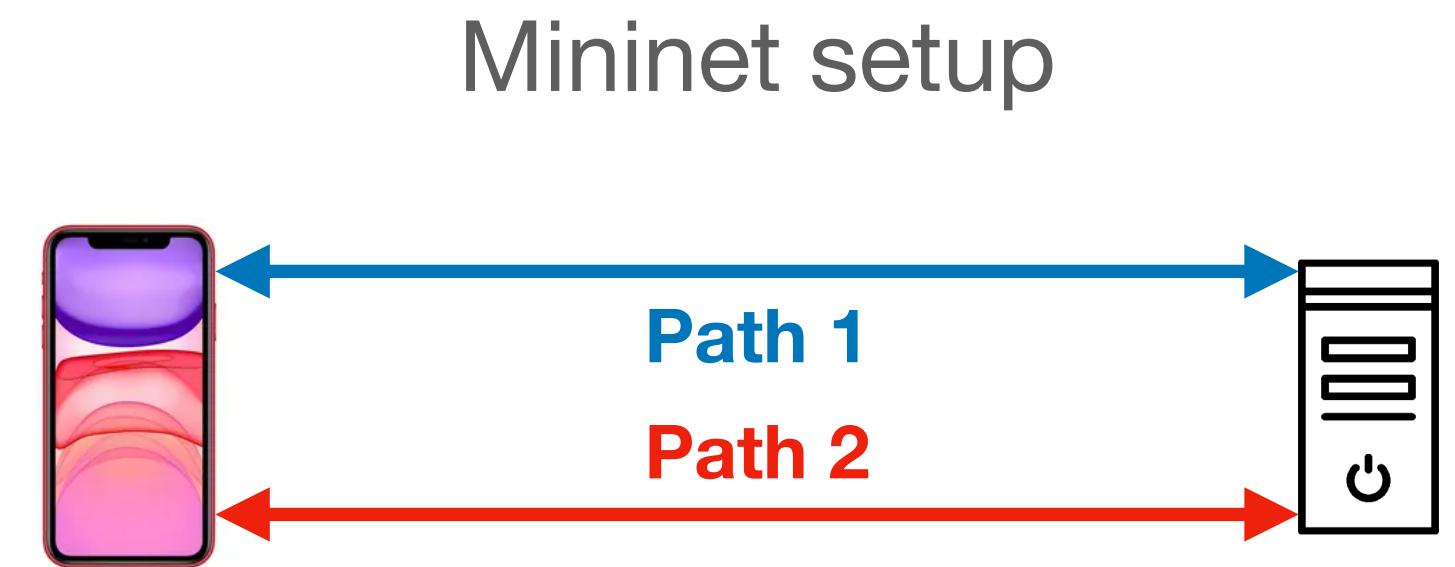
- Path manager
 - Creating paths
 - Packet scheduler
 - Using per-path queueing
 - (Congestion control)
 - Per-path拥塞控制
- Same as Multipath TCP!
- Fullmesh path manager
Lowest-RTT first scheduler
! Also need to schedule control frames, multistreaming...



Revisiting the File Downloader

- Download 20 MB on single stream
 - 9 runs, take median download completion time
 - Compare Linux (Multipath) QUIC with (Multipath) QUIC

Factor	Minimum	Maximum
Capacity [Mbps]	1	50
RTT [ms]	0	50
Queuing Delay [ms]	0	100
Random Loss [%]	0	2.5



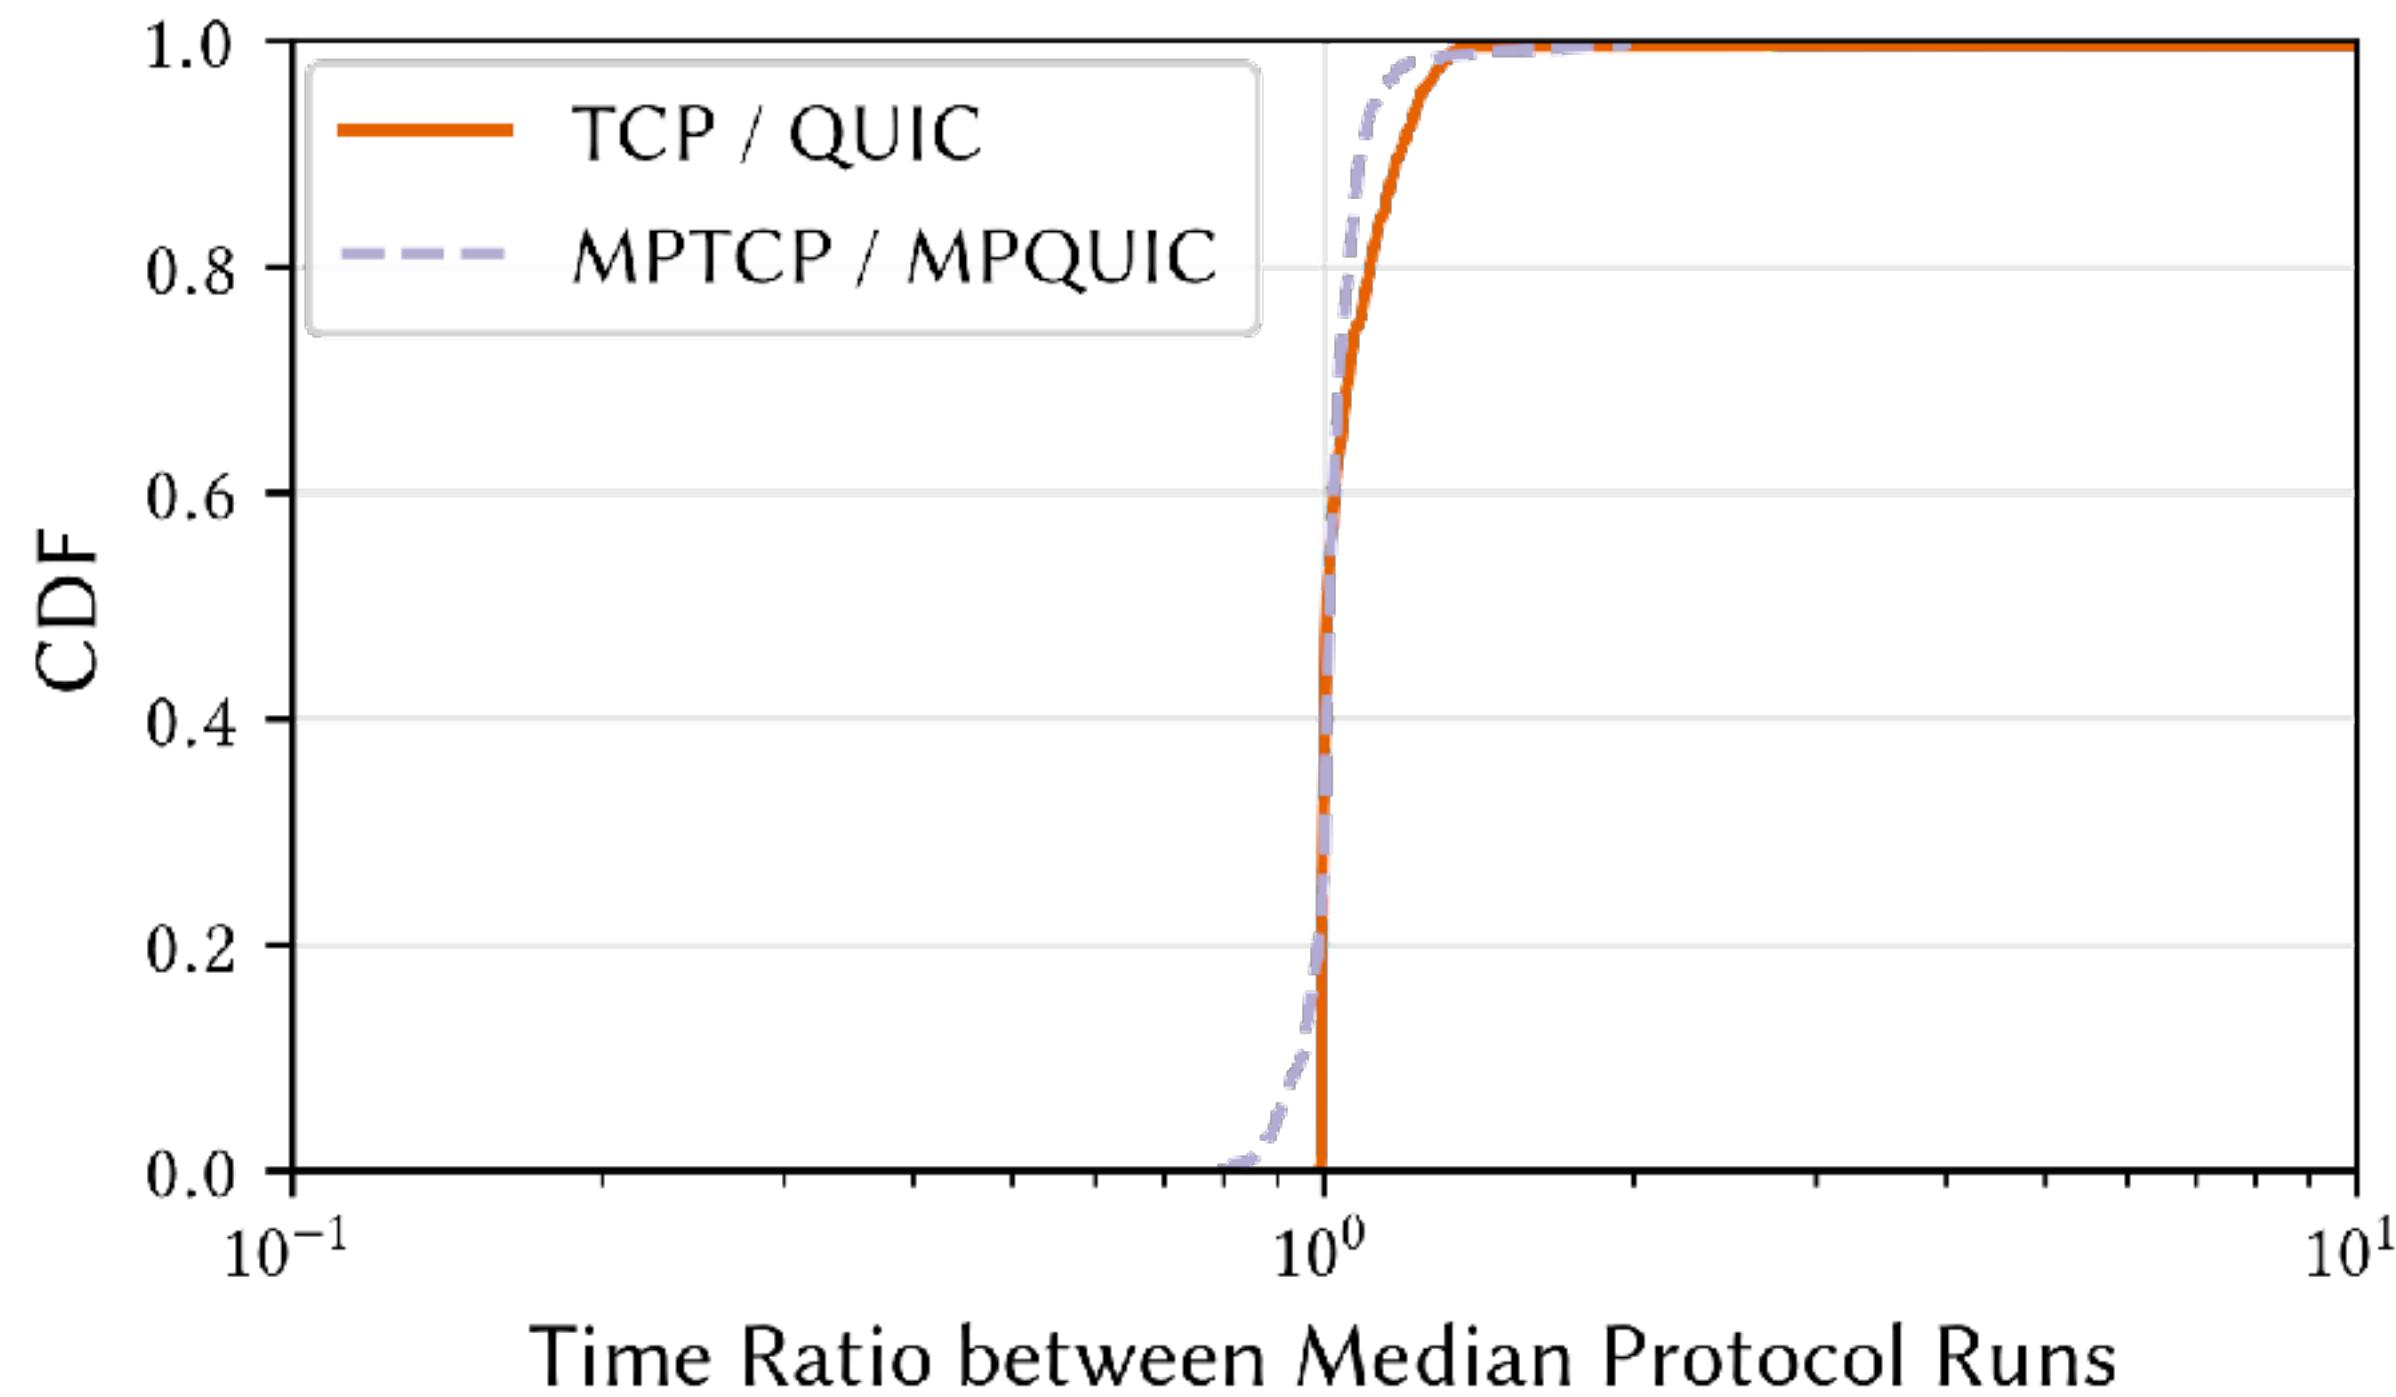
253 network scenarios
Vary initial path

Speedup Ratio of QUIC vs. TCP

Completion time (MP)TCP / Completion Time (MP)QUIC

Fullmesh path manager
Lowest-latency first scheduler

No random loss case

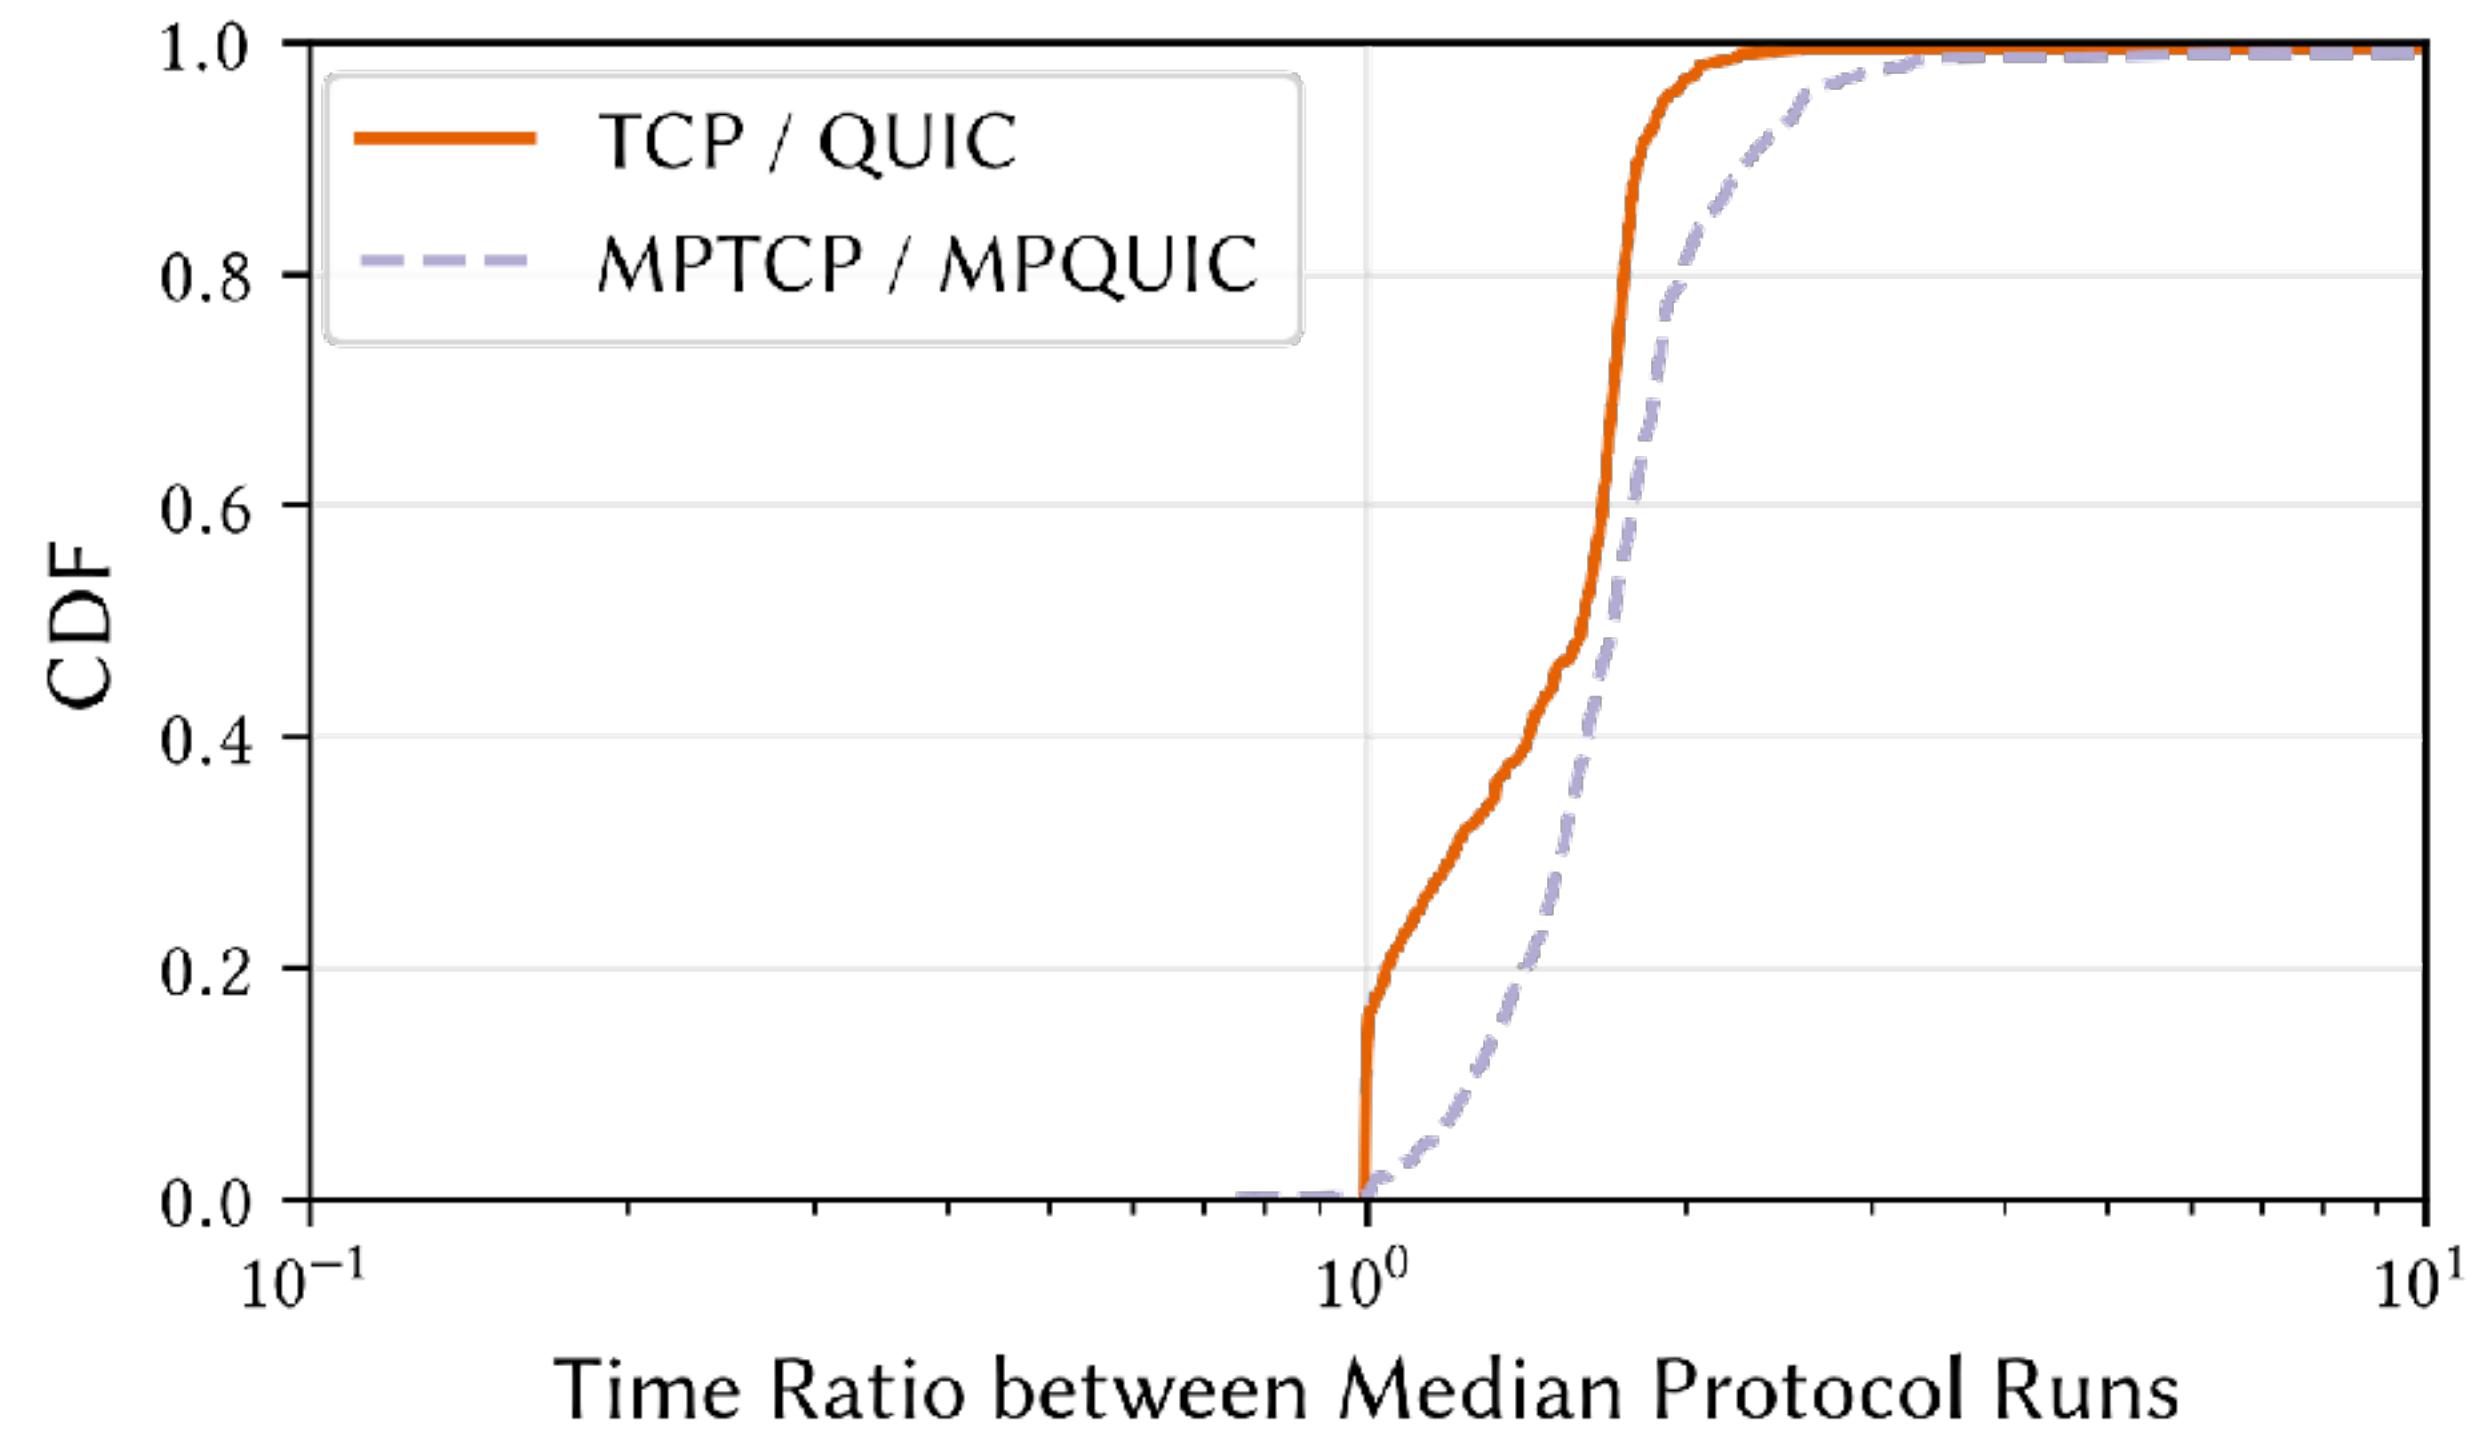


Same algorithms, same results
Hanshake slightly faster in QUIC

Speedup Ratio of QUIC vs. TCP

Completion time (MP)TCP / Completion Time (MP)QUIC

With random losses

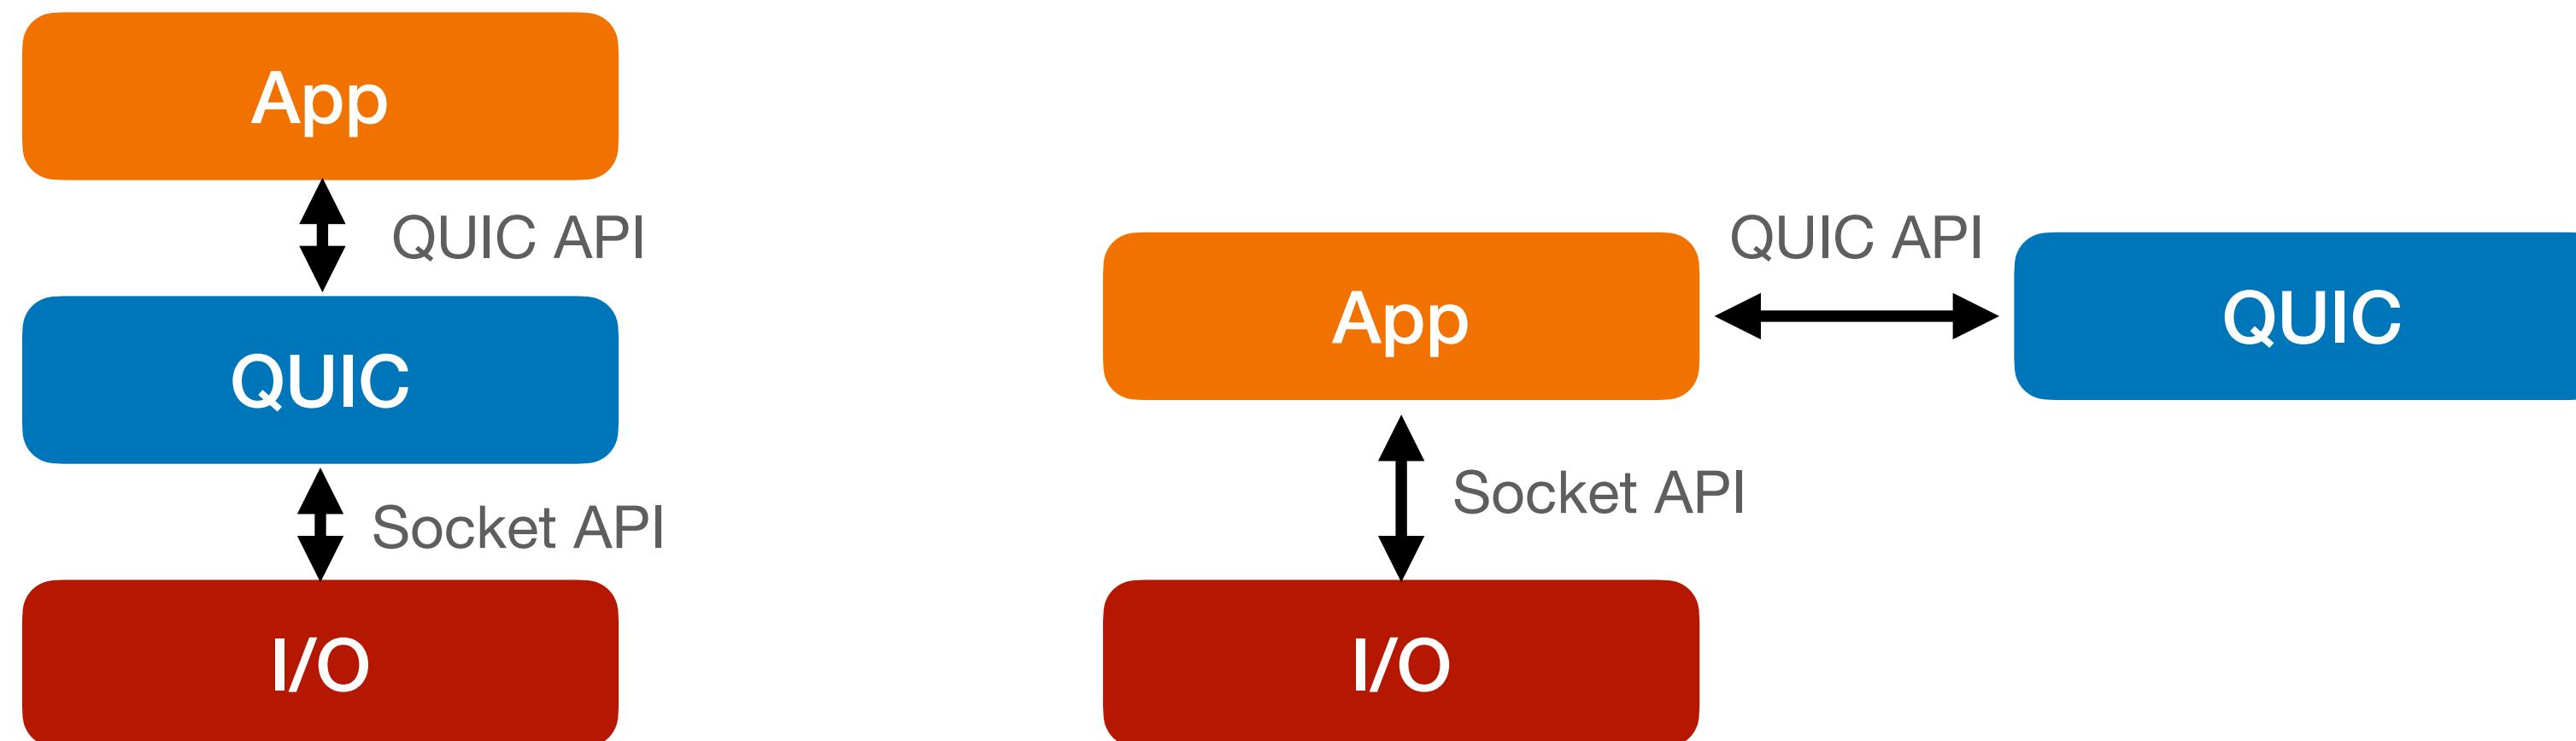


Fullmesh path manager
Lowest-latency first scheduler

- QUIC ACK range (>400 entries) >>> TCP SACK range (2-3 due to TCP options)
- QUIC does not need to (re)transmit original packet if delivered on another path

What About Short Files?

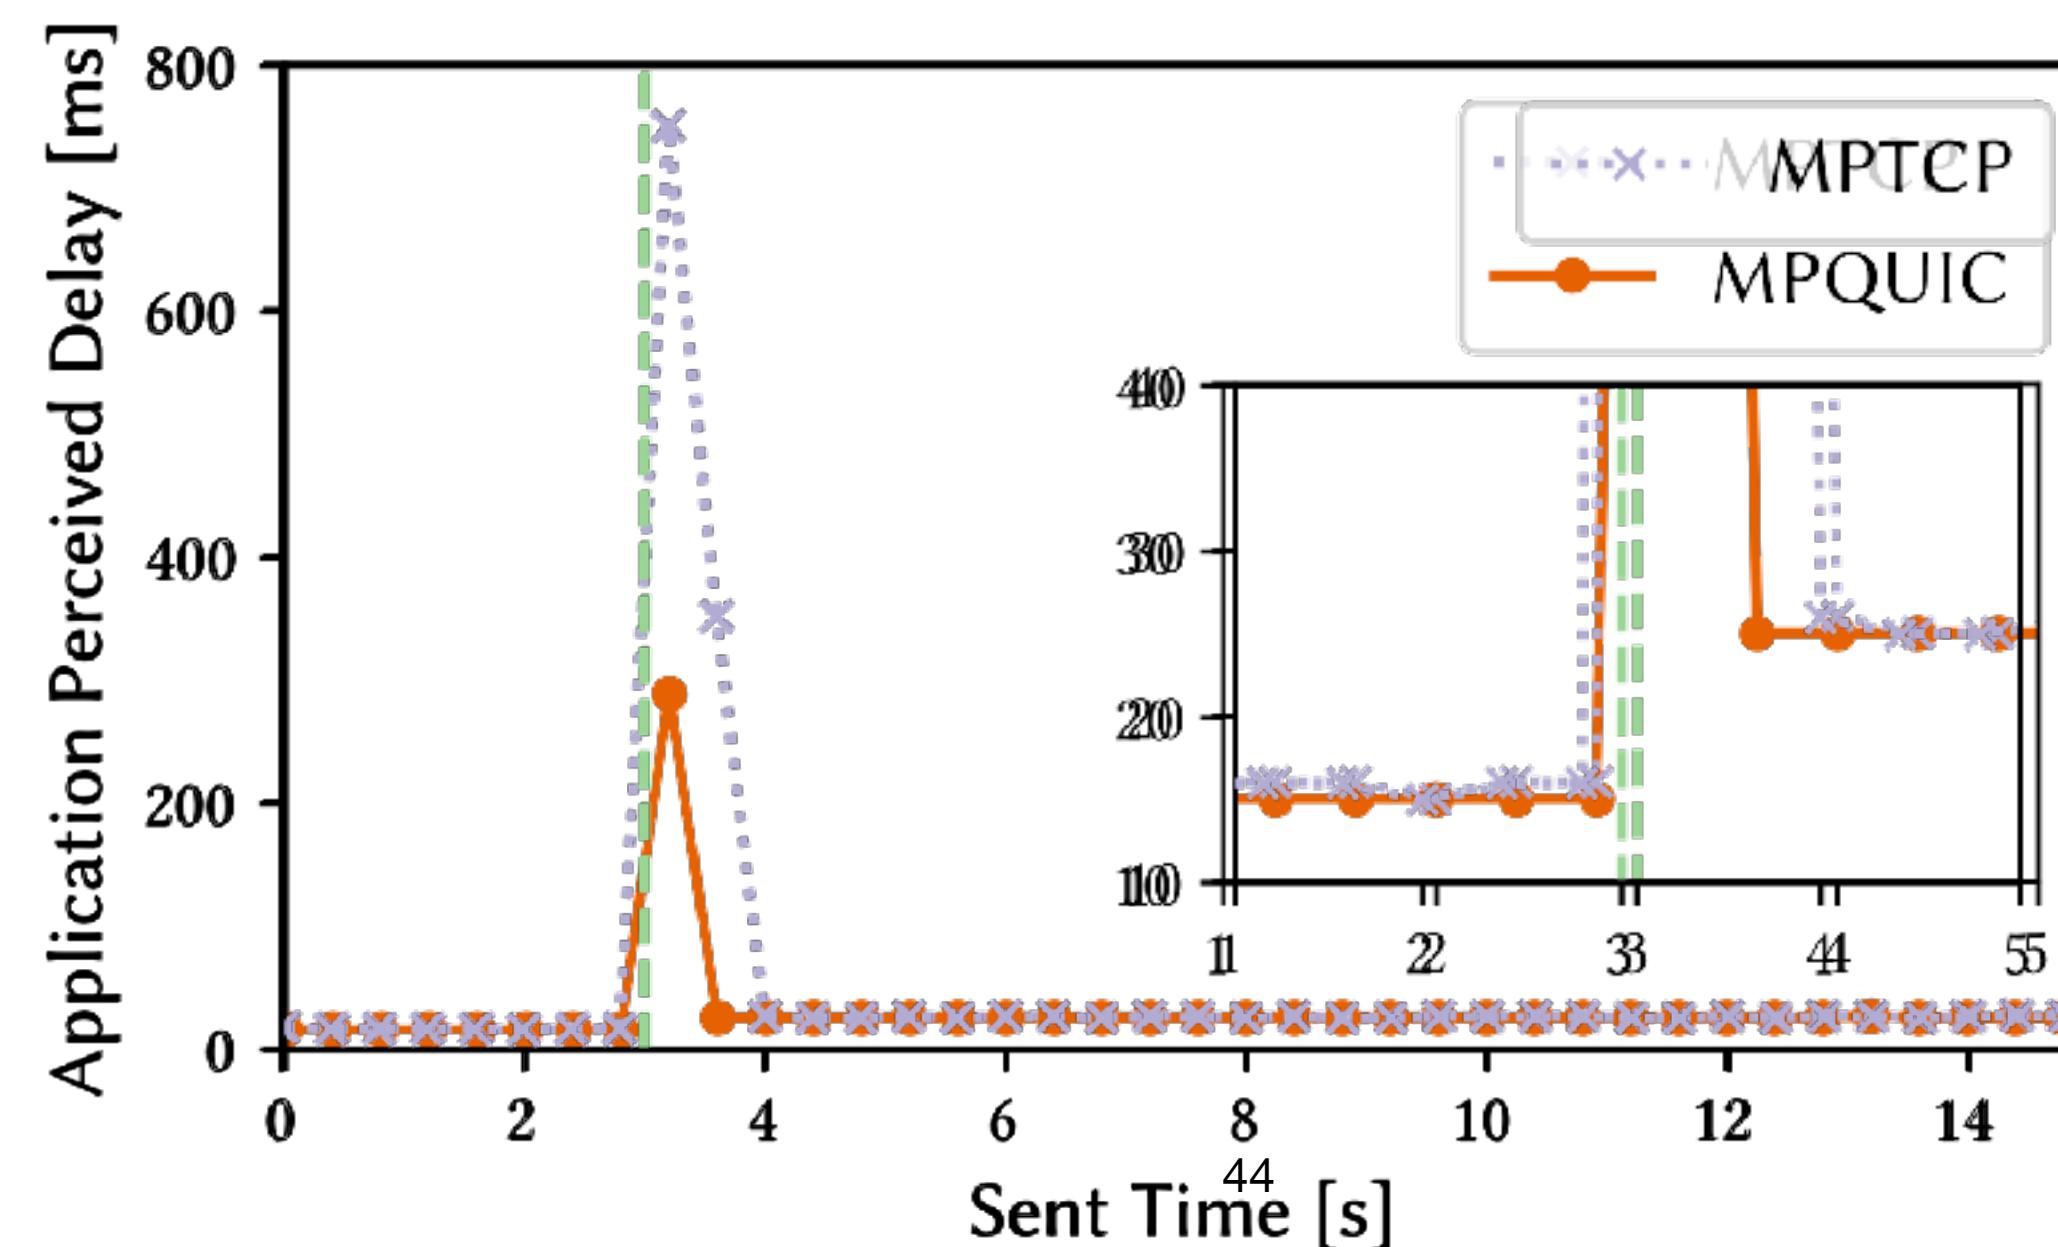
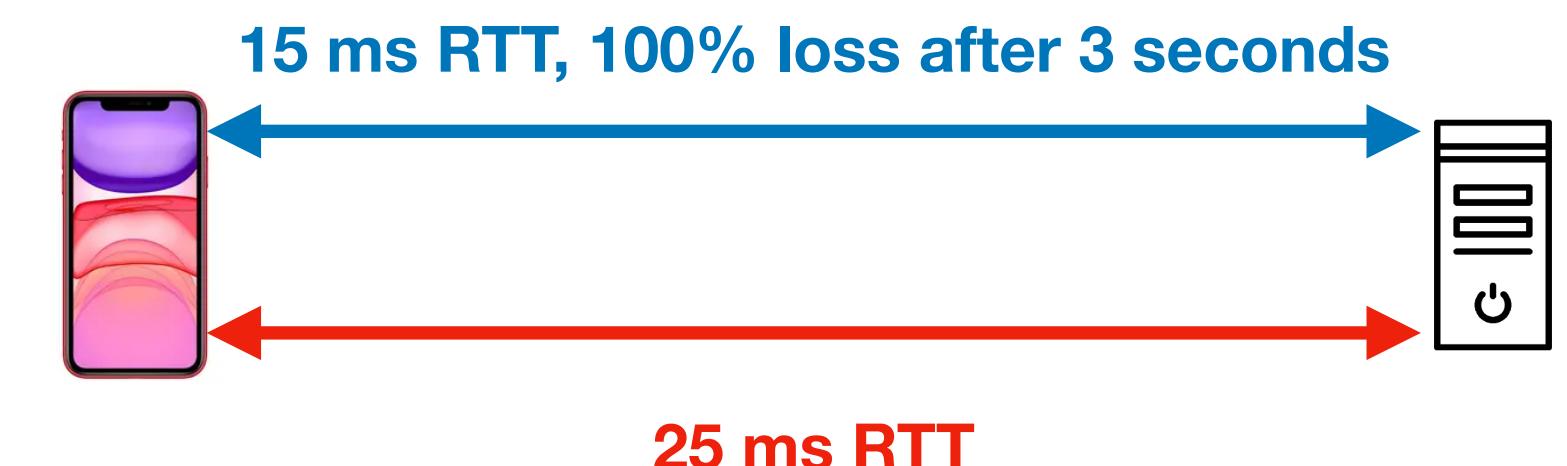
- If a slow path is used at the end of a transfer, still have the issue
 - But application can more easily integrate/tweak the MPQUIC scheduler
 - If transfer size is known, avoid using the slow path when not advantageous



Revisiting the Interactive Traffic with Mobility With Multipath QUIC

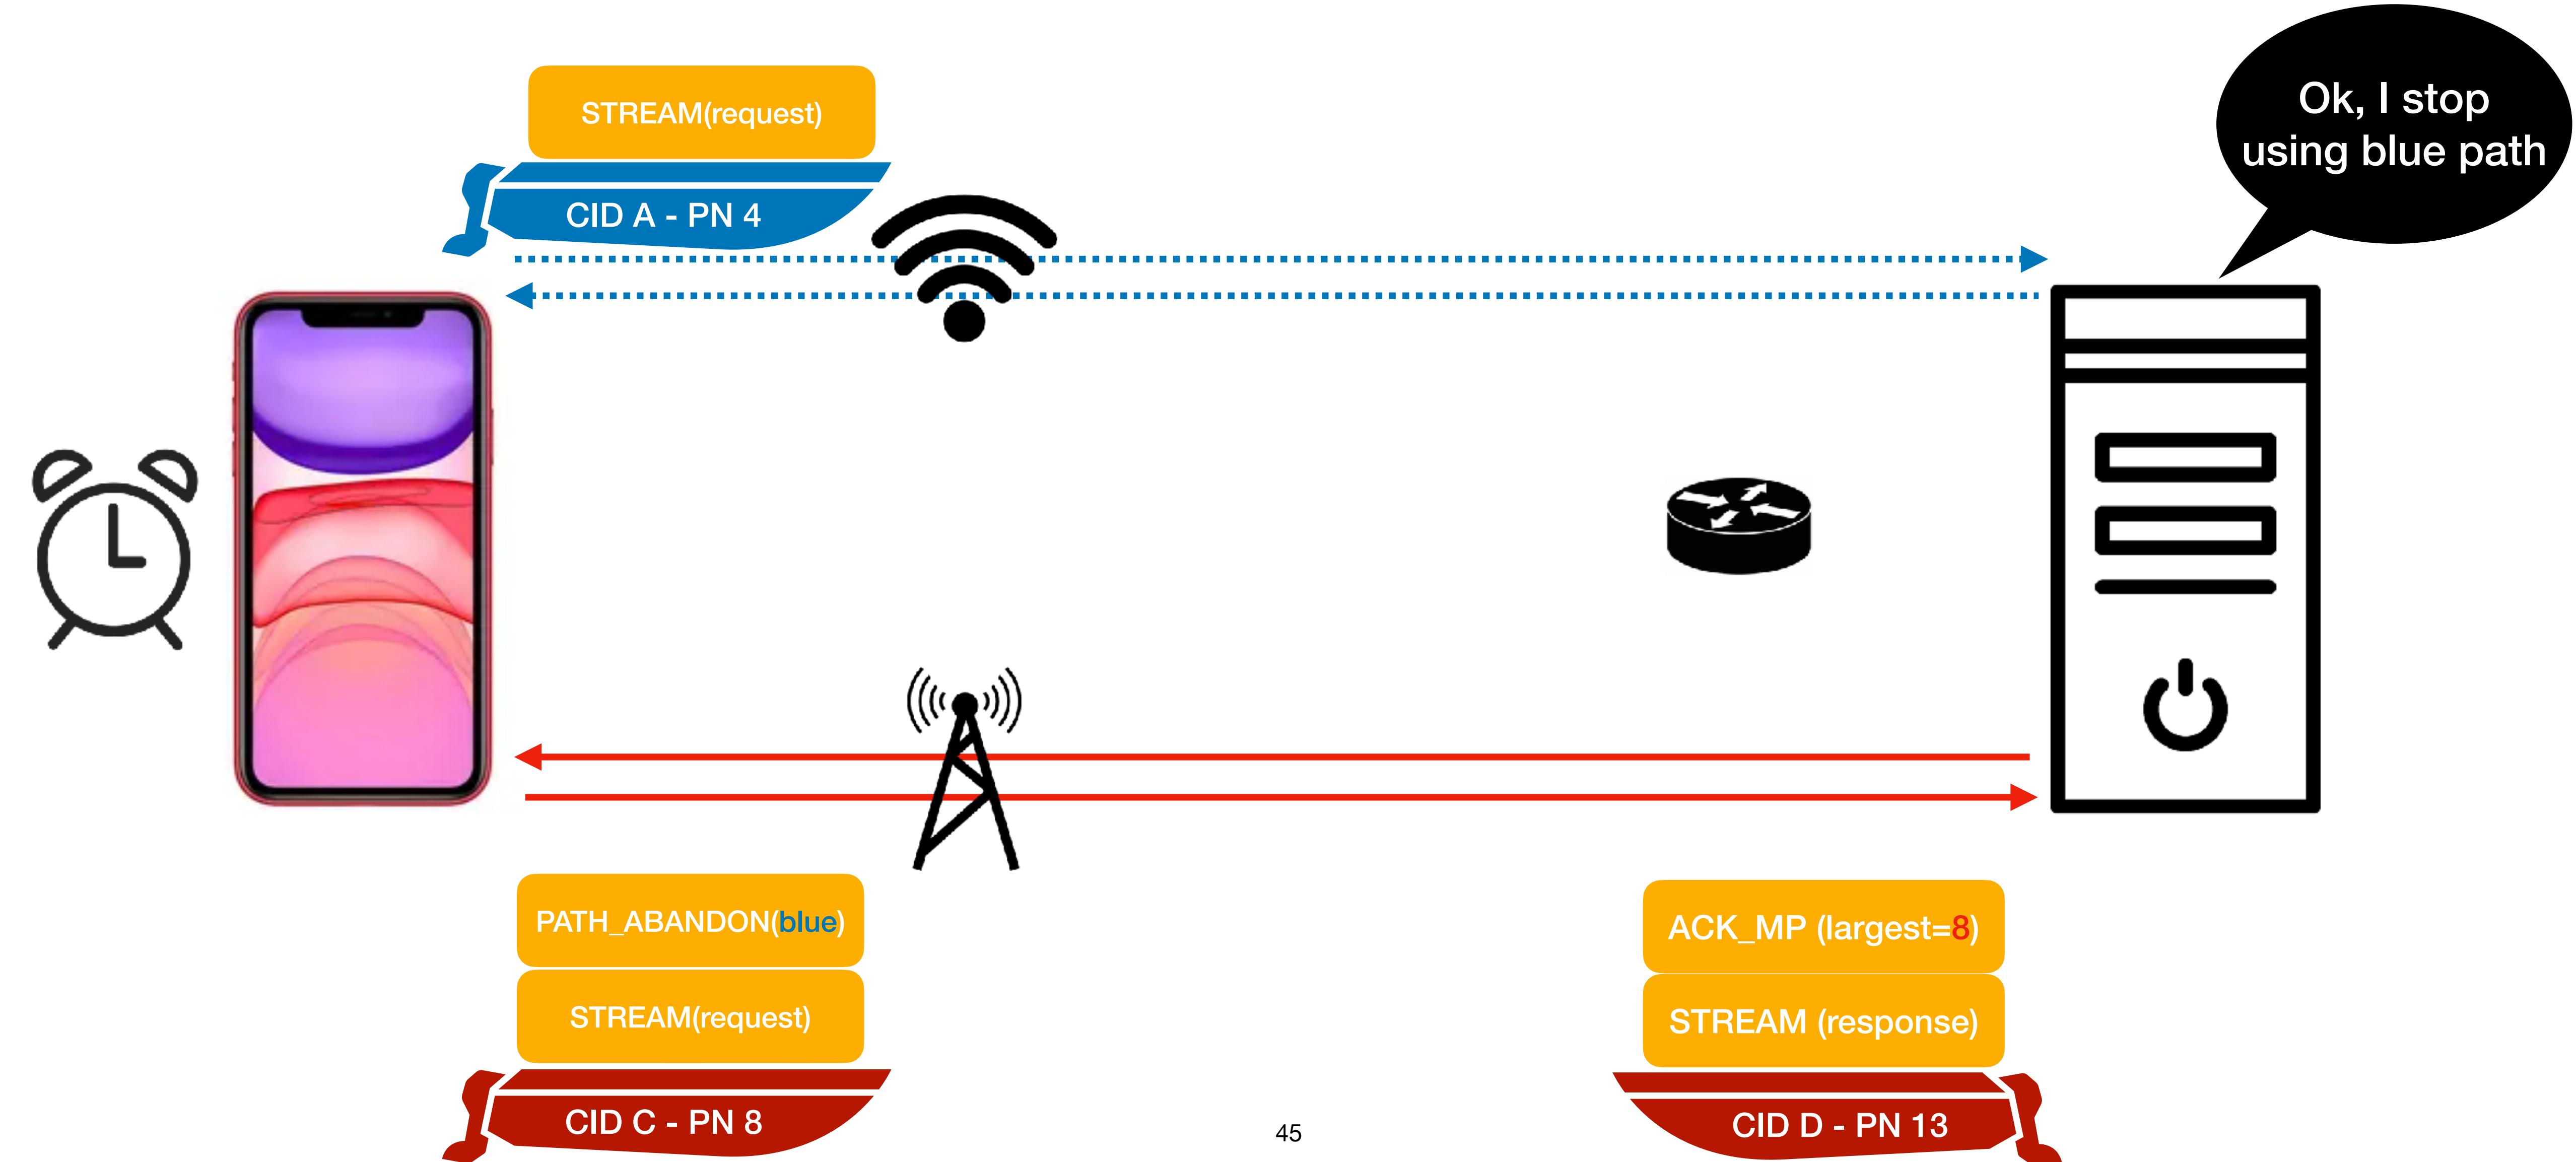
Request/Response traffic

- 750 bytes request/response every 400 ms
- Measure application delay seen by client



Understanding the Difference

Fullmesh path manager
Lowest-latency first scheduler



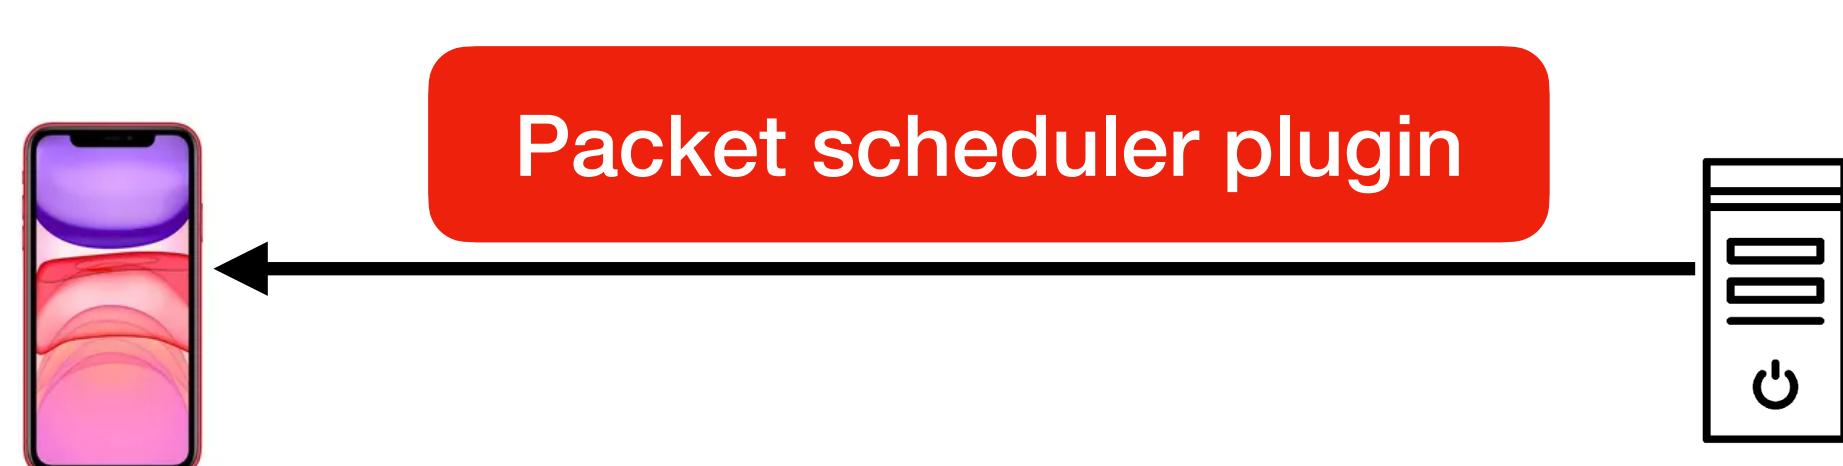
Supporting Different Use Cases?

Cannot
fully control
client side



QUIC

- Closer application integration \neq immediate deployment
- What about letting applications tweak their network stack on a per-connection basis?
- Pluginize your protocol (e.g., QUIC)



WEBASSEMBLY

And What About Video Streaming?

What About Video Streaming?

- Application has both aggregation bandwidth + latency/buffer constraints
 - Especially (near-)live streaming



Where Multipath Transport Can Help

- Audio streaming
 - Handle mobile cases with multiple paths
- Pre-chunked videos (e.g., DASH)
 - Higher video quality through adaptive rate algorithms [CoNext'16]
 - Resilience to network quality changes

Where Multipath Transport Can Help

- Multipath-aware video coding
 - I-frames delivery more important than B- and P-ones
 - Multistreaming of (Multipath) QUIC can help here
 - Doing QUIC Forward Erasure Correction [Networking'19] on side paths

Conclusion



IETF 115 - November 7th, 2022

- Multipath transports now available (more than 10 years of efforts)
- Good application performance requires algorithms adaptations
 - Bandwidth aggregation goal \neq lowest-latency goal
 - Video streaming: identify your goals & implement algorithms
 - Multipath QUIC more powerful/deployable than Multipath TCP
 - At least 3 open-source implementations (Cloudflare's quiche, picoquic,...)